

## Machine Learning & Deep Learning

# 机器学习和深度学习

Machine Learning & Deep Learning

- 人工智能初览
- 深度学习基础
- 初探机器学习
- 深度学习进阶神经网络
- 机器学习基础算法
- 深度学习核心卷积神经网络
- 机器学习进阶算法
- 深度学习网络架构
- 机器学习实战项目
- 机器学习框架TensorFlow应用

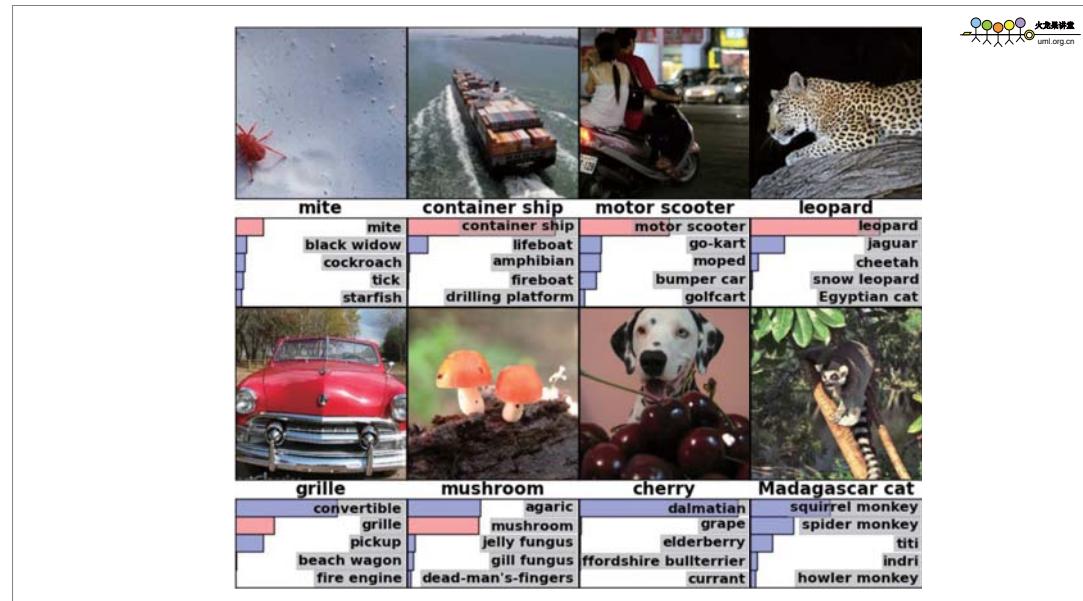
## Gifts from Google and Line



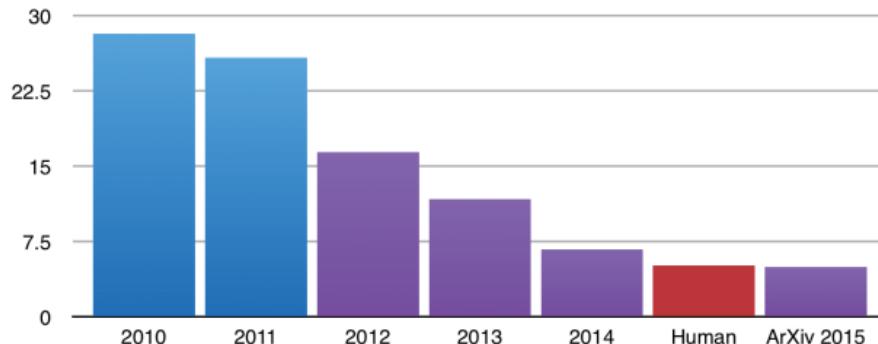
Shocking IT NEWS last years?  
人工智能初览



# Master (AI) vs. Ke Jie 9p

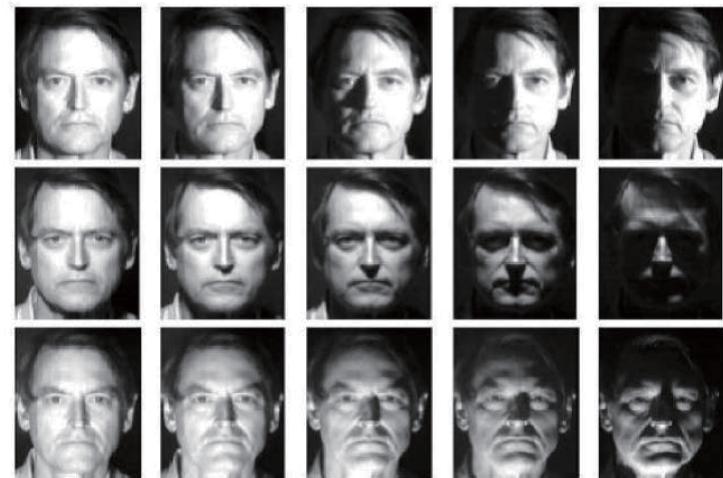


## ILSVRC top-5 error on ImageNet



<https://devblogs.nvidia.com/parallelforall/page/5/>

## Computer Vision is Hard: Illumination Variability



Massachusetts  
Institute of  
Technology

References: [66]

Course 6.S094:  
Deep Learning for Self-Driving Cars

Lex Fridman:  
fridman@mit.edu

Website:  
cars.mit.edu

January  
2017

## Computer Vision is Hard: Pose Variability and Occlusions



Figure 1. The deformable and truncated cat. Cats exhibit (al-

Parkhi et al. "The truth about cats and dogs." 2011.



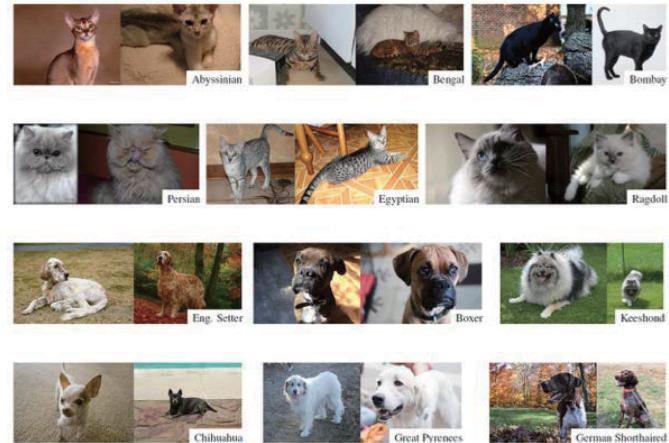
References: [69]

Course 6.S094:  
Deep Learning for Self-Driving Cars

Lex Fridman:  
fridman@mit.edu

Website:  
cars.mit.edu

## Computer Vision is Hard: Intra-Class Variability



Parkhi et al. "Cats and dogs." 2012.



Massachusetts  
Institute of  
Technology

References: [70]

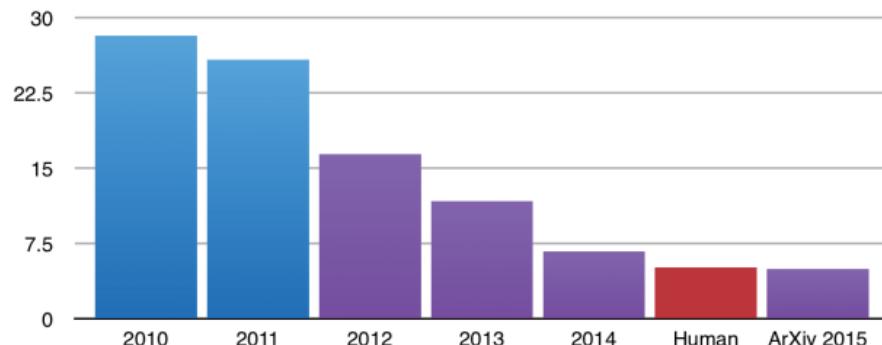
Course 6.S094:  
Deep Learning for Self-Driving Cars

Lex Fridman:  
fridman@mit.edu

Website:  
cars.mit.edu

January  
2017

## ILSVRC top-5 error on ImageNet



<https://devblogs.nvidia.com/parallelforall/page/5/>

# MACHINE LEARNING



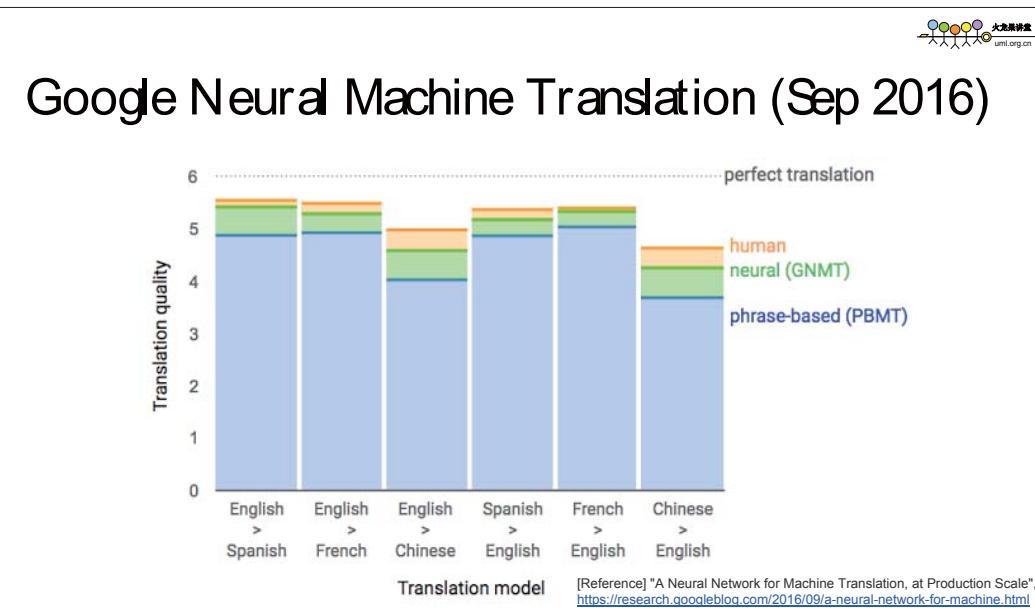
大白鲸讲堂  
uml.org.cn

# Machine Learning is Already Affecting Your World 初探机器学习

<http://respondr.io/7-ways-machine-learning-is-already-affecting-your-world/>

如果要建造一艘船，不要一起鼓勵人們收集  
木材，  
不要分配任務和工作，  
而應該教他們漫長的海洋無限遠。

The screenshot shows the Google Translate interface. At the top, there are language selection buttons for English, Korean, Spanish, Detect language, and a dropdown menu. Below this, a quote in English is displayed: "If you want to build a ship, don't drum up people together to collect wood and don't assign them tasks and work, but rather teach them to long for the endless immensity of the sea." The Chinese translation is shown below it: "如果要建造一艘船，不要一起鼓勵人們收集木材，不要分配任務和工作，而應該教他們漫長的海洋無限遠。". A small note at the bottom indicates the Chinese text is a direct translation of the English sentence.



### TalkType Voice Keyboard

Baidu Research Productivity

★★★★★ 227

This app is incompatible with your device.

Add to Wishlist Install

<https://play.google.com/store/apps/details?id=com.baidu.research.talktype&hl=en>

### Lose It!

FitNow, Inc. Health & Fitness

Top Developer

★★★★★ 55,021

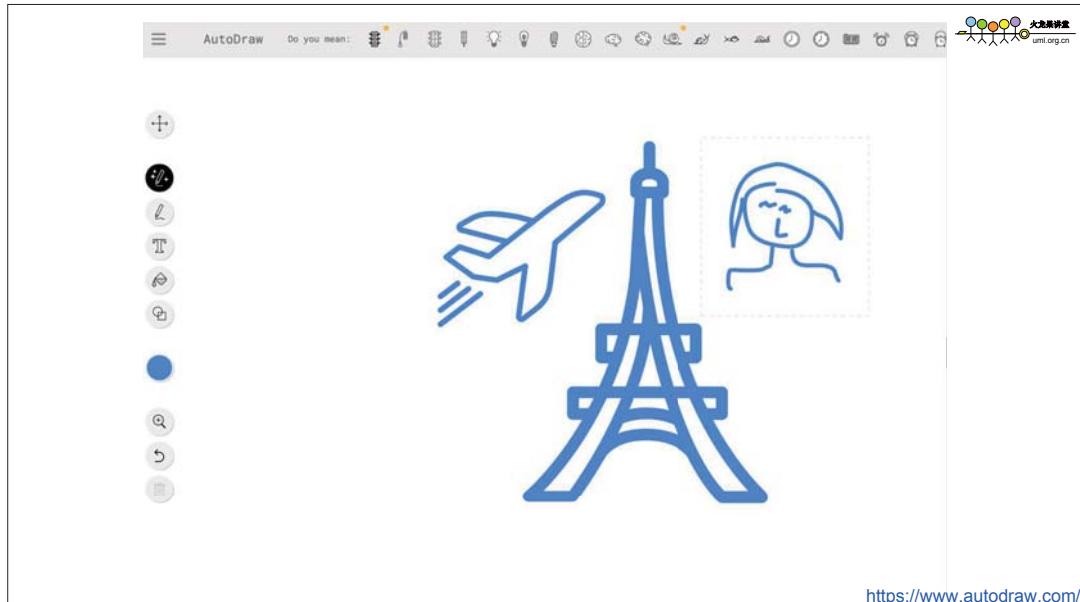
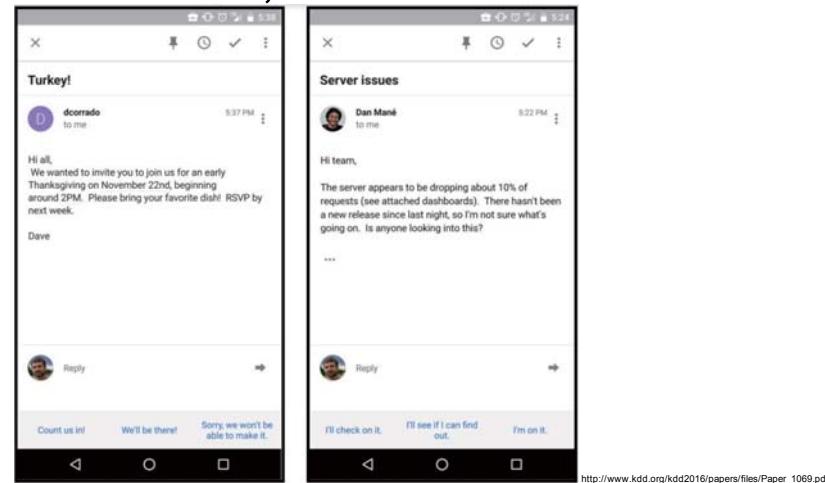
Offers in-app purchases

This app is compatible with your device.

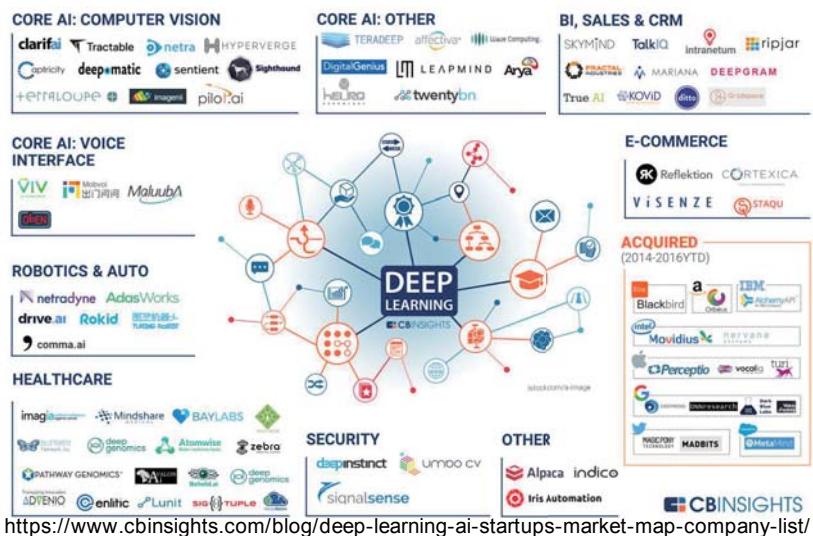
Add to Wishlist Install

<https://play.google.com/store/apps/details?id=com.fitnow.loseit&hl=en>

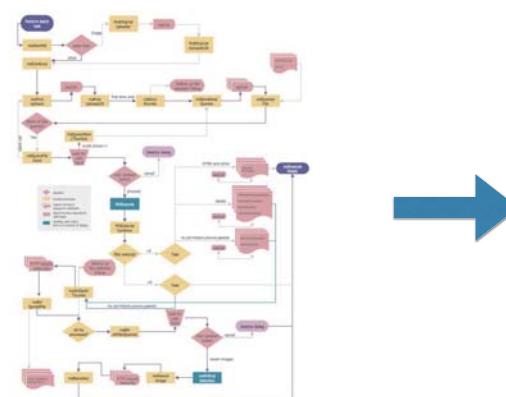
## Smart Reply: Promising results Seen 78% Used 31.9%



## 60+ STARTUPS USING DEEP LEARNING



## Rule based programming VS data driven learning



## How can we develop?



Do you mean:

Google

Translate

English Korean Spanish Detect language

如果要建造一艘船，不要一起鼓勵人們收集木材，  
不要分配任務和工作，而應該教他們漫長的海洋無  
限遠。

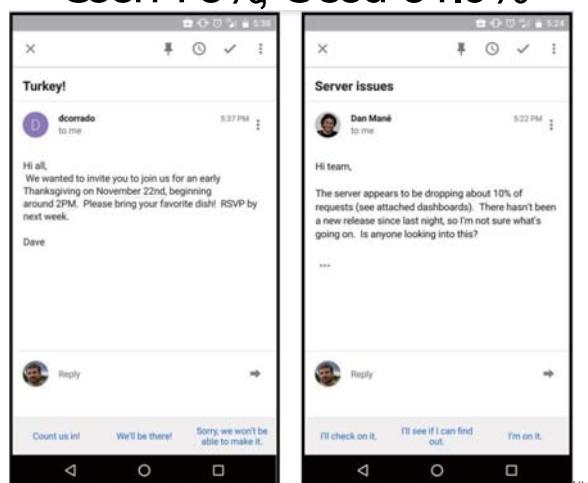
Rúguò yào jiànzhào yí sōu chuán, bùyào yīqǐ gǔlì rénmen shōují mùcái, bùyào fēnpéi rènwù hé  
gōngzuò, ér yǐnggāi jiāo tāmen mǎncháng dì hǎiyáng wúxiànl yuǎn.

Turn off instant translation

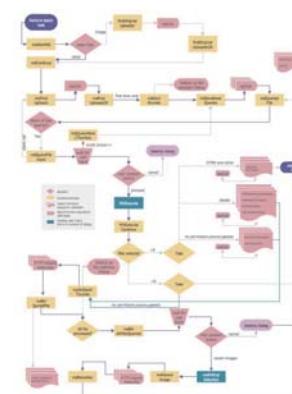
Korean Chinese (Traditional) English

Translate

## Smart Reply: Promising results Seen 78%, Used 31.9%



## Rule based programming VS data driven learning





TensorFlow™ [Install](#) [Develop](#) [API 1.0](#) [Deploy](#) [Extend](#) [Resources](#) [Versions](#) [Search](#) [GitHub](#)

An open-source software library for Machine Intelligence

[GET STARTED](#)

TensorFlow 1.0 has arrived!

We're excited to announce the release of TensorFlow 1.0! Check out the migration guide to upgrade your code with ease.

[UPGRADE NOW](#)

Dynamic graphs in TensorFlow

We've open-sourced TensorFlow Fold to make it easier than ever to work with input data with varying shapes and sizes.

[LEARN MORE](#)

The 2017 TensorFlow Dev Summit

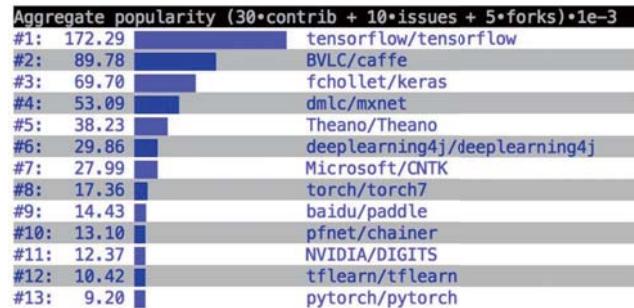
Thousands of people from the TensorFlow community participated in the first flagship event. Watch the keynote and talks.

[WATCH VIDEOS](#)

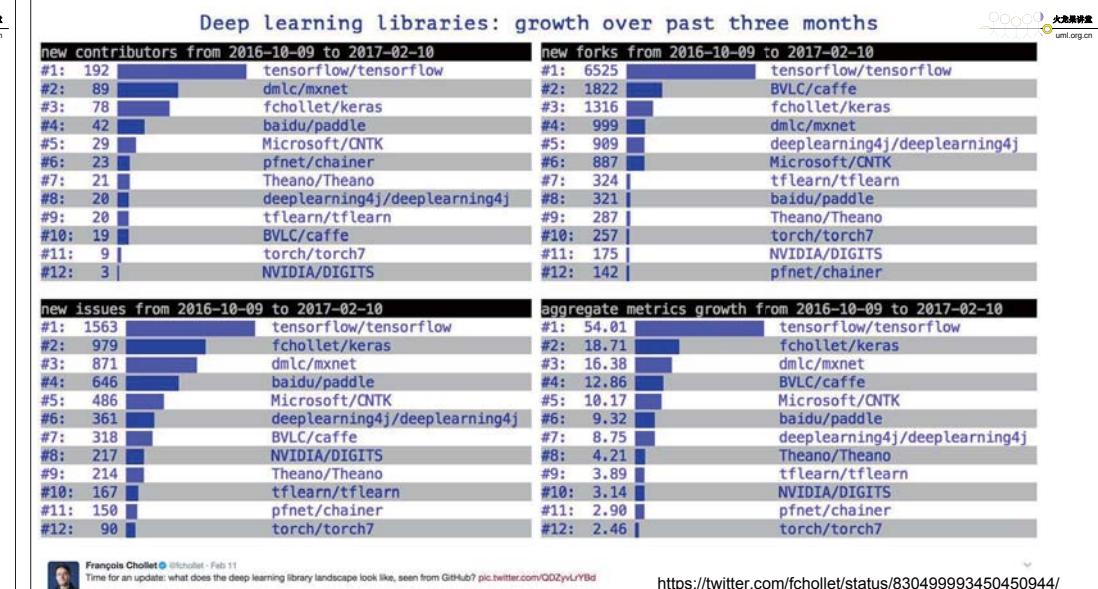
<https://www.tensorflow.org>

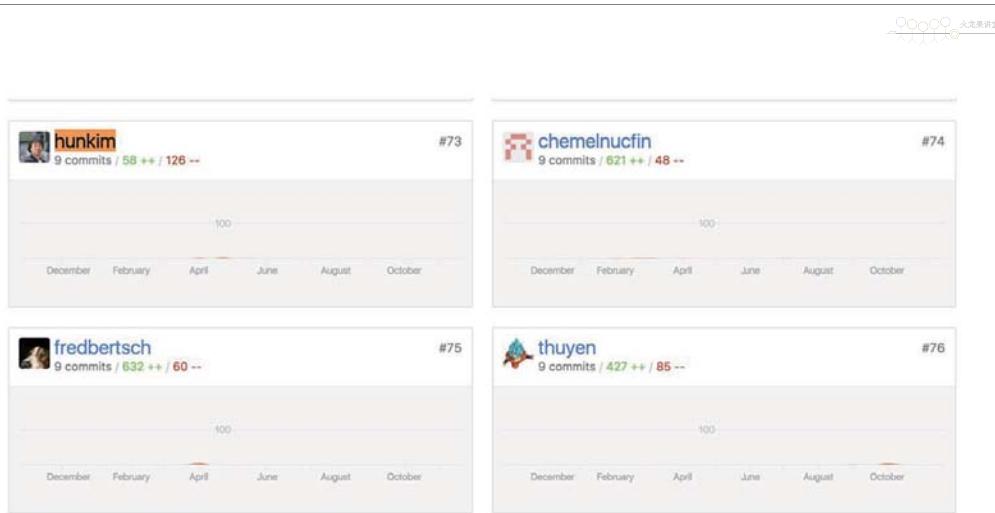
# TensorFlow

## Deep learning libraries: Accumulated GitHub metrics



<https://twitter.com/fchollet/status/830499993450450944/>





## TensorFlow 机器学习框架

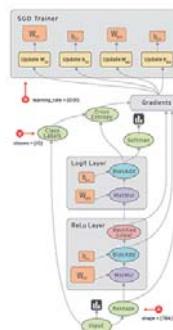
- TensorFlow™ is an open source software library for numerical computation using data flow graphs.
- Python



<https://www.tensorflow.org/>

## What is a Data Flow Graph?

- Nodes in the graph represent mathematical operations
- Edges represent the multidimensional data arrays (tensors) communicated between them.



<https://www.tensorflow.org/>

## Installing TensorFlow

- Linux, Mac OSX, Windows
  - (sudo -H) pip install --upgrade tensorflow
  - (sudo -H) pip install --upgrade tensorflow-gpu
- From source
  - bazel ...
  - [https://www.tensorflow.org/install/install\\_sources](https://www.tensorflow.org/install/install_sources)
- Google search/Community help

<https://www.tensorflow.org/install/>

## Check installation and version

```
$ python3  
Python 3.6.0 (v3.6.0:41df79263a11, Dec 22 2016, 17:23:13)  
[GCC 4.2.1 (Apple Inc. build 5666) (dot 3)] on darwin  
Type "help", "copyright", "credits" or "license" for more  
information.  
>>> import tensorflow as tf  
>>> tf.__version__  
'1.0.0'  
>>>
```

## TensorFlow Hello World!

### Hello TensorFlow!

```
In [2]: # Create a constant op  
# This op is added as a node to the default graph  
hello = tf.constant("Hello, TensorFlow!")  
  
# start a TF session  
sess = tf.Session()  
  
# run the op and get result  
print(sess.run(hello))
```

b'Hello, TensorFlow!'

b'String' 'b' indicates Bytes literals. <http://stackoverflow.com/questions/6269765/>

[lab-01-basics.ipynb](#)

## Computational Graph

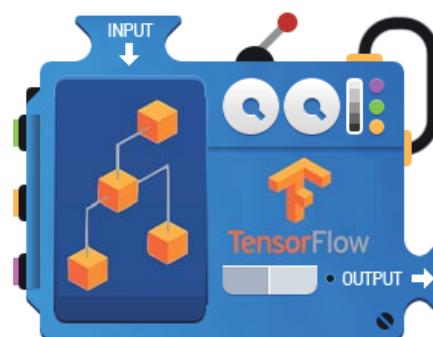
```
In [4]: node1 = tf.constant(3.0, tf.float32)  
node2 = tf.constant(4.0) # also tf.float32 implicitly  
node3 = tf.add(node1, node2)  
  
In [5]: print("node1:", node1, "node2:", node2)  
print("node3: ", node3)  
  
node1: Tensor("Const_1:0", shape=(), dtype=float32) node2: Tensor("Const_2:0", shape=(), dtype=float32)  
node3: Tensor("Add:0", shape=(), dtype=float32)  
  
In [6]: sess = tf.Session()  
print("sess.run(node1, node2): ", sess.run([node1, node2]))  
print("sess.run(node3): ", sess.run(node3))  
  
sess.run(node1, node2): [3.0, 4.0]  
sess.run(node3): 7.0
```

[lab-01-basics.ipynb](#)



## TensorFlow Mechanics

- feed data and run graph (operation)  
`sess.run (op)`



- Build graph using TensorFlow operations

- update variables in the graph (and return values)

[WWW.MATHWAREHOUSE.COM](http://WWW.MATHWAREHOUSE.COM)

# Computational Graph

(1) Build graph (tensors) using TensorFlow operations

```
In [4]: node1 = tf.constant(3.0, tf.float32)
node2 = tf.constant(4.0) # also tf.float32 implicitly
node3 = tf.add(node1, node2)
```

(2) feed data and run graph (operation)  
**sess.run (op)**

(3) update variables in the graph  
(and return values)

```
In [6]: sess = tf.Session()
print("sess.run(node1, node2): ", sess.run([node1, node2]))
print("sess.run(node3): ", sess.run(node3))
```

```
sess.run(node1, node2): [3.0, 4.0]
sess.run(node3): 7.0
```

[lab-01-basics.ipynb](#)



# Placeholder

```
In [7]: a = tf.placeholder(tf.float32)
b = tf.placeholder(tf.float32)
adder_node = a + b # + provides a shortcut for tf.add(a, b)

print(sess.run(adder_node, feed_dict={a: 3, b: 4.5}))
print(sess.run(adder_node, feed_dict={a: [1,3], b: [2, 4]}))
```

```
7.5
[ 3.  7.]
```

[lab-01-basics.ipynb](#)

# TensorFlow Mechanics

2 feed data and run graph (operation)  
**sess.run (op, feed\_dict={x: x\_data})**

1 Build graph using TensorFlow operations



WWW.MATHWAREHOUSE.COM

3 update variables  
in the graph  
(and return values)

# Everything is Tensor

## Tensors

```
In [3]: 3 # a rank 0 tensor; this is a scalar with shape []
[1., 2., 3.] # a rank 1 tensor; this is a vector with shape [3]
[[1., 2., 3.], [4., 5., 6.]] # a rank 2 tensor; a matrix with shape [2, 3]
[[[1., 2., 3.]], [[7., 8., 9.]]] # a rank 3 tensor with shape [2, 1, 3]

Out[3]: [[[1.0, 2.0, 3.0]], [[7.0, 8.0, 9.0]]]
```

```
t = tf.Constant([1., 2., 3.])
```

## Tensor Ranks, Shapes, and Types

```
t = [[1, 2, 3], [4, 5, 6], [7, 8, 9]]
```

Rank	Math entity	Python example
0	Scalar (magnitude only)	s = 483
1	Vector (magnitude and direction)	v = [1.1, 2.2, 3.3]
2	Matrix (table of numbers)	m = [[1, 2, 3], [4, 5, 6], [7, 8, 9]]
3	3-Tensor (cube of numbers)	t = [[[2, 4, 6], [8, 10, 12], [14, 16, 18]]]
n	n-Tensor (you get the idea)	....

[https://www.tensorflow.org/programmers\\_guide/dims\\_types](https://www.tensorflow.org/programmers_guide/dims_types)

## Tensor Ranks, Shapes, and Types

```
t = [[1, 2, 3], [4, 5, 6], [7, 8, 9]]
```

Rank	Shape	Dimension number	Example
0	[]	0-D	A 0-D tensor. A scalar.
1	[D0]	1-D	A 1-D tensor with shape [5].
2	[D0, D1]	2-D	A 2-D tensor with shape [3, 4].
3	[D0, D1, D2]	3-D	A 3-D tensor with shape [1, 4, 3].
n	[D0, D1, ..., Dn-1]	n-D	A tensor with shape [D0, D1, ..., Dn-1].

[https://www.tensorflow.org/programmers\\_guide/dims\\_types](https://www.tensorflow.org/programmers_guide/dims_types)

## Tensor Ranks, Shapes, and Types

```
t = [[1, 2, 3], [4, 5, 6], [7, 8, 9]]
```

Data type	Python type	Description
DT_FLOAT	tf.float32	32 bits floating point.
DT_DOUBLE	tf.float64	64 bits floating point.
DT_INT8	tf.int8	8 bits signed integer.
DT_INT16	tf.int16	16 bits signed integer.
DT_INT32	tf.int32	32 bits signed integer.
DT_INT64	tf.int64	64 bits signed integer.
...		

<https://www.quora.com/When-should-I-use-tf-float32-vs-tf-float64-in-TensorFlow>

## TensorFlow Mechanics

- 2 feed data and run graph (operation)  
`sess.run (op, feed_dict={x: x_data})`

- 1 Build graph using  
TensorFlow operations



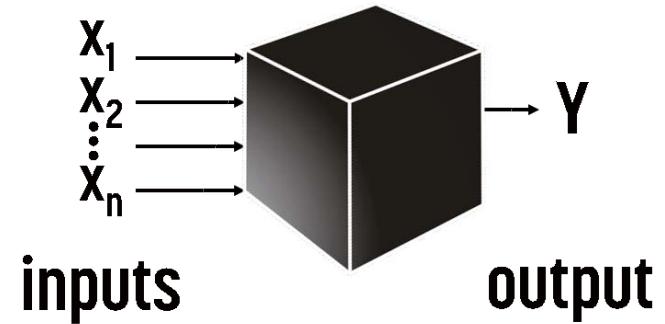
WWW.MATHWAREHOUSE.COM

# Machine Learning Basics

## 机器学习&深度学习 基础与进阶算法

- Linear Regression
- Logistic Regression (Binary classification)
- Softmax Classification
- Neural Networks

## Regression



inputs

output

## Predicting exam score: regression

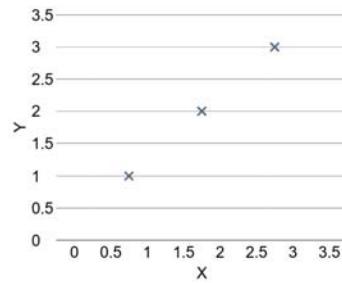
x (hours)	y (score)
10	90
9	80
3	50
2	30

## Regression (data)

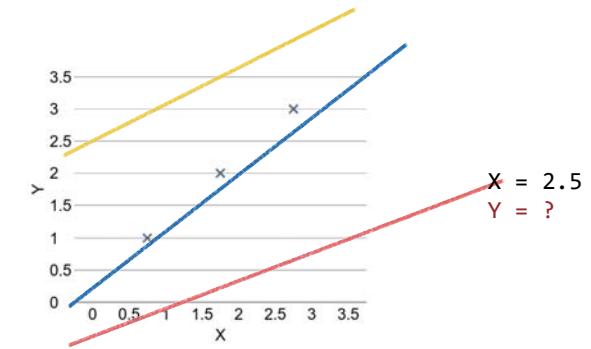
x	y
1	1
2	2
3	3

## Regression (presentation)

x	y
1	1
2	2
3	3

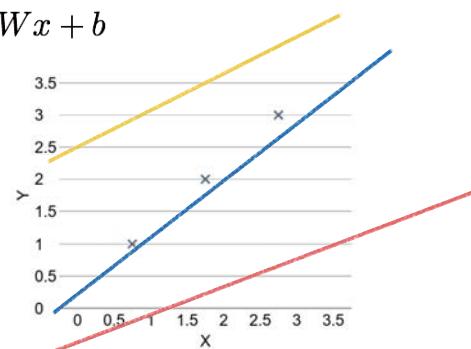


## (Linear) Hypothesis

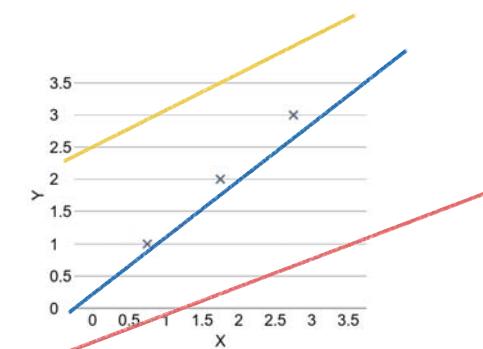


## (Linear) Hypothesis

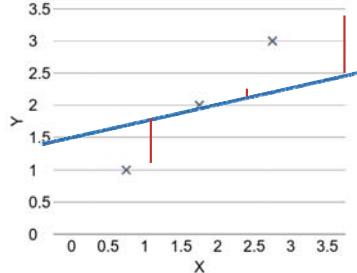
$$H(x) = Wx + b$$



Which hypothesis is better?



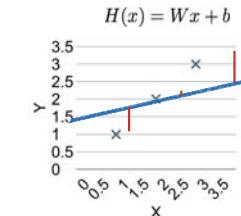
## Which hypothesis is better?



## Cost function

- How fit the line to our (training) data

$$H(x) - y$$



Cost function

- How fit the line to our (training) data

$$\frac{(H(x^{(1)}) - y^{(1)})^2 + (H(x^{(2)}) - y^{(2)})^2 + (H(x^{(3)}) - y^{(3)})^2}{3}$$

$$cost = \frac{1}{m} \sum_{i=1}^m (H(x^{(i)}) - y^{(i)})^2$$

x	y	H(x)
0	1	0
1	2	2
2	3	4

Cost function

$$cost = \frac{1}{m} \sum_{i=1}^m (H(x^{(i)}) - y^{(i)})^2$$

$$H(x) = Wx + b$$

$$cost(W, b) = \frac{1}{m} \sum_{i=1}^m (H(x^{(i)}) - y^{(i)})^2$$

## Simplified hypothesis in TF

$$H(x) = Wx$$

```
y_model = tf.mul(X, w)
```

$$cost(W) = \frac{1}{m} \sum_{i=1}^m (Wx^{(i)} - y^{(i)})^2$$

```
cost = tf.square(Y - y_model)
```

## Goal: Minimize cost

$$\underset{W,b}{\text{minimize}} \ cost(W, b)$$

## What $cost(W)$ looks like?

$$cost(W) = \frac{1}{m} \sum_{i=1}^m (Wx^{(i)} - y^{(i)})^2$$

x	y
1	1
2	2
3	3

- $W=1, cost(W)=?$

## What $cost(W)$ looks like?

$$cost(W) = \frac{1}{m} \sum_{i=1}^m (Wx^{(i)} - y^{(i)})^2$$

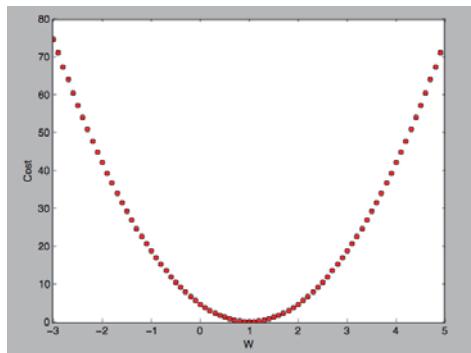
x	y
1	1
2	2
3	3

- $W=1, cost(W)=0$   
 $\frac{1}{3}((1*1-1)^2 + (1*2-2)^2 + (1*3-3)^2)$
- $W=0, cost(W)=4.67$   
 $\frac{1}{3}((0*1-1)^2 + (0*2-2)^2 + (0*3-3)^2)$
- $W=2, cost(W)=?$

## What $cost(W)$ looks like?

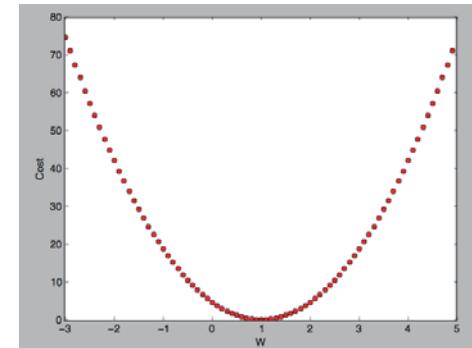
$$cost(W) = \frac{1}{m} \sum_{i=1}^m (Wx^{(i)} - y^{(i)})^2$$

- $W=1, cost(W)=0$
- $W=0, cost(W)=4.67$
- $W=2, cost(W)=4.67$



## How to minimize cost?

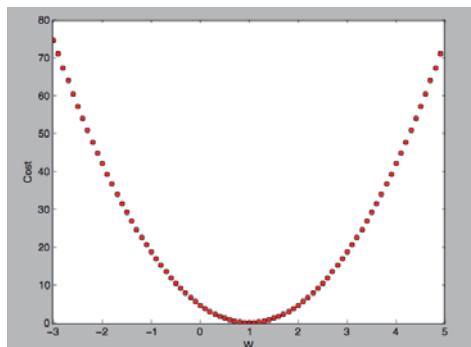
$$cost(W) = \frac{1}{m} \sum_{i=1}^m (Wx^{(i)} - y^{(i)})^2$$



## Formal definition

$$cost(W) = \frac{1}{2m} \sum_{i=1}^m (Wx^{(i)} - y^{(i)})^2$$

$$W := W - \alpha \frac{\partial}{\partial W} cost(W)$$



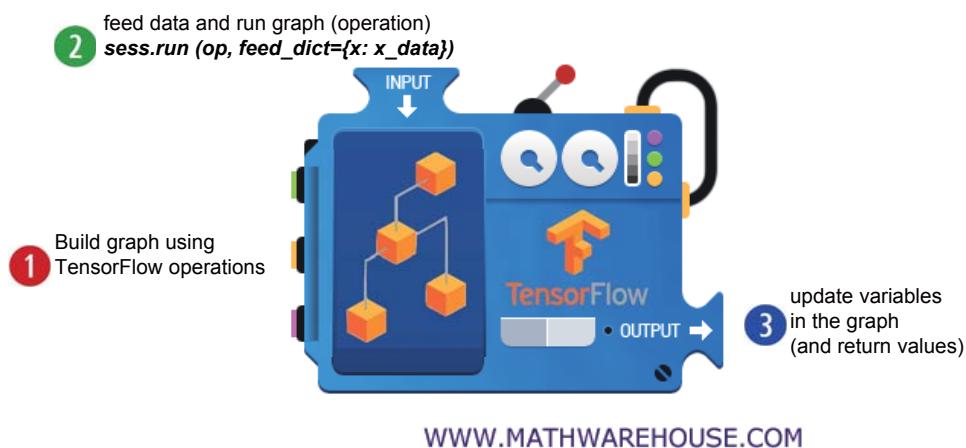
## Gradient descent algorithm in TF

$$W := W - \alpha \frac{\partial}{\partial W} cost(W)$$

$$W := W - \alpha \frac{1}{m} \sum_{i=1}^m (Wx^{(i)} - y^{(i)})x^{(i)}$$

```
cost = tf.square(Y - y_model) # use square error for cost function
train_op = tf.train.GradientDescentOptimizer(0.01).minimize(cost)
```

## TensorFlow Mechanics



### ① Build graph using TF operations

$$H(x) = Wx + b$$

```
# X and Y data
x_train = [1, 2, 3]
y_train = [1, 2, 3]

W = tf.Variable(tf.random_normal([1]), name='weight')
b = tf.Variable(tf.random_normal([1]), name='bias')
# Our hypothesis XW+b
hypothesis = x_train * W + b

cost(W, b) = 
$$\frac{1}{m} \sum_{i=1}^m (H(x^{(i)}) - y^{(i)})^2$$


# cost/loss function
cost = tf.reduce_mean(tf.square(hypothesis - y_train))
```

### ① Build graph using TF operations

$$cost(W, b) = \frac{1}{m} \sum_{i=1}^m (H(x^{(i)}) - y^{(i)})^2$$

```
t = [1., 2., 3., 4.]
tf.reduce_mean(t) => 2.5
```

```
# cost/loss function
cost = tf.reduce_mean(tf.square(hypothesis - y_train))
```

### GradientDescent

```
# Minimize
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.01)
train = optimizer.minimize(cost)
```

[https://www.tensorflow.org/api\\_docs/python/tf/reduce\\_mean](https://www.tensorflow.org/api_docs/python/tf/reduce_mean)

### ② Run/update graph and get results

```
# Launch the graph in a session.
sess = tf.Session()
# Initializes global variables in the graph.
sess.run(tf.global_variables_initializer())
```

```
# Fit the line
for step in range(2001):
    sess.run(train)
    if step % 20 == 0:
        print(step, sess.run(cost), sess.run(W), sess.run(b))
```

```

import tensorflow as tf

# X and Y data
x_train = [1, 2, 3]
y_train = [1, 2, 3]

W = tf.Variable(tf.random_normal([1]), name='weight')
b = tf.Variable(tf.random_normal([1]), name='bias')

# Our hypothesis XW+b
hypothesis = x_train * W + b

# cost/loss function
cost = tf.reduce_mean(tf.square(hypothesis - y_train))

# Minimize
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.01)
train = optimizer.minimize(cost)

# Launch the graph in a session.
sess = tf.Session()
# Initializes global variables in the graph.
sess.run(tf.global_variables_initializer())

# Fit the line
for step in range(2001):
    sess.run(train)
    if step % 20 == 0:
        print(step, sess.run(cost), sess.run(W), sess.run(b))

```

...
 0 2.82329 [ 2.12867713] [-0.85235667]  
 20 0.190351 [ 1.53392804] [-1.05059612]  
 40 0.151357 [ 1.45725465] [-1.02391243]  
 ...  
 1920 1.77484e-05 [ 1.00489295] [-0.01112291]  
 1940 1.61197e-05 [ 1.00466311] [-0.01060018]  
 1960 1.46397e-05 [ 1.0044444] [-0.01010205]  
 1980 1.32962e-05 [ 1.00423515] [-0.00962736]  
 2000 1.20761e-05 [ 1.00403607] [-0.00917497]  
 ...

## Full code (less than 20 lines)


大数讲堂  
umi.org.cn

## Placeholders

```

In [7]: a = tf.placeholder(tf.float32)
b = tf.placeholder(tf.float32)
adder_node = a + b # + provides a shortcut for tf.add(a, b)

print(sess.run(adder_node, feed_dict={a: 3, b: 4.5}))
print(sess.run(adder_node, feed_dict={a: [1,3], b: [2, 4]}))

7.5
[ 3.  7.]

```

[lab-01-basics.ipynb](#)


大数讲堂  
umi.org.cn

## Placeholders

```

# X and Y data
x_train = [1, 2, 3]
y_train = [1, 2, 3]

# Now we can use X and Y in place of x_data and y_data
# # placeholders for a tensor that will be always fed using feed_dict
# See http://stackoverflow.com/questions/36693740/
X = tf.placeholder(tf.float32)
Y = tf.placeholder(tf.float32)
...
# Fit the Line
for step in range(2001):
    cost_val, W_val, b_val, _ = \
        sess.run([cost, W, b, train],
                feed_dict={X: [1, 2, 3], Y: [1, 2, 3]})
    if step % 20 == 0:
        print(step, cost_val, W_val, b_val)

```


大数讲堂  
umi.org.cn

## Full code with placeholders

```

import tensorflow as tf
W = tf.Variable(tf.random_normal([1]), name='weight')
b = tf.Variable(tf.random_normal([1]), name='bias')

X = tf.placeholder(tf.float32, shape=[None])
Y = tf.placeholder(tf.float32, shape=[None])

# Our hypothesis XW+b
hypothesis = X * W + b
# cost/loss function
cost = tf.reduce_mean(tf.square(hypothesis - Y))
# Minimize
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.01)
train = optimizer.minimize(cost)

# Launch the graph in a session.
sess = tf.Session()
# Initializes global variables in the graph.
sess.run(tf.global_variables_initializer())

# Fit the line
for step in range(2001):
    cost_val, W_val, b_val, _ = sess.run([cost, W, b, train],
                                         feed_dict={X: [1, 2, 3], Y: [1, 2, 3]})
    if step % 20 == 0:
        print(step, cost_val, W_val, b_val)

```

...
 1980 1.32962e-05 [ 1.00423515] [-0.00962736]  
 2000 1.20761e-05 [ 1.00403607] [-0.00917497]  
 ...  
 # Testing our model
 print(sess.run(hypothesis, feed\_dict={X: [5]}))
 print(sess.run(hypothesis, feed\_dict={X: [2.5]}))
 print(sess.run(hypothesis,
 feed\_dict={X: [1.5, 3.5]}))

 [ 5.0110054]
 [ 2.50091505]
 [ 1.49687922 3.50495124]

```

import tensorflow as tf
W = tf.Variable(tf.random_normal([1]), name='weight')
b = tf.Variable(tf.random_normal([1]), name='bias')
X = tf.placeholder(tf.float32, shape=[None])
Y = tf.placeholder(tf.float32, shape=[None])

# Our hypothesis XW+b
hypothesis = X * W + b
# cost/loss function
cost = tf.reduce_mean(tf.square(hypothesis - Y))
# Minimize
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.01)
train = optimizer.minimize(cost)

# Launch the graph in a session.
sess = tf.Session()
# Initializes global variables in the graph.
sess.run(tf.global_variables_initializer())

# Fit the Line with new training data
for step in range(2001):
    cost_val, W_val, b_val, _ = sess.run([cost, W, b, train],
                                         feed_dict={X: [1, 2, 3, 4, 5],
                                                    Y: [2.1, 3.1, 4.1, 5.1, 6.1]})
    if step % 20 == 0:
        print(step, cost_val, W_val, b_val)

```

## Full code with placeholders

```

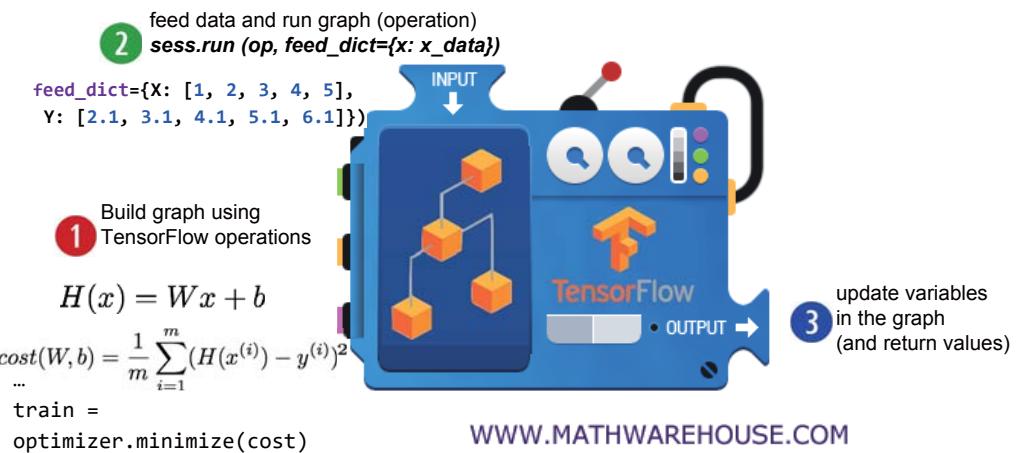
1960 3.32396e-07 [ 1.00037301] [ 1.09886529]
1980 2.90429e-07 [ 1.00034881] [ 1.09874094]
2000 2.5373e-07 [ 1.00032604] [ 1.09882331]

# Testing our model
print(sess.run(hypothesis, feed_dict={X: [5]}))
print(sess.run(hypothesis, feed_dict={X: [2.5]}))
print(sess.run(hypothesis,
               feed_dict={X: [1.5, 3.5]}))

[ 6.10045338]
[ 3.59963846]
[ 2.59931231  4.59996414]

```

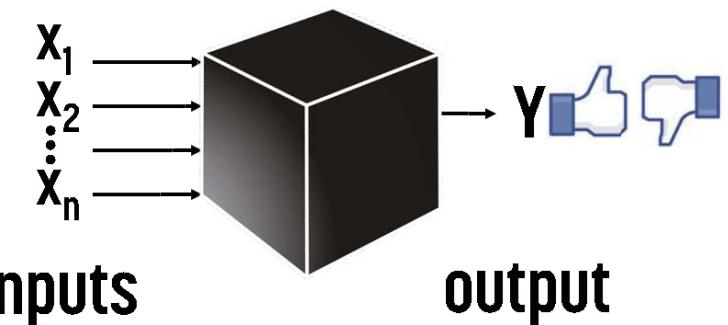
## TensorFlow Mechanics



## Machine Learning Basics

- Linear Regression
- Logistic Regression (Binary classification)
- Softmax Classification
- Neural Networks

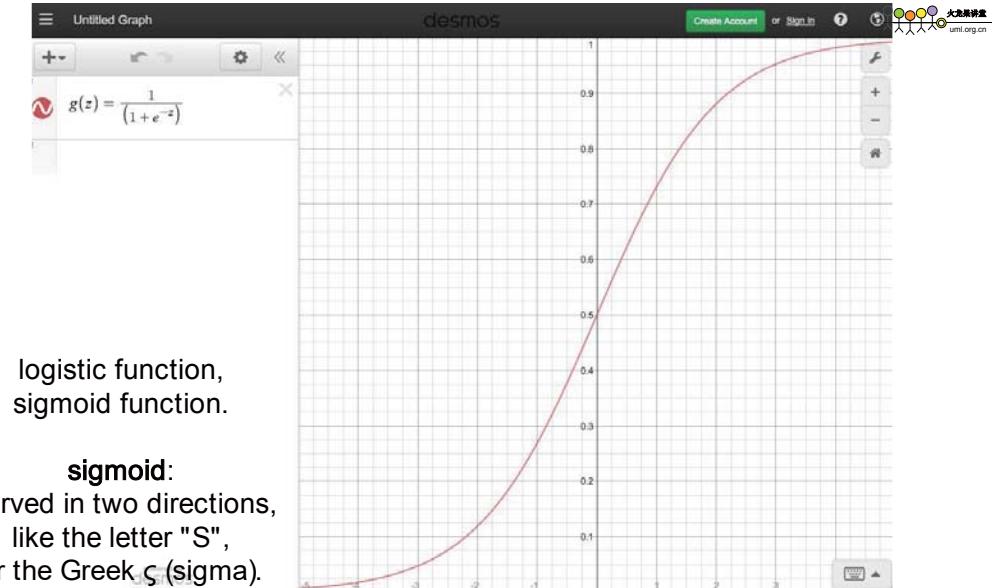
## Classification



## Linear Regression Hypothesis

$$H(x) = Wx$$

$$cost(W) = \frac{1}{m} \sum_{i=1}^m (Wx^{(i)} - y^{(i)})^2$$



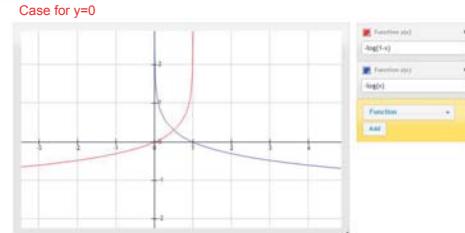
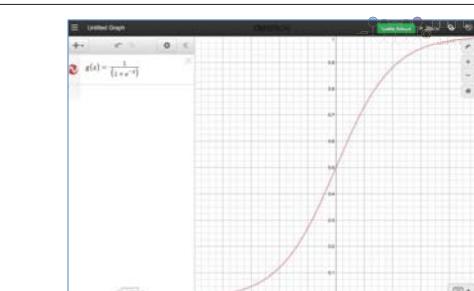
## Logistic Regression

$$H(X) = \text{sigmoid}(XW) = \frac{1}{1 + e^{-XW}}$$

$$cost(W) = -\frac{1}{m} \sum [y \log(H(x)) + (1-y)(\log(1-H(x))]$$

Case for y=1

$$W := W - \alpha \frac{\partial}{\partial W} cost(W)$$



## Logistic Classification in TF

$$H(X) = \text{sigmoid}(XW) = \frac{1}{1 + e^{-XW}}$$

hypothesis =  
`tf.sigmoid(tf.matmul(X, W) + b)`

$$cost(W) = -\frac{1}{m} \sum [y \log(H(x)) + (1-y)(\log(1-H(x))]$$

# cost/loss function

```
cost = -tf.reduce_mean(Y * tf.log(hypothesis) + (1 - Y) * tf.log(1 - hypothesis))
```

## Training Data

```
x_data = [[1, 2], [2, 3], [3, 1], [4, 3], [5, 3], [6, 2]]
y_data = [[0], [0], [0], [1], [1], [1]]
```

# placeholders for a tensor that will be always fed.

```
X = tf.placeholder(tf.float32, shape=[None, 2])
Y = tf.placeholder(tf.float32, shape=[None, 1])
```

[lab-05-1-logistic\\_regression.py](#)

```
X = tf.placeholder(tf.float32, shape=[None, 2])
Y = tf.placeholder(tf.float32, shape=[None, 1])
W = tf.Variable(tf.random_normal([2, 1]), name='weight')
b = tf.Variable(tf.random_normal([1]), name='bias')
```

# Hypothesis using sigmoid:  $tf.div(1., 1. + tf.exp(tf.matmul(X, W) + b))$

```
hypothesis = tf.sigmoid(tf.matmul(X, W) + b)
```

$$H(X) = \frac{1}{1 + e^{-WX}}$$

# cost/loss function

```
cost = -tf.reduce_mean(Y * tf.log(hypothesis) + (1 - Y) *
tf.log(1 - hypothesis))
```

train = tf.train.GradientDescentOptimizer(learning\_rate=0.01).minimize(cost)

$$W := W - \alpha \frac{\partial}{\partial W} cost(W)$$

# Accuracy computation

# True if hypothesis>0.5 else False

```
predicted = tf.cast(hypothesis > 0.5, dtype=tf.float32)
```

```
accuracy = tf.reduce_mean(tf.cast(tf.equal(predicted, Y), dtype=tf.float32))
```

[lab-05-1-logistic\\_regression.py](#)

## Train the model

```
# Launch graph
with tf.Session() as sess:
    # Initialize TensorFlow variables
    sess.run(tf.global_variables_initializer())

for step in range(10001):
    cost_val, _ = sess.run([cost, train], feed_dict={X: x_data, Y: y_data})
    if step % 200 == 0:
        print(step, cost_val)

# Accuracy report
h, c, a = sess.run([hypothesis, predicted, accuracy],
                  feed_dict={X: x_data, Y: y_data})
print("\nHypothesis: ", h, "\nCorrect (Y): ", c, "\nAccuracy: ", a)
```

[lab-05-1-logistic\\_regression.py](#)

```
x_data = [[1, 2], [2, 3], [3, 1], [4, 3], [5, 3], [6, 2]]
y_data = [[0], [0], [0], [1], [1], [1]]
```

# placeholders for a tensor that will be always fed.

```
X = tf.placeholder(tf.float32, shape=[None, 2])
Y = tf.placeholder(tf.float32, shape=[None, 1])
```

```
W = tf.Variable(tf.random_normal([2, 1]), name='weight')
```

```
b = tf.Variable(tf.random_normal([1]), name='bias')
```

# Hypothesis using sigmoid:  $tf.div(1., 1. + tf.exp(tf.matmul(X, W)))$

```
hypothesis = tf.sigmoid(tf.matmul(X, W) + b)
```

# cost/Loss function

```
cost = -tf.reduce_mean(Y * tf.log(hypothesis) + (1 - Y) * tf.log(1 - hypothesis))
```

```
train = tf.train.GradientDescentOptimizer(learning_rate=0.01).minimize(cost)
```

# Accuracy computation

# True if hypothesis>0.5 else False

```
predicted = tf.cast(hypothesis > 0.5, dtype=tf.float32)
```

```
accuracy = tf.reduce_mean(tf.cast(tf.equal(predicted, Y), dtype=tf.float32))
```

# Launch graph

```
with tf.Session() as sess:
```

# Initialize TensorFlow variables

```
sess.run(tf.global_variables_initializer())
```

for step in range(10001):

```
    cost_val, _ = sess.run([cost, train], feed_dict={X: x_data, Y: y_data})
```

if step % 200 == 0:

```
        print(step, cost_val)
```

# Accuracy report

```
h, c, a = sess.run([hypothesis, predicted, accuracy],
                  feed_dict={X: x_data, Y: y_data})
```

```
print("\nHypothesis: ", h, "\nCorrect (Y): ", c, "\nAccuracy: ", a)
```

```
# step, cost
0 1.73078
200 0.571512
400 0.507414
...
9600 0.154132
9800 0.151778
10000 0.149496
```

Hypothesis:

```
[[ 0.03074029]
 [ 0.15884677]
 [ 0.30486736]
 [ 0.78138196]
 [ 0.93957496]
 [ 0.98016882]]
```

Correct (Y):

```
[[ 0.]
 [ 0.]
 [ 0.]
 [ 1.]
 [ 1.]
 [ 1.]]
```

Accuracy: 1.0

# Classifying diabetes

## 机器学习实战项目



-0.411765	0.165829	0.213115	0	0	-0.23696	-0.894962	-0.7	1
-0.647059	-0.21608	-0.180328	-0.353535	-0.791962	-0.0760059	-0.854825	-0.833333	0
0.176471	0.155779	0	0	0	0.052161	-0.952178	-0.733333	1
-0.764706	0.979899	0.147541	-0.0909091	0.283688	-0.0909091	-0.931682	0.0666667	0
-0.0588235	0.256281	0.57377	0	0	0	-0.868488	0.1	0
-0.529412	0.105528	0.508197	0	0	0.120715	-0.903501	-0.7	1
0.176471	0.688442	0.213115	0	0	0.132638	-0.608027	-0.566667	0
0.176471	0.396985	0.311475	0	0	-0.19225	0.163962	0.2	1

```
xy = np.loadtxt('data-03-diabetes.csv', delimiter=',', dtype=np.float32)
x_data = xy[:, :-1]
y_data = xy[:, [-1]]
```

[lab-05-2-logistic\\_regression\\_diabetes.py](#)

```
xy = np.loadtxt('data-03-diabetes.csv', delimiter=',', dtype=np.float32)
x_data = xy[:, :-1]
y_data = xy[:, [-1]]

# placeholders for a tensor that will be always fed.
X = tf.placeholder(tf.float32, shape=[None, 8])
Y = tf.placeholder(tf.float32, shape=[None, 1])

W = tf.Variable(tf.random_normal([8, 1]), name='weight')
b = tf.Variable(tf.random_normal([1]), name='bias')

# Hypothesis using sigmoid: tf.div(1., 1. + tf.exp(tf.matmul(X, W)))
hypothesis = tf.sigmoid(tf.matmul(X, W) + b)
# cost/loss function
cost = -tf.reduce_mean(Y * tf.log(hypothesis) + (1 - Y) * tf.log(1 - hypothesis))
train = tf.train.GradientDescentOptimizer(learning_rate=0.01).minimize(cost)

# Accuracy computation
# True if hypothesis>0.5 else False
predicted = tf.cast(hypothesis > 0.5, dtype=tf.float32)
accuracy = tf.reduce_mean(tf.cast(tf.equal(predicted, Y), dtype=tf.float32))

# Launch graph
with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())

    feed = {X: x_data, Y: y_data}
    for step in range(1000):
        sess.run(train, feed_dict=feed)
        if step % 200 == 0:
            print(step, sess.run(cost, feed_dict=feed))

    # Accuracy report
h, c, a = sess.run([hypothesis, predicted, accuracy], feed_dict=feed)
print("\nHypothesis: ", h, "\nCorrect (Y): ", c, "\nAccuracy: ", a)
```

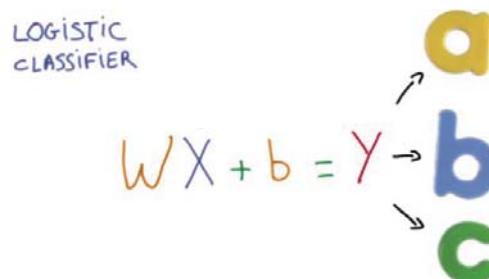
[lab-05-2-logistic\\_regression\\_diabetes.py](#)

火龙果讲堂	umi.org.cn
0 0.82794	[ 0.82794 ]
200 0.755181	[ 0.755181 ]
400 0.726355	[ 0.726355 ]
600 0.705179	[ 0.705179 ]
800 0.686631	[ 0.686631 ]
...	[ ... ]
9600 0.492056	[ 0.492056 ]
9800 0.491396	[ 0.491396 ]
10000 0.490767	[ 0.490767 ]
	[ 0.7461012 ]
	[ 0.79919308 ]
	[ 0.72995949 ]
	[ 0.88297188 ]
	[ 1.]
	[ 1.]
	[ 1.]
	Accuracy:
	0.762846

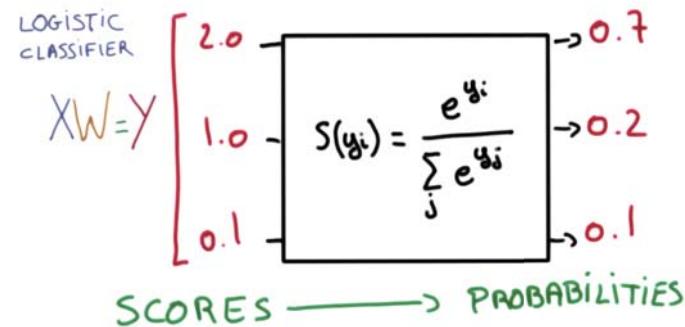
## Machine Learning Basics

- Linear Regression
- Logistic Regression (Binary classification)
- Softmax Classification
- Neural Networks

## Multiple labels (a, b, c)

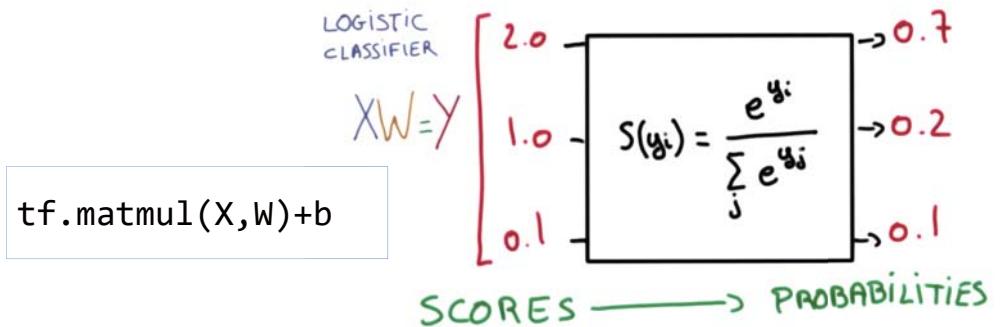


## Softmax function



<https://www.udacity.com/course/viewer#!/c-ud730/l-6370362152/m-6379811817>

```
hypothesis = tf.nn.softmax(tf.matmul(X,w)+b)
```



<https://www.udacity.com/course/viewer#!/c-ud730/l-6370362152/m-6379811817>

## Cost function: cross entropy

LOSS

$$\mathcal{L} = \frac{1}{N} \sum_i D(s(wx_i + b), l_i)$$

TRAINING SET

STEP

$$-\alpha \Delta \mathcal{L}(w_i, w_2)$$

DERIVATIVE

```
# Cross entropy cost/loss
cost = tf.reduce_mean(-tf.reduce_sum(Y * tf.log(hypothesis), axis=1))
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.1).minimize(cost)
```

<https://www.udacity.com/course/viewer#!/c-ud730/l-6370362152/m-6379811817>

## Cost function: cross entropy

```
hypothesis = tf.nn.softmax(tf.matmul(X,w)+b)

# Cross entropy cost/loss
cost = tf.reduce_mean(
    -tf.reduce_sum(Y * tf.log(hypothesis), axis=1))

optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.1).minimize(cost)
```

<https://www.udacity.com/course/viewer#!/c-ud730/l-6370362152/m-6379811817>

```
x_data = [[1, 2, 1, 1], [2, 1, 3, 2], [3, 1, 3, 4], [4, 1, 5, 5], [1, 7, 5, 5],
           [1, 2, 5, 6], [1, 6, 6, 6], [1, 7, 7, 7]] x x x x x x x x
y_data = [[0, 0, 1], [0, 0, 1], [0, 0, 1], [0, 1, 0], [0, 1, 0], [0, 1, 0], [1, 0, 0], [1, 0, 0]]
X = tf.placeholder("float", [None, 4])
Y = tf.placeholder("float", [None, 3])
nb_classes = 3

W = tf.Variable(tf.random_normal([4, nb_classes]), name='weight')
b = tf.Variable(tf.random_normal([nb_classes]), name='bias')

# tf.nn.softmax computes softmax activations
# softmax = exp(Logits) / reduce_sum(exp(Logits), dim)
hypothesis = tf.nn.softmax(tf.matmul(X, W) + b)

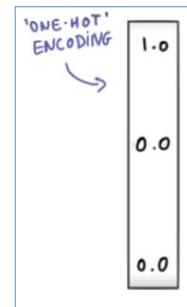
# Cross entropy cost/loss
cost = tf.reduce_mean(-tf.reduce_sum(Y * tf.log(hypothesis), axis=1))
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.1).minimize(cost)

# Launch graph
with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())

    for step in range(2001):
        sess.run(optimizer, feed_dict={X: x_data, Y: y_data})
        if step % 200 == 0:
            print(step, sess.run(cost, feed_dict={X: x_data, Y: y_data}))
lab-06-1-softmax_classifier.py
```

'ONE-HOT' ENCODING

1.0  
0.0  
0.0

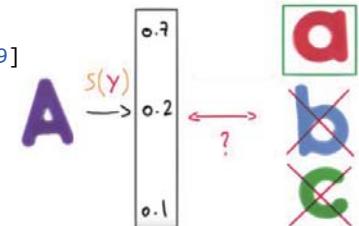


## Test & one-hot encoding

```
hypothesis = tf.nn.softmax(tf.matmul(X,W)+b)
```

```
# Testing & One-hot encoding  
a = sess.run(hypothesis, feed_dict={X: [[1, 11, 7, 9]]})  
print(a, sess.run(tf.argmax(a, 1)))
```

```
[[ 1.38904958e-03  9.98601854e-01  9.06129117e-06]] [1]
```



lab-06-1-softmax\_classifier.py

## Test & one-hot encoding

```
hypothesis = tf.nn.softmax(tf.matmul(x,w)+b)
```

```
all = sess.run(hypothesis, feed_dict={X: [[1, 11, 7, 9],  
                                         [1, 3, 4, 3],  
                                         [1, 1, 0, 1]]})  
print(all, sess.run(tf.argmax(all, 1)))
```

```
[1] 1.38904958e-03 9.98601854e-01 9.06129117e-06]
[2] 9.31192040e-01 6.29020557e-02 5.90589503e-03]
[3] 1.27327668e-08 3.34112905e-04 9.99665856e-01]]
```

[1 0 2]

[lab-06-1-softmax\\_classifier.py](#)

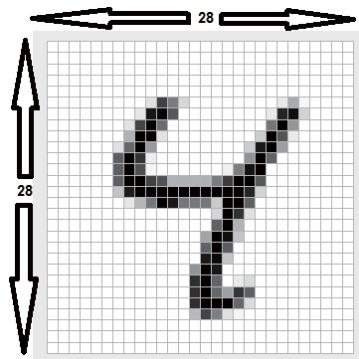
MNIST Dataset 机器学习实战项目

000000000000000000000000  
111111111111111111111111  
222222222222222222222222  
333333333333333333333333  
494444444444444444444444  
555555555555555555555555  
666666666666666666666666  
777777777777777777777777  
888888888888888888888888  
999999999999999999999999

[train-images-idx3-ubyte.gz](#): training set images (9912422 bytes)  
[train-labels-idx1-ubyte.gz](#): training set labels (28881 bytes)  
[t10k-images-idx3-ubyte.gz](#): test set images (1648877 bytes)  
[t10k-labels-idx1-ubyte.gz](#): test set labels (4542 bytes)

<http://yann.lecun.com/exdb/mnist/>

28x28x1 image



```
# MNIST data image of shape 28 * 28 = 784
X = tf.placeholder(tf.float32, [None, 784])
# 0 - 9 digits recognition = 10 classes
Y = tf.placeholder(tf.float32, [None, nb_classes])
```

<http://derindelimavi.blogspot.hk/2015/04/mnist-el-yazs-rakam-veri-seti.html>

# MNIST Dataset



```
from tensorflow.examples.tutorials.mnist import input_data
# Check out https://www.tensorflow.org/get_started/mnist/beginners for
# more information about the mnist dataset
mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)
...
batch_xs, batch_ys = mnist.train.next_batch(100)
...
print("Accuracy: ", accuracy.eval(session=sess,
    feed_dict={x: mnist.test.images, y: mnist.test.labels}))
```

lab-07-4-mnist introduction.py

## Reading data and set variables

```
from tensorflow.examples.tutorials.mnist import input_data
# Check out https://www.tensorflow.org/get_started/mnist/beginners for
# more information about the mnist dataset
mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)

nb_classes = 10

# MNIST data image of shape 28 * 28 = 784
X = tf.placeholder(tf.float32, [None, 784])
# 0 - 9 digits recognition = 10 classes
Y = tf.placeholder(tf.float32, [None, nb_classes])

W = tf.Variable(tf.random_normal([784, nb_classes]))
b = tf.Variable(tf.random_normal([nb_classes]))
```

lab-07-4-mnist introduction.py

## Softmax!

```

# Hypothesis (using softmax)
hypothesis = tf.nn.softmax(tf.matmul(X, W) + b)

cost = tf.reduce_mean(-tf.reduce_sum(Y * tf.log(hypothesis), axis=1))
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.1).minimize(cost)

# Test model
is_correct = tf.equal(tf.argmax(hypothesis, 1), tf.argmax(Y, 1))
# Calculate accuracy
accuracy = tf.reduce_mean(tf.cast(is_correct, tf.float32))

```

lab-07-4-mnist introduction.py



## Training epoch/batch

```
# parameters
training_epochs = 15
batch_size = 100

with tf.Session() as sess:
    # Initialize TensorFlow variables
    sess.run(tf.global_variables_initializer())
    # Training cycle
    for epoch in range(training_epochs):
        avg_cost = 0
        total_batch = int(mnist.train.num_examples / batch_size)

        for i in range(total_batch):
            batch_xs, batch_ys = mnist.train.next_batch(batch_size)
            c, _ = sess.run([cost, optimizer], feed_dict={X: batch_xs, Y: batch_ys})
            avg_cost += c / total_batch

        print('Epoch:', '%04d' % (epoch + 1), 'cost =', '{:.9f}'.format(avg_cost))
lab-07-4-mnist_introduction.py
```

## Training epoch/batch

In the neural network terminology:

one **epoch** = one forward pass and one backward pass of *all* the training examples

**batch size** = the number of training examples in one forward/backward pass. The higher the batch size, the more memory space you'll need.

number of **iterations** = number of passes, each pass using [batch size] number of examples.

To be clear, one pass = one forward pass + one backward pass (we do not count the forward pass and backward pass as two different passes).

Example: if you have *1000 training examples*, and your *batch size is 500*, then it will take 2 iterations to complete 1 epoch.

<http://stackoverflow.com/questions/4752626/epoch-vs-iteration-when-training-neural-networks>



## Training epoch/batch

```
# parameters
training_epochs = 15
batch_size = 100

with tf.Session() as sess:
    # Initialize TensorFlow variables
    sess.run(tf.global_variables_initializer())
    # Training cycle
    for epoch in range(training_epochs):
        avg_cost = 0
        total_batch = int(mnist.train.num_examples / batch_size)

        for i in range(total_batch):
            batch_xs, batch_ys = mnist.train.next_batch(batch_size)
            c, _ = sess.run([cost, optimizer], feed_dict={X: batch_xs, Y: batch_ys})
            avg_cost += c / total_batch

        print('Epoch:', '%04d' % (epoch + 1), 'cost =', '{:.9f}'.format(avg_cost))
lab-07-4-mnist_introduction.py
```



## Report results on test dataset

```
# Test the model using test sets
print("Accuracy: ", accuracy.eval(session=sess,
    feed_dict={X: mnist.test.images, Y: mnist.test.labels}))
```

[lab-07-4-mnist\\_introduction.py](lab-07-4-mnist_introduction.py)

```

hypothesis = tf.nn.softmax(tf.matmul(X, W) + b)
cost = tf.reduce_mean(-tf.reduce_sum(Y * tf.log(hypothesis), axis=1))
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.1).minimize(cost)

is_correct = tf.equal(tf.argmax(hypothesis, 1), tf.argmax(Y, 1))
accuracy = tf.reduce_mean(tf.cast(is_correct, tf.float32))

# parameters
training_epochs = 15
batch_size = 100

with tf.Session() as sess:
    # Initialize TensorFlow variables
    sess.run(tf.global_variables_initializer())
    # Training cycle
    for epoch in range(training_epochs):
        avg_cost = 0
        total_batch = int(mnist.train.num_examples / batch_size)

        for i in range(total_batch):
            batch_xs, batch_ys = mnist.train.next_batch(batch_size)
            c, _ = sess.run([cost, optimizer],
                           feed_dict={X: batch_xs, Y: batch_ys})
            avg_cost += c / total_batch

        print('Epoch:', '%04d' % (epoch + 1),
              'cost =', '{:.9f}'.format(avg_cost))

```

[lab-07-4-mnist\\_introduction.py](#)



大漠讲堂



大漠讲堂  
umi.org.cn

## Sample image show and prediction

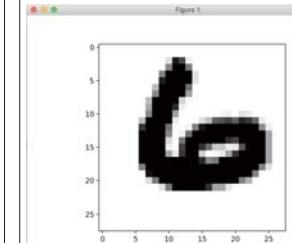


Figure 1

```

import matplotlib.pyplot as plt
import random

# Get one and predict
r = random.randint(0, mnist.test.num_examples - 1)
print("Label:", sess.run(tf.argmax(mnist.test.labels[r:r+1], 1)))
print("Prediction:", sess.run(tf.argmax(hypothesis, 1),
                             feed_dict={X: mnist.test.images[r:r + 1]}))

plt.imshow(mnist.test.images[r:r + 1].
           reshape(28, 28), cmap='Greys', interpolation='nearest')
plt.show()

```

[lab-07-4-mnist\\_introduction.py](#)

## Machine Learning Basics

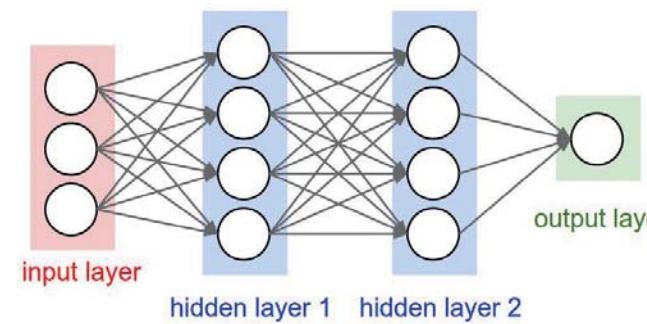


大漠讲堂



- Linear Regression
- Logistic Regression (Binary classification)
- Softmax Classification
- Neural Networks

## Neural Net 深度学习进阶 神经网络



```

# weights & bias for nn layers
W = tf.Variable(tf.random_normal([784, 10]))
b = tf.Variable(tf.random_normal([10]))
hypothesis = tf.matmul(X, W) + b
# define cost/loss & optimizer
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=hypothesis, labels=Y))
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost)
# initialize
sess = tf.Session()
sess.run(tf.global_variables_initializer())
# train my model
for epoch in range(training_epochs):
    avg_cost = 0
    total_batch = int(mnist.train.num_examples / batch_size)

    for i in range(total_batch):
        batch_xs, batch_ys = mnist.train.next_batch(batch_size)
        feed_dict = {X: batch_xs, Y: batch_ys}
        c, _ = sess.run([cost, optimizer], feed_dict=feed_dict)
        avg_cost += c / total_batch

    print('Epoch:', '%04d' % (epoch + 1), 'cost =', '{:.9f}'.format(avg_cost))

print('Learning Finished!')

# Test model and check accuracy
correct_prediction = tf.equal(tf.argmax(hypothesis, 1), tf.argmax(Y, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
print('Accuracy:', sess.run(accuracy, feed_dict={X: mnist.test.images, Y: mnist.test.labels}))

```

## Softmax classifier for MNIST

```

cost = tf.reduce_mean(-tf.reduce_sum(Y *
tf.log(hypothesis), axis=1))

```

```

Epoch: 0001 cost = 5.888845987
Epoch: 0002 cost = 1.860620173
Epoch: 0003 cost = 1.159035648
Epoch: 0004 cost = 0.892340870
Epoch: 0005 cost = 0.751155428
Epoch: 0006 cost = 0.662484806
Epoch: 0007 cost = 0.601544010
Epoch: 0008 cost = 0.556526115
Epoch: 0009 cost = 0.521186961
Epoch: 0010 cost = 0.493068354
Epoch: 0011 cost = 0.469686249
Epoch: 0012 cost = 0.449967254
Epoch: 0013 cost = 0.433519321
Epoch: 0014 cost = 0.419000337
Epoch: 0015 cost = 0.406490815

```

Learning Finished!

Accuracy: 0.9035

[lab-10-1-mnist\\_softmax.py](#)

## NN for MNIST

```

Epoch: 0001 cost = 141.207671860
Epoch: 0002 cost = 38.788445864
Epoch: 0003 cost = 23.977515479
Epoch: 0004 cost = 16.315132428
Epoch: 0005 cost = 11.702554882
Epoch: 0006 cost = 8.573139748
Epoch: 0007 cost = 6.370995680
Epoch: 0008 cost = 4.537178684
Epoch: 0009 cost = 3.216900532
Epoch: 0010 cost = 2.329708954
Epoch: 0011 cost = 1.715552875
Epoch: 0012 cost = 1.189857912
Epoch: 0013 cost = 0.820965160
Epoch: 0014 cost = 0.624131458
Epoch: 0015 cost = 0.454633765

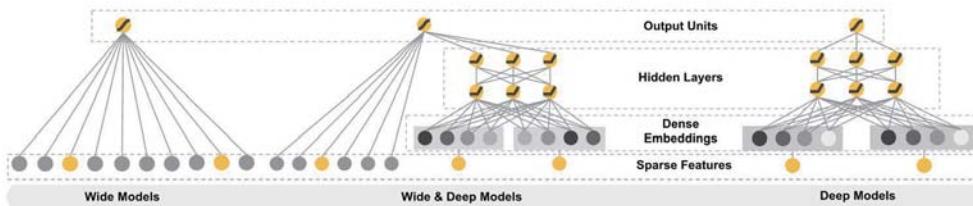
```

Learning Finished!

Accuracy: 0.9455

[lab-10-2-mnist\\_nn.py](#)

## Wide & Deep



[https://www.tensorflow.org/versions/r0.11/tutorials/wide\\_and\\_deep/index.html](https://www.tensorflow.org/versions/r0.11/tutorials/wide_and_deep/index.html)



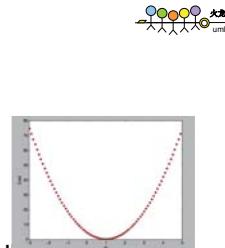
# Homework

- Tensor manipulation
  - [https://docs.google.com/presentation/d/1gQ7Xxrhykr5Kk5pG15yvX3yOln\\_hk2-H6jrQeXqKmU/](https://docs.google.com/presentation/d/1gQ7Xxrhykr5Kk5pG15yvX3yOln_hk2-H6jrQeXqKmU/)
- Titanic with NN 改进泰坦尼克获救预测
  - [k0-01-titanic](#)



- CNN
- RNN
- Example code: [DeepLearningZeroToAll](#)

## How to minimize cost?



- Start with initial guesses
  - Start at 0,0 (or any other value)
  - Keeping changing W and b a little bit to try and reduce  $\text{cost}(W, b)$
- Each time you change the parameters, you select the gradient which reduces  $\text{cost}(W, b)$  the most possible
- Repeat
- Do so until you converge to a local minimum
- Has an interesting property
  - Where you start can determine which minimum you end up

[http://www.holehouse.org/mlclass/01\\_02\\_Introduction\\_regression\\_analysis\\_and\\_gr.html](http://www.holehouse.org/mlclass/01_02_Introduction_regression_analysis_and_gr.html)

## Formal definition

$$\text{cost}(W) = \frac{1}{m} \sum_{i=1}^m (Wx^{(i)} - y^{(i)})^2$$



$$\text{cost}(W) = \frac{1}{2m} \sum_{i=1}^m (Wx^{(i)} - y^{(i)})^2$$

## Formal definition

$$W := W - \alpha \frac{\partial}{\partial W} \frac{1}{2m} \sum_{i=1}^m (Wx^{(i)} - y^{(i)})^2$$

$$W := W - \alpha \frac{1}{2m} \sum_{i=1}^m 2(Wx^{(i)} - y^{(i)})x^{(i)}$$

$$W := W - \alpha \frac{1}{m} \sum_{i=1}^m (Wx^{(i)} - y^{(i)})x^{(i)}$$

## Derivative Calculator

Calculate derivatives online  
— with steps and graphing!

**Hello there!**  
Was this calculator helpful to you? Then I would highly appreciate small donations via PayPal.  
[DONATE](#) 300.00 € to Get  
... or use this link for shopping on Amazon, without affecting your order.  
Thank you!

Calculate the Derivative of ...  
 $(xa-y)^2$

This will be calculated:  
 $\frac{d}{dx} [(xa-y)^2]$

Not what you mean? Use parentheses! Set differentiation variable and order in "Options".

About | Help | Examples | Options

The Derivative Calculator lets you calculate derivatives of functions with one or more variables.  
Our calculator allows you to check your solutions to calculus exercises. It helps you practice by showing you the full working (step by step differentiation).  
The Derivative Calculator supports computing first, second, ..., fifth derivatives as well as differentiating functions with many variables (partial derivatives).  
Interactive graphs/plots help visualize and better understand the functions.  
For more about how to use the Derivative Calculator, go to "Help" or take a look at the examples.  
And now: Happy differentiating!

## Gradient descent algorithm

$$W := W - \alpha \frac{1}{m} \sum_{i=1}^m (Wx^{(i)} - y^{(i)})x^{(i)}$$

## CNN

深度学习核心卷积神经网络



# Gifts from Google and Line

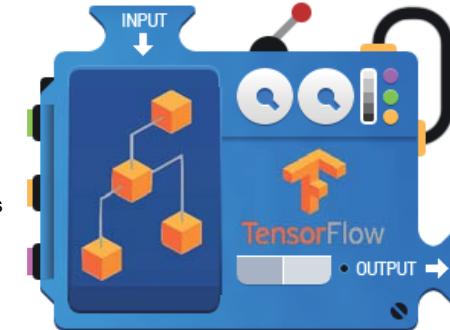


火龙果讲堂  
uml.org.cn

## TensorFlow Mechanics

- feed data and run graph (operation)  
`sess.run (op, feed_dict={x: x_data})`

- Build graph using TensorFlow operations



火龙果讲堂  
uml.org.cn

- update variables in the graph (and return values)

WWW.MATHWAREHOUSE.COM

# Machine Learning Basics



火龙果讲堂  
uml.org.cn

- Linear Regression
- Logistic Regression (Binary classification)
- Softmax Classification
- Neural Networks

## Linear Regression

- $H(x) = Wx + b$

```
W = tf.Variable(tf.random_normal([1]), name='weight')
b = tf.Variable(tf.random_normal([1]), name='bias')
# Our hypothesis XW+b
hypothesis = X * W + b
```

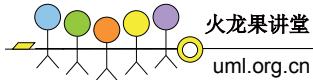
- $\text{cost}(W, b) = \frac{1}{m} \sum_{i=1}^m (H(x^{(i)}) - y^{(i)})^2$

```
# cost/loss function
cost = tf.reduce_mean(tf.square(hypothesis - y_train))
```

- GradientDescent  $W := W - \alpha \frac{\partial}{\partial W} \text{cost}(W)$

```
# Minimize
train = tf.train.GradientDescentOptimizer(learning_rate=0.01).minimize(cost)
```

# Logistic Regression



火龙果讲堂

uml.org.cn

① Model? (Hypothesis?)

# Logistic Regression



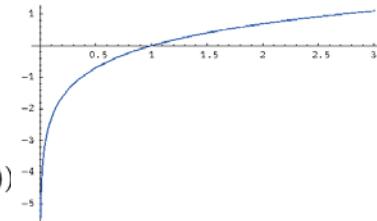
火龙果讲堂

uml.org.cn

① Model? (Hypothesis?)

$$H(X) = \text{sigmoid}(XW) = \frac{1}{1 + e^{-XW}}$$

```
hypothesis = tf.sigmoid(tf.matmul(X, W) + b)
```



② Cost?

② Cost?

$$\text{cost}(W) = -\frac{1}{m} \sum y \log(H(x)) + (1 - y) \log(1 - H(x))$$

# cost/Loss function

```
cost = -tf.reduce_mean(Y * tf.log(hypothesis) + (1 - Y) * tf.log(1 - hypothesis))
```

③ GradientDescent  $W := W - \alpha \frac{\partial}{\partial W} \text{cost}(W)$

```
train = tf.train.GradientDescentOptimizer(learning_rate=0.01).minimize(cost)
```

③ GradientDescent  $W := W - \alpha \frac{\partial}{\partial W} \text{cost}(W)$

```
train = tf.train.GradientDescentOptimizer(learning_rate=0.01).minimize(cost)
```

# Softmax Classifier



火龙果讲堂

uml.org.cn

① Model? (Hypothesis?)

```
logits = tf.matmul(X,W)+b
hypothesis = tf.nn.softmax(logits)
```

② Cost?

```
# Cross entropy cost/loss
cost = tf.reduce_mean (-tf.reduce_sum(Y * tf.log(hypothesis), axis=1))
```

③ GradientDescent  $W := W - \alpha \frac{\partial}{\partial W} \text{cost}(W)$

```
train = tf.train.GradientDescentOptimizer(learning_rate=0.01).minimize(cost)
```

# Softmax Classifier



火龙果讲堂

uml.org.cn

① Model? (Hypothesis?)

```
logits = tf.matmul(X,W)+b
hypothesis = tf.nn.softmax(logits)
```

② Cost?

```
# Cross entropy cost/loss
cost = tf.reduce_mean (-tf.reduce_sum(Y * tf.log(hypothesis), axis=1))
```

```
cost = tf.reduce_mean(
    tf.nn.softmax_cross_entropy_with_logits(logits=logits, labels=Y))
```

③ GradientDescent  $W := W - \alpha \frac{\partial}{\partial W} \text{cost}(W)$

```
train = tf.train.GradientDescentOptimizer(learning_rate=0.01).minimize(cost)
```

# (Deep) Neural Nets

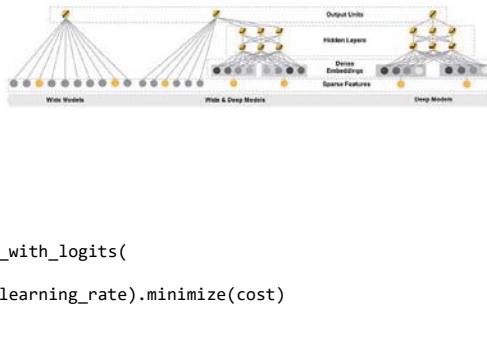
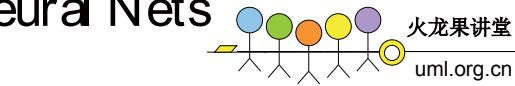
```
# input placeholders
X = tf.placeholder(tf.float32, [None, 784])
Y = tf.placeholder(tf.float32, [None, 10])

# weights & bias for nn layers
W1 = tf.Variable(tf.random_normal([784, 256]))
b1 = tf.Variable(tf.random_normal([256]))
L1 = tf.nn.relu(tf.matmul(X, W1) + b1)

W2 = tf.Variable(tf.random_normal([256, 256]))
b2 = tf.Variable(tf.random_normal([256]))
L2 = tf.nn.relu(tf.matmul(L1, W2) + b2)

W3 = tf.Variable(tf.random_normal([256, 10]))
b3 = tf.Variable(tf.random_normal([10]))
hypothesis = tf.matmul(L2, W3) + b3

# define cost/loss & optimizer
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(
    logits=hypothesis, labels=Y))
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost)
```



## NN tips



- Initializing weights
- Activation functions
- Regularization
- Optimizers

Random?

```
W = tf.Variable(tf.random_normal([1]), name='weight')
```



### How to do Xavier initialization on TensorFlow

I'm porting my Caffe network over to TensorFlow but it doesn't seem to have xavier initialization. I'm using `truncated_normal` but this seems to be making it a lot harder to train.

32

share edit flag

10

edited Dec 19 '15 at 22:25

asked Nov 10 '15 at 22:07  
 Hooked  
32.7k ● 13 ▲ 108 ● 169

Aleph7  
2,661 ● 1 ▲ 14 ● 24

Since version 0.8 there is a Xavier initializer, [see here for the docs](#).

56 You can use something like this:

```
W = tf.get_variable("W", shape=[784, 256],
                    initializer=tf.contrib.layers.xavier_initializer())
```

share edit delete flag

edited Dec 1 '16 at 16:32

fabian789  
5,351 ● 3 ▲ 32 ● 75

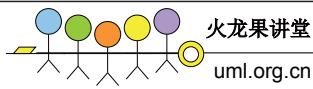
answered Apr 22 '16 at 4:23

Sung Kim  
2,451 ● 1 ▲ 16 ● 24

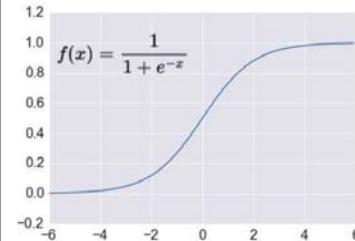
<http://stackoverflow.com/questions/33640581/>

## NN tips

- Initializing weights
- Activation functions
- Regularization
- Optimizers

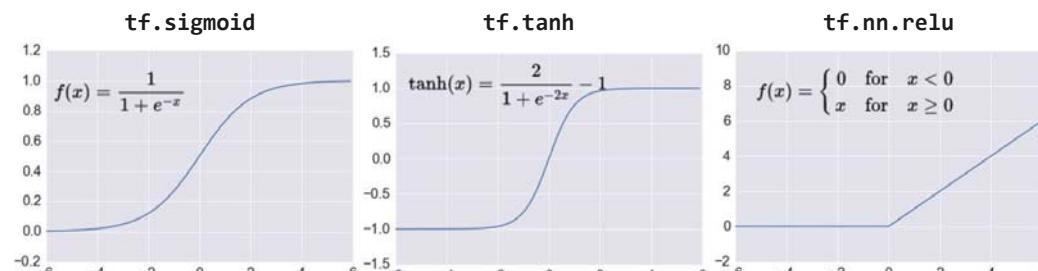
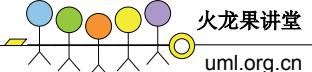


## Activation functions



<http://adilmoujahid.com/posts/2016/06/introduction-deep-learning-python-cafe/>

## Activation functions

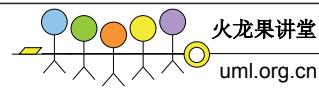
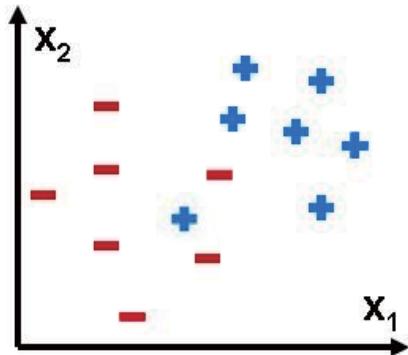


<http://adilmoujahid.com/posts/2016/06/introduction-deep-learning-python-cafe/>

## NN tips

- Initializing weights
- Activation functions
- Regularization
- Optimizers

## Overfitting

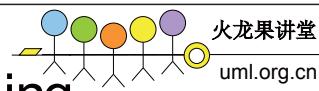


## Am I overfitting?

- Very high accuracy on the training dataset (eg: 0.99)
- Poor accuracy on the test data set (0.85)

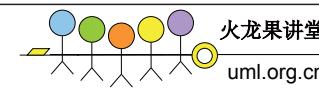
## Solutions for overfitting

- More training data!
- Reduce the number of features
- Regularization



## Regularization

- Let's not have too big numbers in the weight



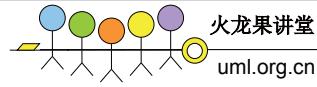
# Regularization

- Let's not have too big numbers in the weight

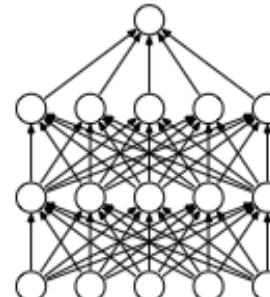
$$\cos t + \lambda \sum W^2$$

regulation strength

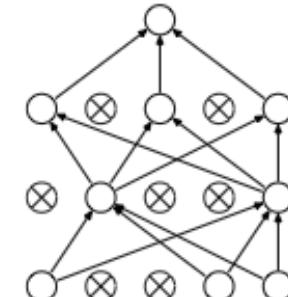
```
l2reg = 0.001 * tf.reduce_sum(tf.square(W))
```



## Dropout: A Simple Way to Prevent Neural Networks from Overfitting [Srivastava et al. 2014]



(a) Standard Neural Net

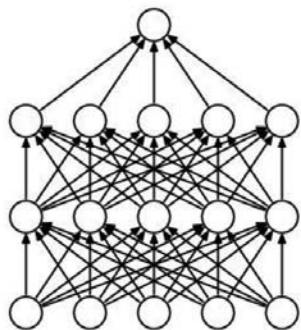


(b) After applying dropout.

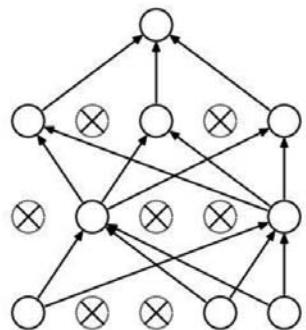


## Regularization: Dropout

“randomly set some neurons to zero in the forward pass”



(a) Standard Neural Net

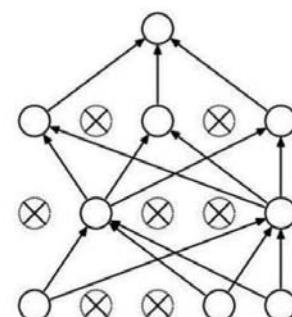


(b) After applying dropout.

[Srivastava et al., 2014]

Waaaait a second...

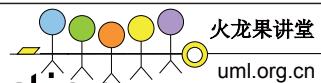
How could this possibly be a good idea?



Forces the network to have a redundant representation.



# TensorFlow implementation



火龙果讲堂  
uml.org.cn

```
keep_prob = tf.placeholder("float")  
  
L1 = ...  
  
L1_d = tf.nn.dropout(L1, keep_prob)
```

TRAIN:

```
sess.run(optimizer, feed_dict={X: batch_xs, Y: batch_ys, keep_prob: 0.7})
```

EVALUATION:

```
print "Accuracy:", accuracy.eval({X: mnist.test.images, Y: mnist.test.labels, keep_prob: 1})
```

```
# dropout (keep_prob) rate 0.7 on training, but should be 1 for testing  
keep_prob = tf.placeholder(tf.float32)
```

```
W1 = tf.get_variable("W1", shape=[784, 512])  
b1 = tf.Variable(tf.random_normal([512]))  
L1 = tf.nn.relu(tf.matmul(X, W1) + b1)  
L1 = tf.nn.dropout(L1, keep_prob=keep_prob)
```

```
W2 = tf.get_variable("W2", shape=[512, 512])  
b2 = tf.Variable(tf.random_normal([512]))  
L2 = tf.nn.relu(tf.matmul(L1, W2) + b2)  
L2 = tf.nn.dropout(L2, keep_prob=keep_prob)
```

```
...  
# train my model  
for epoch in range(training_epochs):  
    ...  
    for i in range(total_batch):  
        batch_xs, batch_ys = mnist.train.next_batch(batch_size)  
        feed_dict = {X: batch_xs, Y: batch_ys, keep_prob: 0.7}  
        c, _ = sess.run([cost, optimizer], feed_dict=feed_dict)  
        avg_cost += c / total_batch
```

```
# Test model and check accuracy  
correct_prediction = tf.equal(tf.argmax(hypothesis, 1), tf.argmax(Y, 1))  
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))  
print('Accuracy:', sess.run(accuracy, feed_dict={  
    X: mnist.test.images, Y: mnist.test.labels, keep_prob: 1}))
```

lab-10-5-mnist\_nn\_dropout.py

## NN tips



火龙果讲堂  
uml.org.cn

- Initializing weights
- Activation functions
- Regularization
- Optimizers

## Optimizers

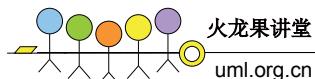
```
train = tf.train.GradientDescentOptimizer(learning_rate=0.1).minimize(cost)
```

[https://www.tensorflow.org/api\\_guides/python/train](https://www.tensorflow.org/api_guides/python/train)

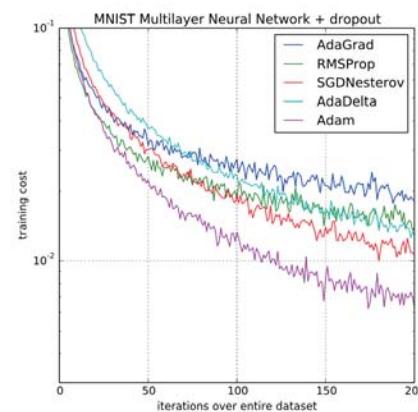
## Optimizers

```
train = tf.train.GradientDescentOptimizer(learning_rate=0.1).minimize(cost)
tf.train.AdadeltaOptimizer
tf.train.AdagradOptimizer
tf.train.AdagradDAOptimizer
tf.train.MomentumOptimizer
tf.train.AdamOptimizer
tf.train.FtrlOptimizer
tf.train.ProximalGradientDescentOptimizer
tf.train.ProximalAdagradOptimizer
tf.train.RMSPropOptimizer
```

[https://www.tensorflow.org/api\\_guides/python/train](https://www.tensorflow.org/api_guides/python/train)



## ADAM: a method for stochastic optimization [Kingma et al. 2015]



## Use Adam Optimizer

```
# define cost/loss & optimizer
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(
    logits=hypothesis, labels=Y))
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost)
```



## Convolutional Neural Networks

<http://cs231n.stanford.edu/>

## A bit of history:

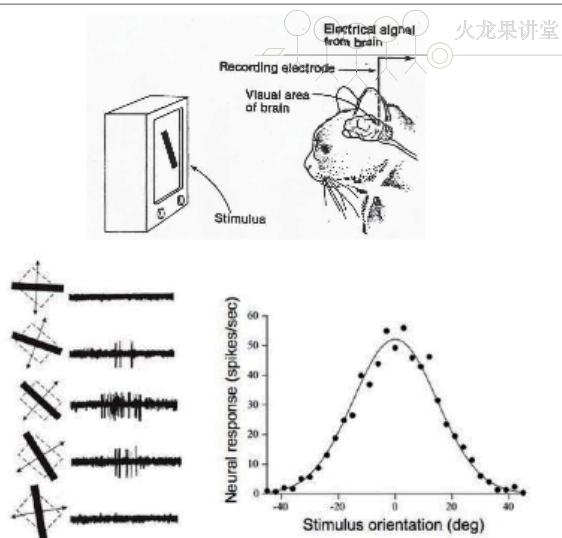
**Hubel & Wiesel,  
1959**

RECEPTIVE FIELDS OF SINGLE NEURONES IN THE CAT'S STRIATE CORTEX

**1962**

RECEPTIVE FIELDS, BINOCULAR INTERACTION AND FUNCTIONAL ARCHITECTURE IN THE CAT'S VISUAL CORTEX

**1968...**



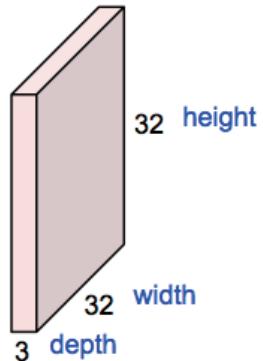
Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 7

27 Jan 2016

## Convolution Layer

32x32x3 image



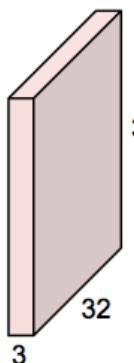
Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 10

27 Jan 2016

## Convolution Layer

32x32x3 image



5x5x3 filter



Convolve the filter with the image  
i.e. "slide over the image spatially,  
computing dot products"

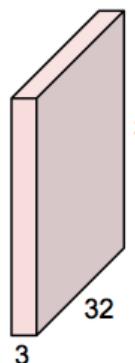
Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 11

27 Jan 2016

## Convolution Layer

32x32x3 image



5x5x3 filter



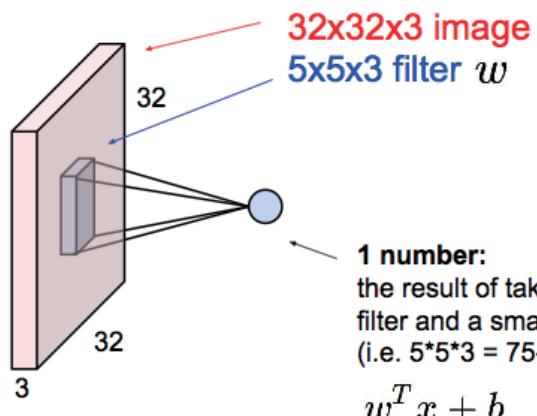
Convolve the filter with the image  
i.e. "slide over the image spatially,  
computing dot products"

Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 12

27 Jan 2016

## Convolution Layer

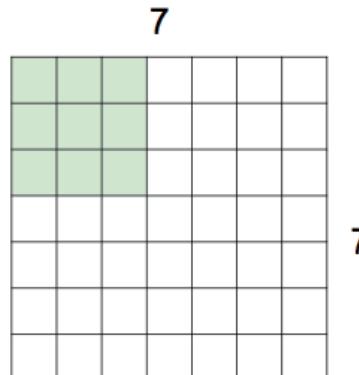


Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 13

27 Jan 2016

## A closer look at spatial dimensions:

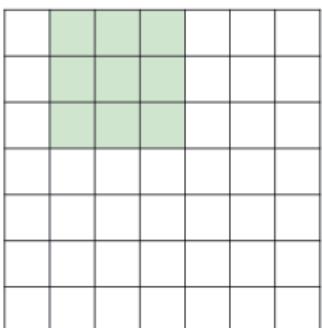


7x7 input (spatially)  
assume 3x3 filter

## A closer look at spatial dimensions:



7



7x7 input (spatially)  
assume 3x3 filter

Fei-Fei Li & Andrej Karpathy & Justin Johnson

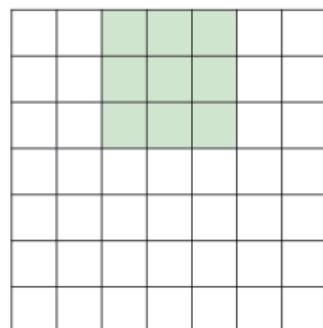
Lecture 7 - 25

27 Jan 2016

## A closer look at spatial dimensions:



7



7x7 input (spatially)  
assume 3x3 filter

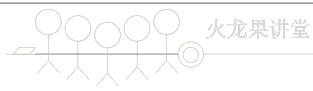
7

Fei-Fei Li & Andrej Karpathy & Justin Johnson

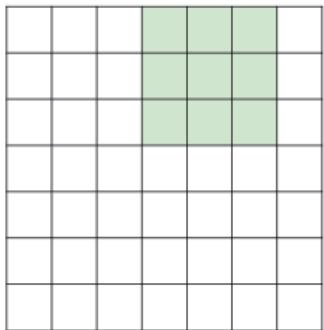
Lecture 7 - 26

27 Jan 2016

A closer look at spatial dimensions:



7

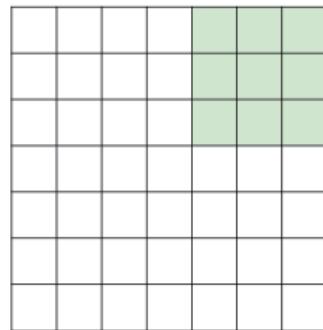


7x7 input (spatially)  
assume 3x3 filter

A closer look at spatial dimensions:



7



7x7 input (spatially)  
assume 3x3 filter

=> 5x5 output

Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 27

27 Jan 2016

Fei-Fei Li & Andrej Karpathy & Justin Johnson

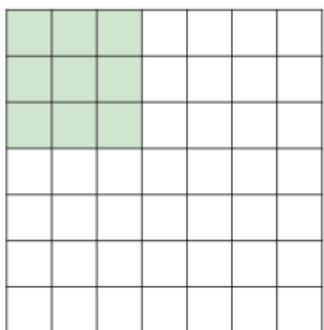
Lecture 7 - 28

27 Jan 2016

A closer look at spatial dimensions:



7

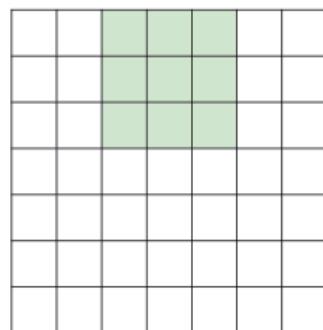


7x7 input (spatially)  
assume 3x3 filter  
applied **with stride 2**

A closer look at spatial dimensions:



7



7x7 input (spatially)  
assume 3x3 filter  
applied **with stride 2**

Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 29

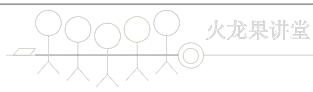
27 Jan 2016

Fei-Fei Li & Andrej Karpathy & Justin Johnson

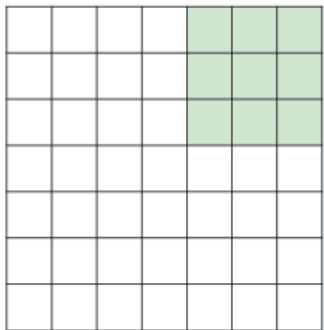
Lecture 7 - 30

27 Jan 2016

A closer look at spatial dimensions:



7



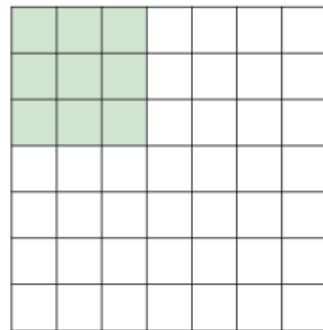
7

7x7 input (spatially)  
assume 3x3 filter  
applied **with stride 2**  
**=> 3x3 output!**

A closer look at spatial dimensions:



7



7

7x7 input (spatially)  
assume 3x3 filter  
applied **with stride 3?**

Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 31

27 Jan 2016

Fei-Fei Li & Andrej Karpathy & Justin Johnson

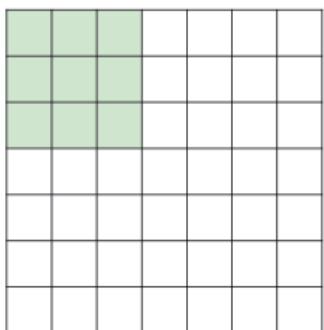
Lecture 7 - 32

27 Jan 2016

A closer look at spatial dimensions:



7

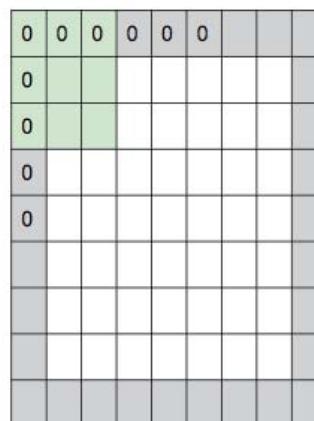


7

7x7 input (spatially)  
assume 3x3 filter  
applied **with stride 3?**

**doesn't fit!**  
cannot apply 3x3 filter on  
7x7 input with stride 3.

In practice: Common to zero pad the border



e.g. input 7x7  
3x3 filter, applied with **stride 1**  
**pad with 1 pixel border => what is the output?**

Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 33

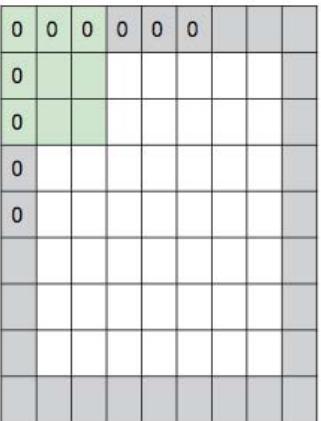
27 Jan 2016

Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 35

27 Jan 2016

## In practice: Common to zero pad the border



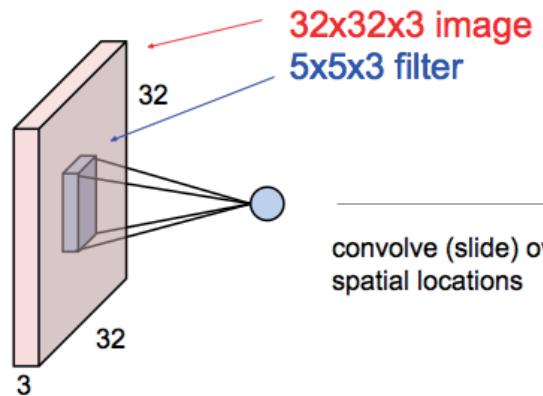
e.g. input 7x7

3x3 filter, applied with **stride 1**

pad with 1 pixel border => what is the output?

7x7 output!

## Convolution Layer



Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 36

27 Jan 2016

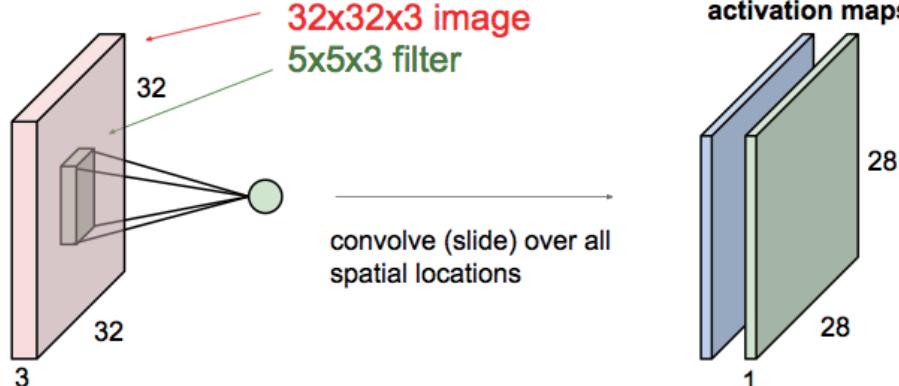
Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 14

27 Jan 2016

## Convolution Layer

consider a second, green filter

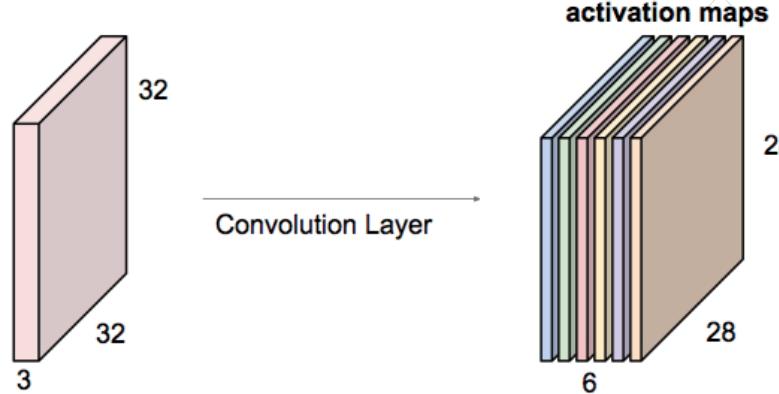


Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 15

27 Jan 2016

For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:



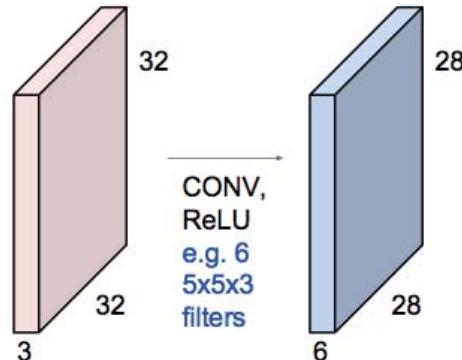
We stack these up to get a "new image" of size 28x28x6!

Fei-Fei Li & Andrej Karpathy & Justin Johnson

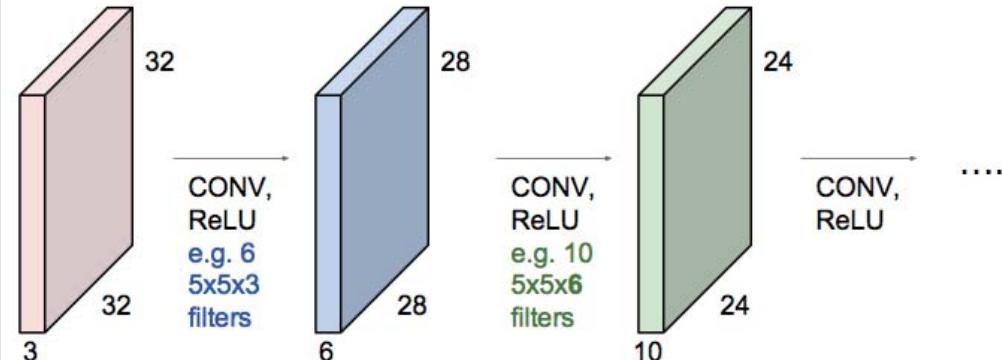
Lecture 7 - 16

27 Jan 2016

Preview: ConvNet is a sequence of Convolution Layers, interspersed with activation functions



Preview: ConvNet is a sequence of Convolutional Layers, interspersed with activation functions



Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 17

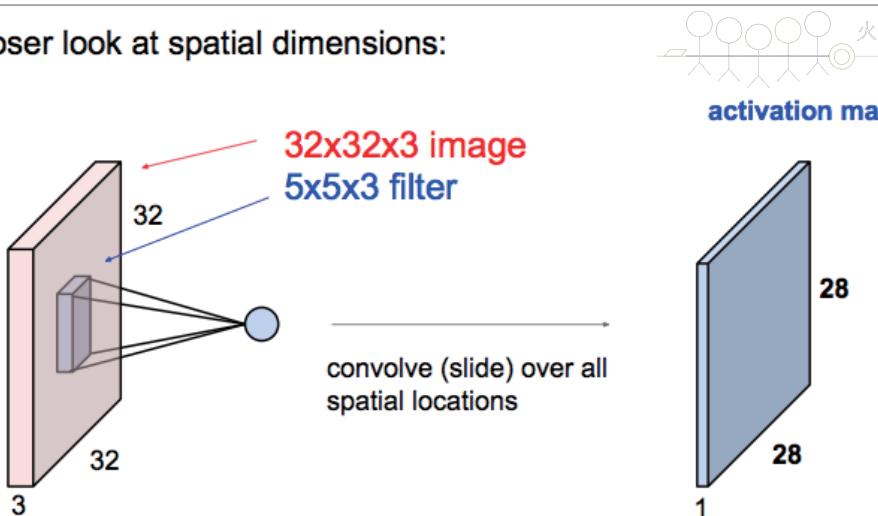
27 Jan 2016

Fei-Fei Li & Andrej Karpathy & Justin Johnson

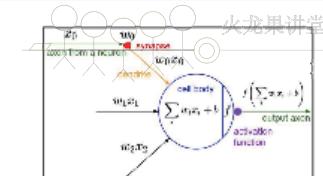
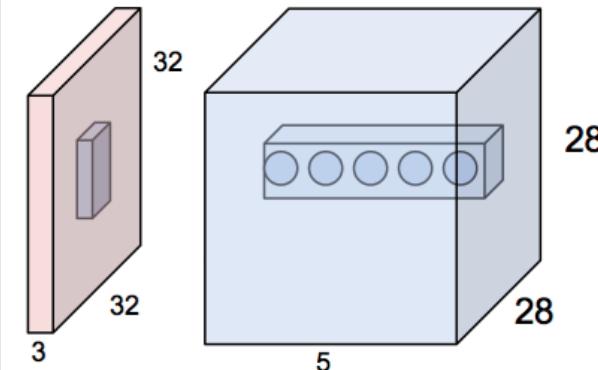
Lecture 7 - 18

27 Jan 2016

A closer look at spatial dimensions:



The brain/neuron view of CONV Layer



E.g. with 5 filters,  
CONV layer consists of  
neurons arranged in a 3D grid  
(28x28x5)

There will be 5 different  
neurons all looking at the same  
region in the input volume

Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 23

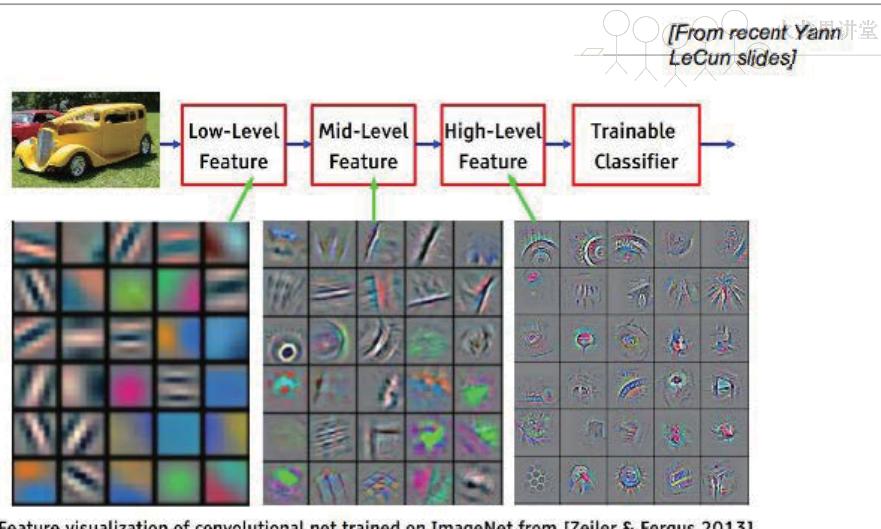
27 Jan 2016

Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 52

27 Jan 2016

## Preview



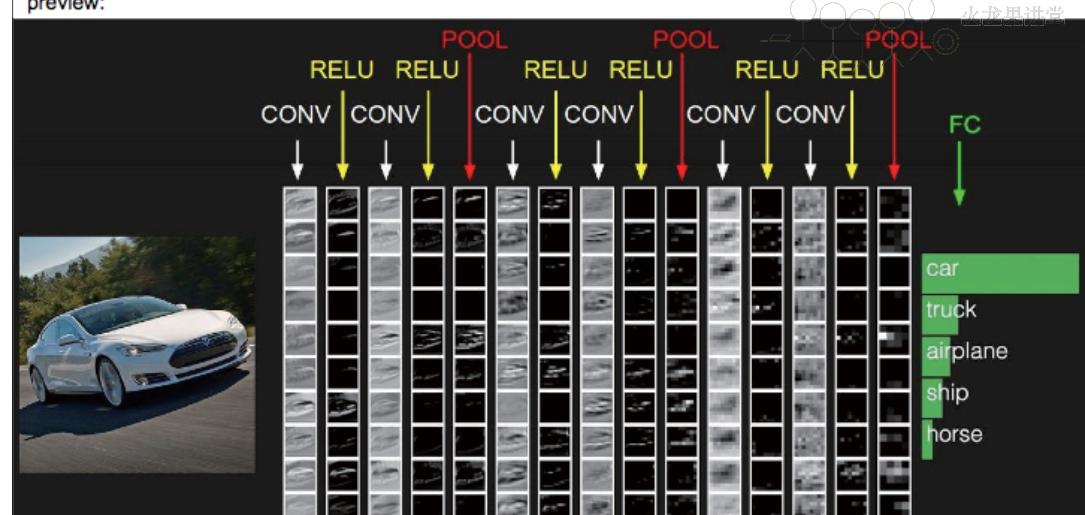
Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 19

27 Jan 2016

preview:



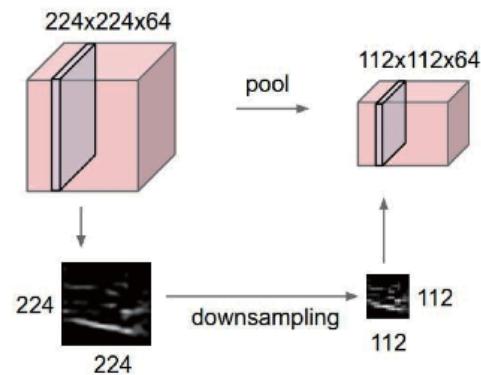
Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 22

27 Jan 2016

## Pooling layer

- makes the representations smaller and more manageable
- operates over each activation map independently:



Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 54

27 Jan 2016

## MAX POOLING

Single depth slice

x	1	1	2	4
5	6	7	8	
3	2	1	0	
1	2	3	4	

max pool with 2x2 filters  
and stride 2

6	8
3	4

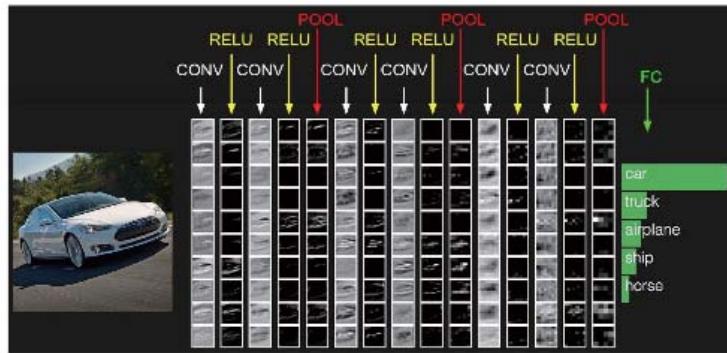
Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 55

27 Jan 2016

## Fully Connected Layer (FC layer)

- Contains neurons that connect to the entire input volume, as in ordinary Neural Networks



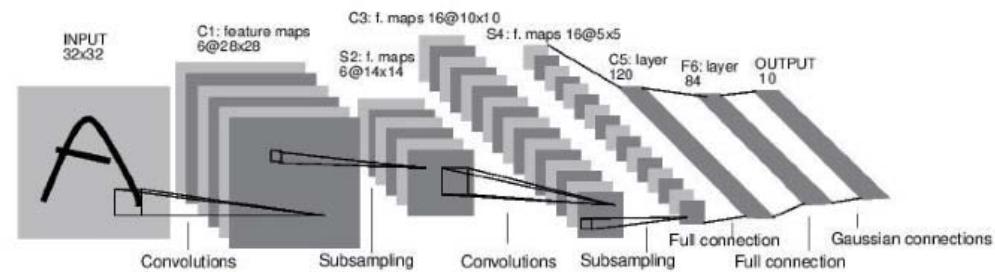
Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 58

27 Jan 2016

## Case Study: LeNet-5

[LeCun et al., 1998]



Conv filters were 5x5, applied at stride 1

Subsampling (Pooling) layers were 2x2 applied at stride 2  
i.e. architecture is [CONV-POOL-CONV-POOL-CONV-FC]

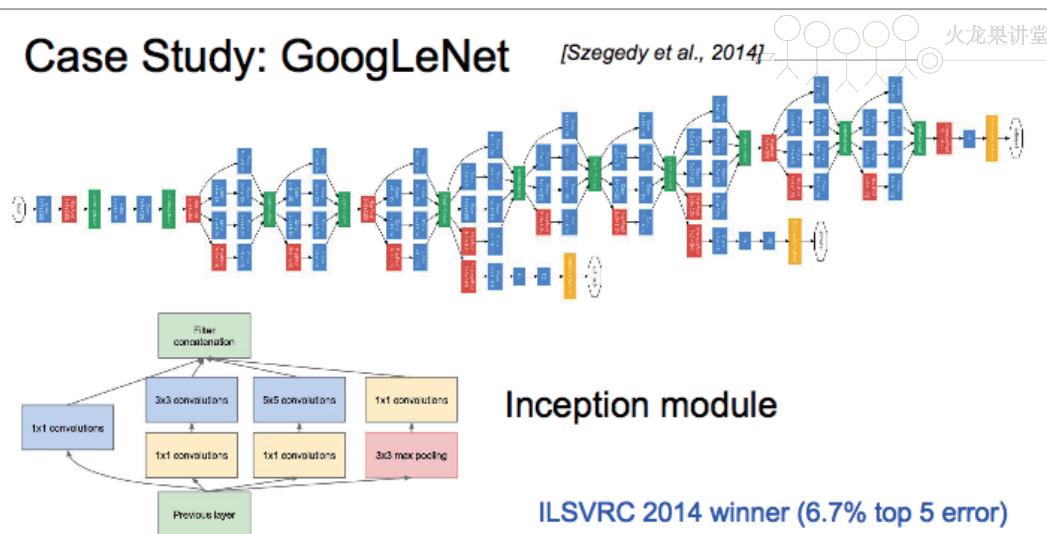
Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 60

27 Jan 2016

## Case Study: GoogLeNet

[Szegedy et al., 2014]



Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 75

27 Jan 2016

## Convolutional Neural Networks for Sentence Classification

[Yoon Kim, 2014]

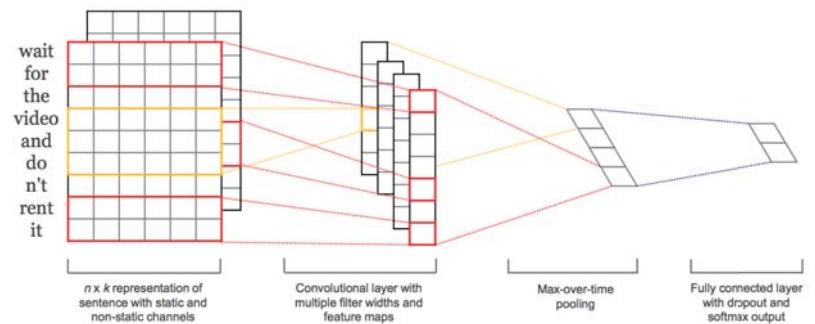
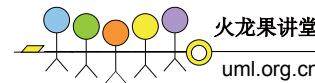
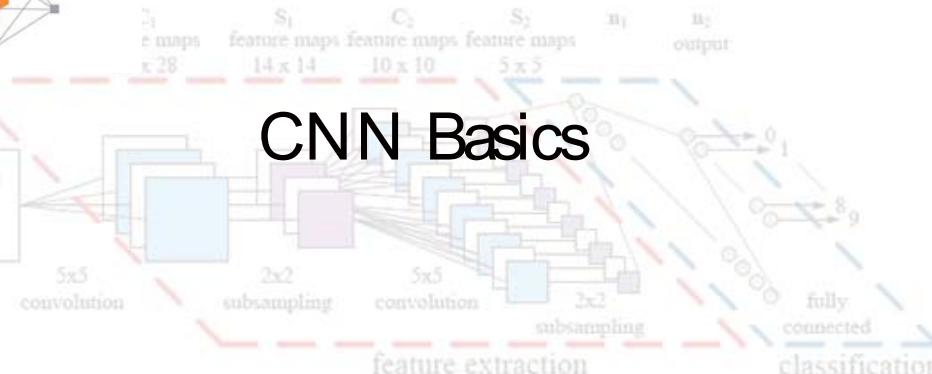


Figure 1: Model architecture with two channels for an example sentence.

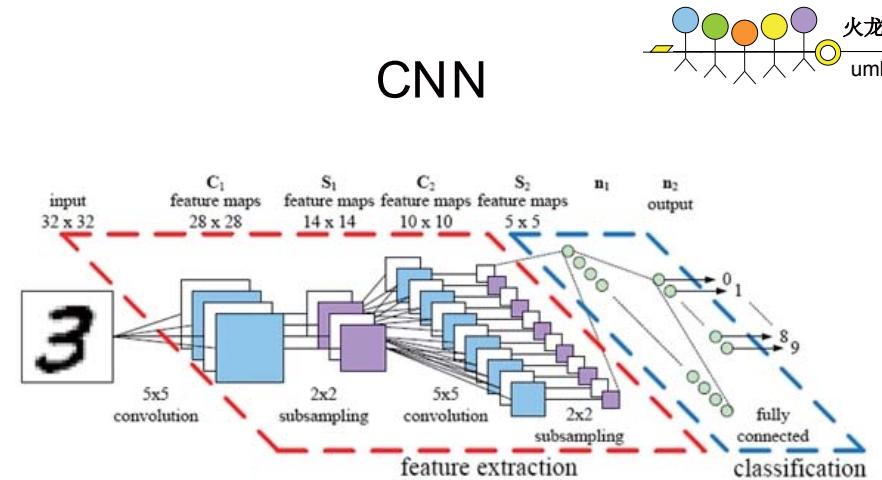
With TF 1.0!



## CNN Basics

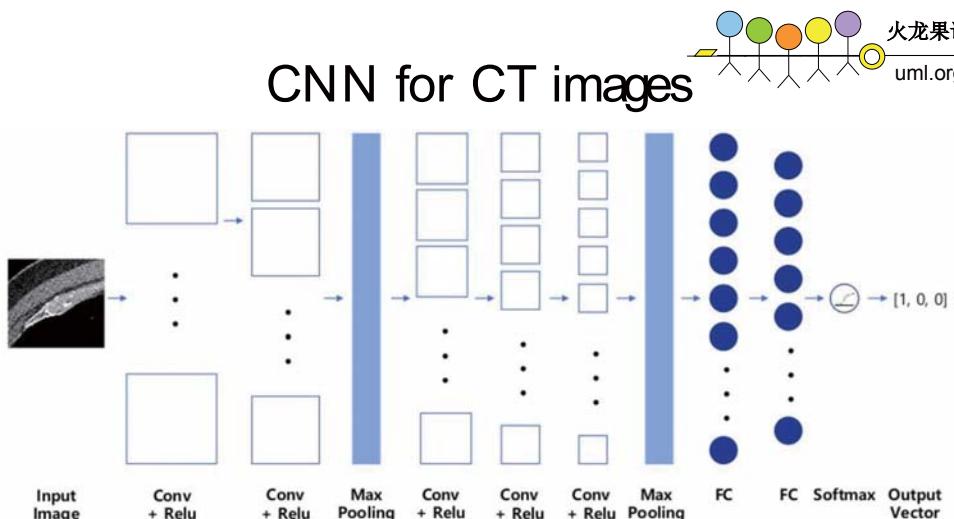


## CNN



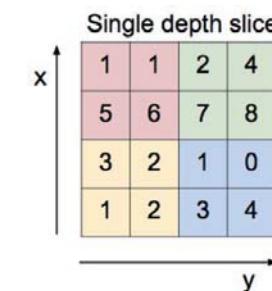
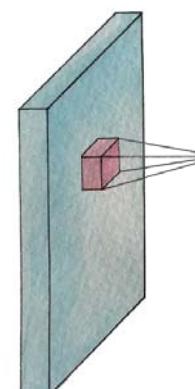
<http://parse.ele.tue.nl/cluster/2/CNNArchitecture.jpg>

## CNN for CT images



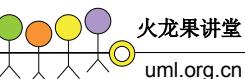
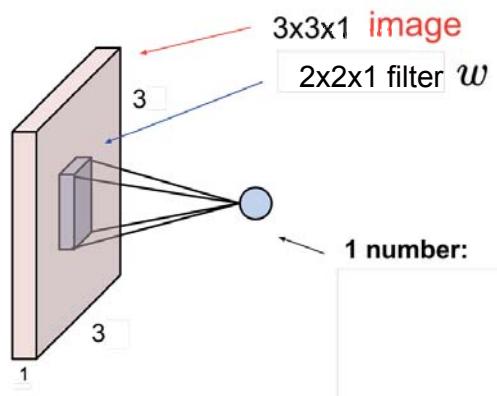
Asan Medical Center & Microsoft Medical Bigdata Contest Winner by GeunYoung Lee and Alex Kim  
<https://www.slideshare.net/GYLee3/ss-72966495>

## Convolution layer and max pooling



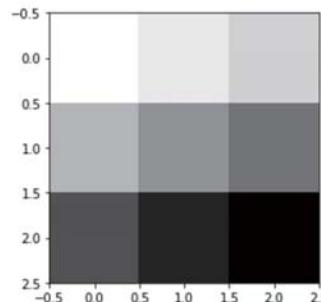
# Simple convolution layer

Stride: 1x1



```
In [2]: sess = tf.InteractiveSession()
image = np.array([[[[1],[2],[3]],
                  [[4],[5],[6]],
                  [[7],[8],[9]]]], dtype=np.float32)
print(image.shape)
plt.imshow(image.reshape(3,3), cmap='Greys')
```

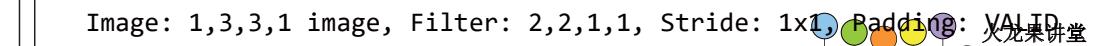
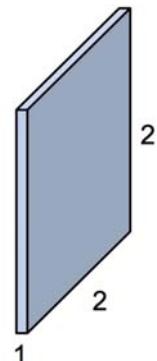
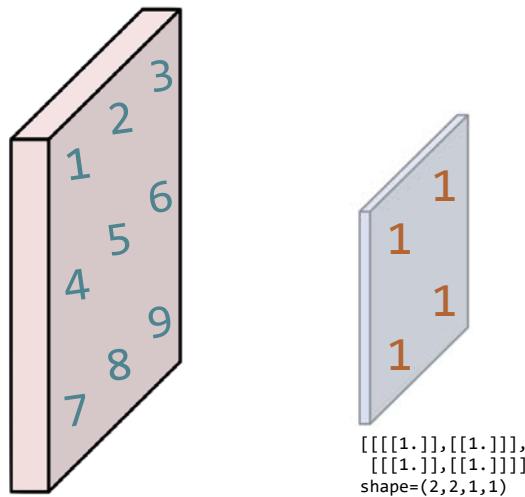
```
Out[2]: <matplotlib.image.AxesImage at 0x10db67dd8>
```



# Toy image

# Simple convolution layer

Image: 1,3,3,1 image, Filter: 2,2,1,1, Stride: 1x1, Padding: VALID



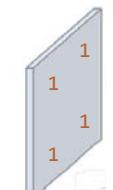
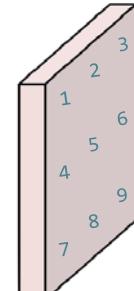
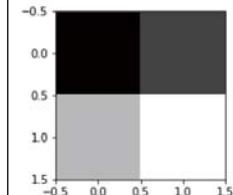
```

# print("image:\n", image)
print("image.shape", image.shape)
weight = tf.constant([[[[1.]], [[1.]]],  

                     [[[1.]], [[1.]]]])
print("weight.shape", weight.shape)
conv2d = tf.nn.conv2d(image, weight, strides=[1, 1, 1, 1], padding='VALID')
conv2d_img = conv2d.eval()
print("conv2d_img.shape", conv2d_img.shape)
conv2d_img = np.swapaxes(conv2d_img, 0, 3)
for i, one_img in enumerate(conv2d_img):
    print(one_img.reshape(2,2))
    plt.subplot(1,2,i+1), plt.imshow(one_img.reshape(2,2), cmap='gray')

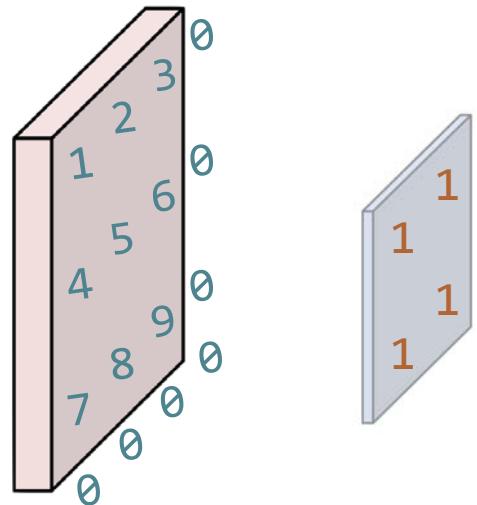
```

```
image.shape (1, 3, 3, 1)
weight.shape (2, 2, 1, 1)
conv2d_img.shape (1, 2, 2, 1)
[[ 12.  16.]
 [ 24.  28.]]
```



## Simple convolution layer

Image: 1,3,3,1 image, Filter: 2,2,1,1, Stride: 1x1, Padding: SAME



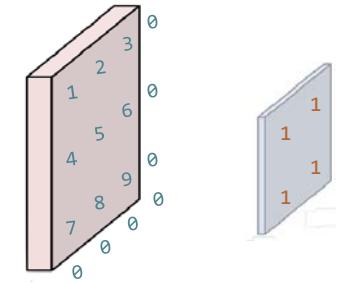
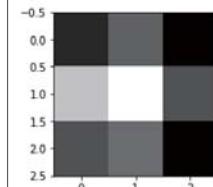
火龙果讲堂  
uml.org.cn

Image: 1,3,3,1 image, Filter: 2,2,1,1, Stride: 1x1, Padding: SAME

```
# print("image:\n", image)
print("image.shape", image.shape)

weight = tf.constant([[[[1.]], [[1.]]], [[[1.]], [[1.]]]])
print("weight.shape", weight.shape)
conv2d = tf.nn.conv2d(image, weight, strides=[1, 1, 1, 1], padding='SAME')
conv2d_img = conv2d.eval()
print("conv2d_img.shape", conv2d_img.shape)
conv2d_img = np.swapaxes(conv2d_img, 0, 3)
for i, one_img in enumerate(conv2d_img):
    print(one_img.reshape(3,3))
    plt.subplot(1,2,i+1), plt.imshow(one_img.reshape(3,3), cmap='gray')
```

image.shape (1, 3, 3, 1)  
weight.shape (2, 2, 1, 1)  
conv2d\_img.shape (1, 3, 3, 1)  
[[ 12. 16. 9.]  
[ 24. 28. 15.]  
[ 15. 17. 9.]]



## 3 filters (2,2,1,3)

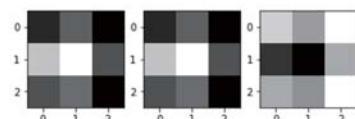
```
# print("image:\n", image)
print("image.shape", image.shape)
```

```
weight = tf.constant([[[[1.,10.,-1.],[1.,10.,-1.]],
[[1.,10.,-1.],[1.,10.,-1.]]], [[[1.,10.,-1.],[1.,10.,-1.]]])
print("weight.shape", weight.shape)
conv2d = tf.nn.conv2d(image, weight, strides=[1, 1, 1, 1], padding='SAME')
conv2d_img = conv2d.eval()
```

```
print("conv2d_img.shape", conv2d_img.shape)
conv2d_img = np.swapaxes(conv2d_img, 0, 3)
for i, one_img in enumerate(conv2d_img):
    print(one_img.reshape(3,3))
    plt.subplot(1,3,i+1), plt.imshow(one_img.reshape(3,3), cmap='gray')
```

image.shape (1, 3, 3, 1)  
weight.shape (2, 2, 1, 3)  
conv2d\_img.shape (1, 3, 3, 3)

[ 12. 16. 9.]  
[ 24. 28. 15.]  
[ 15. 17. 9.]  
[ 120. 160. 90.]  
[ 240. 280. 150.]  
[ 150. 170. 90.]]  
[-12. -16. -9.]  
[-24. -28. -15.]  
[-15. -17. -9.]]



火龙果讲堂  
uml.org.cn

## Max Pooling

4	3
2	1

```
In [19]: image = np.array([[[[4],[3]],
[[2],[1]]]], dtype=np.float32)
pool = tf.nn.max_pool(image, ksize=[1, 2, 2, 1],
strides=[1, 1, 1, 1], padding='SAME')
print(pool.shape)
print(pool.eval())
```

(1, 2, 2, 1)  
[[[[ 4.]
[ 3.]]]

SAME: Zero paddings

[[ 2.]
[ 1.]]]

4	3	0
2	1	0
0	0	0

4	3	0
2	1	0
0	0	0

4	3	0
2	1	0
0	0	0

4	3	0
2	1	0
0	0	0

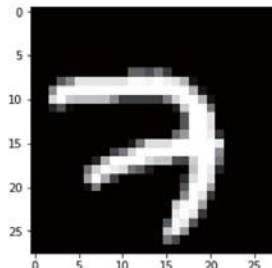
lab-11-0-cnn\_basics.ipynb

```
In [6]: from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)
# Check out https://www.tensorflow.org/get_started/mnist/beginners for
# more information about the mnist dataset
```

Extracting MNIST\_data/train-images-idx3-ubyte.gz  
 Extracting MNIST\_data/train-labels-idx1-ubyte.gz  
 Extracting MNIST\_data/t10k-images-idx3-ubyte.gz  
 Extracting MNIST\_data/t10k-labels-idx1-ubyte.gz

```
In [7]: img = mnist.train.images[0].reshape(28,28)
plt.imshow(img, cmap='gray')
```

```
Out[7]: <matplotlib.image.AxesImage at 0x115029ac8>
```



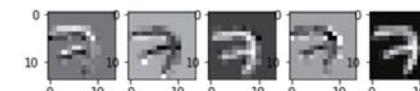
## MNIST image loading

[lab-11-0-cnn\\_basics.ipynb](#)

火龙果讲堂  
[uml.org.cn](http://uml.org.cn)

## MNIST Convolution layer

```
In [8]: sess = tf.InteractiveSession()
img = img.reshape(-1,28,28,1)
W1 = tf.Variable(tf.random_normal([3, 3, 1, 5], stddev=0.01))
conv2d = tf.nn.conv2d(img, W1, strides=[1, 2, 2, 1], padding='SAME')
print(conv2d)
sess.run(tf.global_variables_initializer())
conv2d_img = conv2d.eval()
conv2d_img = np.swapaxes(conv2d_img, 0, 3)
for i, one_img in enumerate(conv2d_img):
    plt.subplot(1,5,i+1), plt.imshow(one_img.reshape(14,14), cmap='gray')
```



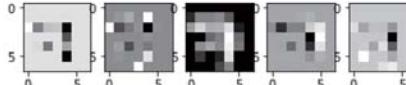
[lab-11-0-cnn\\_basics.ipynb](#)

## MNIST Max pooling



```
In [9]: pool = tf.nn.max_pool(conv2d, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding='SAME')
print(pool)
sess.run(tf.global_variables_initializer())
pool_img = pool.eval()
pool_img = np.swapaxes(pool_img, 0, 3)
for i, one_img in enumerate(pool_img):
    plt.subplot(1,5,i+1), plt.imshow(one_img.reshape(7, 7), cmap='gray')
```

```
Tensor("MaxPool_2:0", shape=(1, 7, 7, 5), dtype=float32)
```



[lab-11-0-cnn\\_basics.ipynb](#)

With TF 1.0!



$\mathbb{N}_1$  input feature maps  $\mathbb{N}_2$  feature maps

$\mathbb{N}_1 \times 28 \times 28$   $\mathbb{N}_2 \times 14 \times 14$

$C_1$  feature maps  $C_2$  feature maps

$\mathbb{N}_2 \times 10 \times 10$   $\mathbb{N}_3 \times 5 \times 5$

$S_1$  feature maps  $S_2$  feature maps

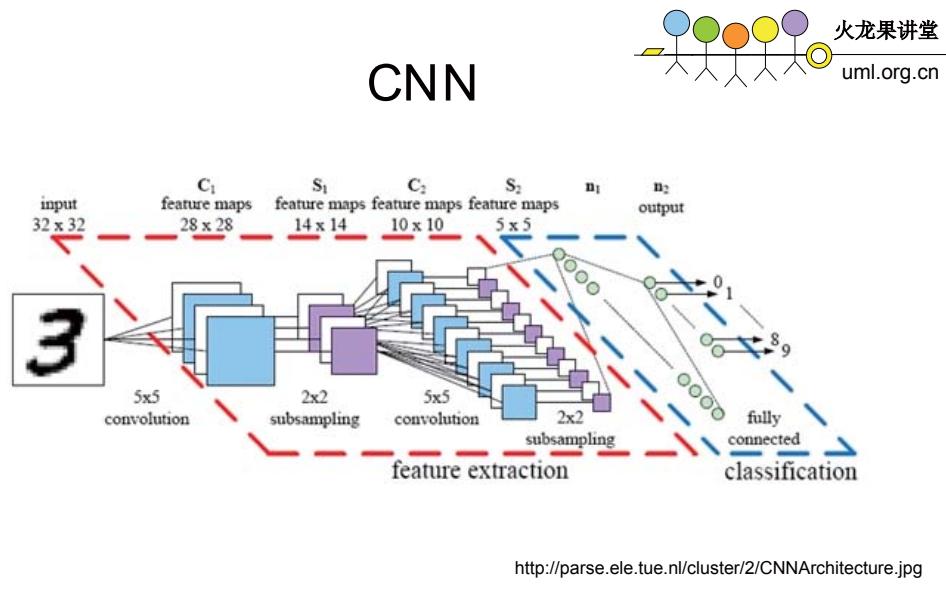
$\mathbb{N}_3 \times 5 \times 5$   $\mathbb{N}_4 \times 1 \times 1$

$n_1$  output  $n_2$  output

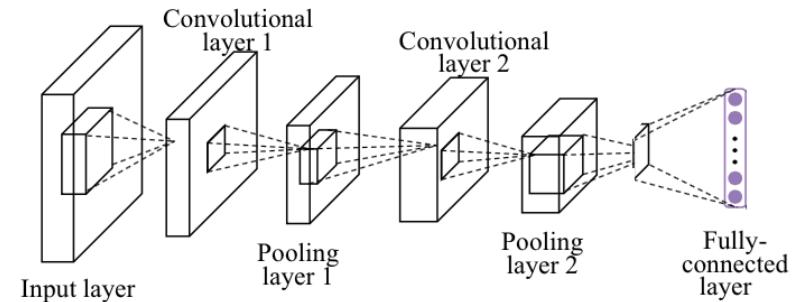
## CNN MNIST 99%



# CNN



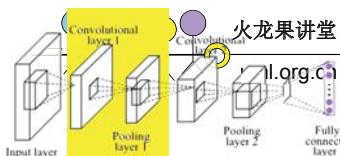
# Simple CNN



## Conv layer 1

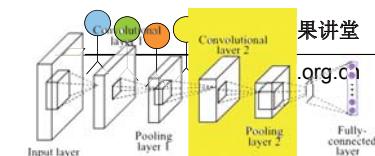
```
# input placeholders
X = tf.placeholder(tf.float32, [None, 784])
X_img = tf.reshape(X, [-1, 28, 28, 1]) # img 28x28x1 (black/white)
Y = tf.placeholder(tf.float32, [None, 10])

# L1 ImgIn shape=(?, 28, 28, 1)
W1 = tf.Variable(tf.random_normal([3, 3, 1, 32], stddev=0.01))
#   Conv      -> (?, 28, 28, 32)
#   Pool      -> (?, 14, 14, 32)
L1 = tf.nn.conv2d(X_img, W1, strides=[1, 1, 1, 1], padding='SAME')
L1 = tf.nn.relu(L1)
L1 = tf.nn.max_pool(L1, ksize=[1, 2, 2, 1],
                    strides=[1, 2, 2, 1], padding='SAME')
...
Tensor("Conv2D:0", shape=(?, 28, 28, 32), dtype=float32)
Tensor("Relu:0", shape=(?, 28, 28, 32), dtype=float32)
Tensor("MaxPool:0", shape=(?, 14, 14, 32), dtype=float32)
...
lab-11-1-mnist_cnn.py
```



## Conv layer 2

```
...
Tensor("Conv2D_0:0", shape=(?, 28, 28, 32), dtype=float32)
Tensor("Relu_0:0", shape=(?, 28, 28, 32), dtype=float32)
Tensor("MaxPool_0:0", shape=(?, 14, 14, 32), dtype=float32)
...
# L2 ImgIn shape=(?, 14, 14, 32)
W2 = tf.Variable(tf.random_normal([3, 3, 32, 64], stddev=0.01))
#   Conv      -> (?, 14, 14, 64)
#   Pool      -> (?, 7, 7, 64)
L2 = tf.nn.conv2d(L1, W2, strides=[1, 1, 1, 1], padding='SAME')
L2 = tf.nn.relu(L2)
L2 = tf.nn.max_pool(L2, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding='SAME')
L2 = tf.reshape(L2, [-1, 7 * 7 * 64])
...
Tensor("Conv2D_1:0", shape=(?, 14, 14, 64), dtype=float32)
Tensor("Relu_1:0", shape=(?, 14, 14, 64), dtype=float32)
Tensor("MaxPool_1:0", shape=(?, 7, 7, 64), dtype=float32)
Tensor("Reshape_1:0", shape=(?, 3136), dtype=float32)
...
lab-11-1-mnist_cnn.py
```



# Fully Connected (FC, Dense) layer

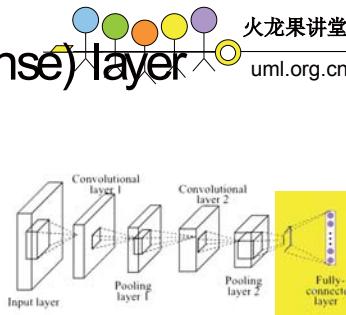
```

...
Tensor("Conv2D_1:0", shape=(?, 14, 14, 64), dtype=float32)
Tensor("Relu_1:0", shape=(?, 14, 14, 64), dtype=float32)
Tensor("MaxPool_1:0", shape=(?, 7, 7, 64), dtype=float32)
Tensor("Reshape_1:0", shape=(?, 3136), dtype=float32)
...
L2 = tf.reshape(L2, [-1, 7 * 7 * 64])

# Final FC 7x7x64 inputs -> 10 outputs
W3 = tf.get_variable("W3", shape=[7 * 7 * 64, 10],
initializer=tf.contrib.layers.xavier_initializer())
b = tf.Variable(tf.random_normal([10]))
hypothesis = tf.matmul(L2, W3) + b

# define cost/loss & optimizer
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=hypothesis,
labels=Y))
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost)
lab-11-1-mnist\_cnn.py

```



```

# initialize
sess = tf.Session()
sess.run(tf.global_variables_initializer())

```

```

# train my model
print('Learning stared. It takes sometime.')
for epoch in range(training_epochs):
    avg_cost = 0
    total_batch = int(mnist.train.num_examples / batch_size)
    for i in range(total_batch):
        batch_xs, batch_ys = mnist.train.next_batch(batch_size)
        feed_dict = {X: batch_xs, Y: batch_ys}
        c, _ = sess.run([cost, optimizer], feed_dict=feed_dict)
        avg_cost += c / total_batch
    print('Epoch:', '%04d' % (epoch + 1), 'cost =', '{:.9f}'.format(avg_cost))

print('Learning Finished!')

# Test model and check accuracy
correct_prediction = tf.equal(tf.argmax(hypothesis, 1), tf.argmax(Y, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
print('Accuracy:', sess.run(accuracy, feed_dict={X: mnist.test.images, Y:
mnist.test.labels})) lab-11-1-mnist\_cnn.py

```

# Training and Evaluation

```

# initialize
sess = tf.Session()
sess.run(tf.global_variables_initializer())

# train my model
print('Learning stared. It takes sometime.')
for epoch in range(training_epochs):
    avg_cost = 0
    total_batch = int(mnist.train.num_examples / batch_size)
    for i in range(total_batch):
        batch_xs, batch_ys = mnist.train.next_batch(batch_size)
        feed_dict = {X: batch_xs, Y: batch_ys}
        c, _ = sess.run([cost, optimizer], feed_dict=feed_dict)
        avg_cost += c / total_batch
    print('Epoch:', '%04d' % (epoch + 1), 'cost =', '{:.9f}'.format(avg_cost))

print('Learning Finished!')

# Test model and check accuracy
correct_prediction = tf.equal(tf.argmax(hypothesis, 1), tf.argmax(Y, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
print('Accuracy:', sess.run(accuracy, feed_dict={X: mnist.test.images, Y:
mnist.test.labels})) lab-11-1-mnist\_cnn.py

```

# Training and Evaluation

```

Epoch: 0001 cost = 0.340291267
Epoch: 0002 cost = 0.090731326
Epoch: 0003 cost = 0.064477619
Epoch: 0004 cost = 0.050683064
...
Epoch: 0011 cost = 0.017758641
Epoch: 0012 cost = 0.014156652
Epoch: 0013 cost = 0.012397016
Epoch: 0014 cost = 0.010693789
Epoch: 0015 cost = 0.009469977
Learning Finished!
Accuracy: 0.9885

```

# Deep CNN

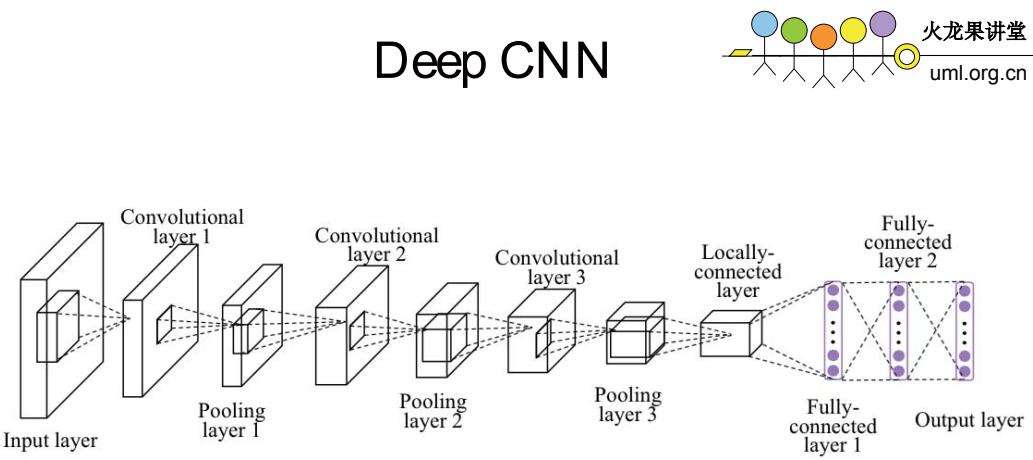


Image credit: [http://personal.ie.cuhk.edu.hk/~ccloy/project\\_target\\_code/index.html](http://personal.ie.cuhk.edu.hk/~ccloy/project_target_code/index.html)

```

# L1 ImgIn shape=(?, 28, 28, 1)
W1 = tf.Variable(tf.random_normal([3, 3, 1, 32], stddev=0.01))
# Conv -> (?, 28, 28, 32)
# Pool -> (?, 14, 14, 32)
L1 = tf.nn.conv2d(X_img, W1, strides=[1, 1, 1, 1], padding='SAME')
L1 = tf.nn.relu(L1)
L1 = tf.nn.max_pool(L1, ksize=[1, 2, 2, 1],
                    strides=[1, 2, 2, 1], padding='SAME')
L1 = tf.nn.dropout(L1, keep_prob=keep_prob)
'''Tensor("Conv2D_0:", shape=(?, 28, 28, 32), dtype=float32)
Tensor("Relu_0:", shape=(?, 28, 28, 32), dtype=float32)
Tensor("MaxPool_0:", shape=(?, 14, 14, 32), dtype=float32)
Tensor("dropout/mul_0:", shape=(?, 14, 14, 32), dtype=float32)'''

# L2 ImgIn shape=(?, 14, 14, 32)
W2 = tf.Variable(tf.random_normal([3, 3, 32, 64], stddev=0.01))
# Conv -> (?, 14, 14, 64)
# Pool -> (?, 7, 7, 64)
L2 = tf.nn.conv2d(L1, W2, strides=[1, 1, 1, 1], padding='SAME')
L2 = tf.nn.relu(L2)
L2 = tf.nn.max_pool(L2, ksize=[1, 2, 2, 1],
                    strides=[1, 2, 2, 1], padding='SAME')
L2 = tf.nn.dropout(L2, keep_prob=keep_prob)
'''Tensor("Conv2D_1:", shape=(?, 14, 14, 64), dtype=float32)
Tensor("Relu_1:", shape=(?, 14, 14, 64), dtype=float32)
Tensor("MaxPool_1:", shape=(?, 7, 7, 64), dtype=float32)
Tensor("dropout_1/mul_0:", shape=(?, 7, 7, 64), dtype=float32)'''

# L3 ImgIn shape=(?, 7, 7, 64)
W3 = tf.Variable(tf.random_normal([3, 3, 64, 128], stddev=0.01))
# Conv -> (?, 7, 7, 128)
# Pool -> (?, 4, 4, 128)
# Reshape -> (?, 4 * 4 * 128) # Flatten them for FC
L3 = tf.nn.conv2d(L2, W3, strides=[1, 1, 1, 1], padding='SAME')
L3 = tf.nn.relu(L3)
L3 = tf.nn.max_pool(L3, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1],
                    padding='SAME')
L3 = tf.nn.dropout(L3, keep_prob=keep_prob)
L3 = tf.reshape(L3, [-1, 128 * 4 * 4])
'''Tensor("Conv2D_2:", shape=(?, 7, 7, 128), dtype=float32)
Tensor("Relu_2:", shape=(?, 7, 7, 128), dtype=float32)
Tensor("MaxPool_2:", shape=(?, 4, 4, 128), dtype=float32)
Tensor("dropout_2/mul_0:", shape=(?, 4, 4, 128), dtype=float32)
Tensor("Reshape_1:", shape=(?, 2048), dtype=float32)'''

# L4 FC 4x4x128 inputs -> 625 outputs
W4 = tf.get_variable("W4", shape=[128 * 4 * 4, 625],
                     initializer=tf.contrib.layers.xavier_initializer())
b4 = tf.Variable(tf.random_normal([625]))
L4 = tf.nn.relu(tf.matmul(L3, W4) + b4)
L4 = tf.nn.dropout(L4, keep_prob=keep_prob)
'''Tensor("Relu_3:", shape=(?, 625), dtype=float32)
Tensor("dropout_3/mul_0:", shape=(?, 625), dtype=float32)
Tensor("Relu_3_0:", shape=(?, 625), dtype=float32)
Tensor("dropout_3/mul_0:", shape=(?, 625), dtype=float32)'''

# L5 Final FC 625 inputs -> 10 outputs
W5 = tf.get_variable("W5", shape=[625, 10],
                     initializer=tf.contrib.layers.xavier_initializer())
b5 = tf.Variable(tf.random_normal([10]))
hypothesis = tf.matmul(L4, W5) + b5
'''Tensor("add_1:", shape=(?, 10), dtype=float32)'''

```

[lab-11-2-mnist\\_deep\\_cnn.py](#)

```

# L1 ImgIn shape=(?, 28, 28, 1)
W1 = tf.Variable(tf.random_normal([3, 3, 1, 32], stddev=0.01))
# Conv -> (?, 28, 28, 32)
# Pool -> (?, 14, 14, 32)
L1 = tf.nn.conv2d(X_img, W1, strides=[1, 1, 1, 1], padding='SAME')
L1 = tf.nn.relu(L1)
L1 = tf.nn.max_pool(L1, ksize=[1, 2, 2, 1],
                    strides=[1, 2, 2, 1], padding='SAME')
L1 = tf.nn.dropout(L1, keep_prob=keep_prob)
'''Tensor("Conv2D_0:", shape=(?, 28, 28, 32), dtype=float32)
Tensor("Relu_0:", shape=(?, 28, 28, 32), dtype=float32)
Tensor("MaxPool_0:", shape=(?, 14, 14, 32), dtype=float32)
Tensor("dropout/mul_0:", shape=(?, 14, 14, 32), dtype=float32)'''

# L2 ImgIn shape=(?, 14, 14, 32)
W2 = tf.Variable(tf.random_normal([3, 3, 32, 64], stddev=0.01))
# Conv -> (?, 14, 14, 64)
# Pool -> (?, 7, 7, 64)
L2 = tf.nn.conv2d(L1, W2, strides=[1, 1, 1, 1], padding='SAME')
L2 = tf.nn.relu(L2)
L2 = tf.nn.max_pool(L2, ksize=[1, 2, 2, 1],
                    strides=[1, 2, 2, 1], padding='SAME')
L2 = tf.nn.dropout(L2, keep_prob=keep_prob)
'''Tensor("Conv2D_1:", shape=(?, 14, 14, 64), dtype=float32)
Tensor("Relu_1:", shape=(?, 14, 14, 64), dtype=float32)
Tensor("MaxPool_1:", shape=(?, 7, 7, 64), dtype=float32)
Tensor("dropout_1/mul_0:", shape=(?, 7, 7, 64), dtype=float32)'''

# L3 ImgIn shape=(?, 7, 7, 64)
W3 = tf.Variable(tf.random_normal([3, 3, 64, 128], stddev=0.01))
# Conv -> (?, 7, 7, 128)
# Pool -> (?, 4, 4, 128)
# Reshape -> (?, 4 * 4 * 128) # Flatten them for FC
L3 = tf.nn.conv2d(L2, W3, strides=[1, 1, 1, 1], padding='SAME')
L3 = tf.nn.relu(L3)
L3 = tf.nn.max_pool(L3, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1],
                    padding='SAME')
L3 = tf.nn.dropout(L3, keep_prob=keep_prob)
L3 = tf.reshape(L3, [-1, 128 * 4 * 4])
'''Tensor("Conv2D_2:", shape=(?, 7, 7, 128), dtype=float32)
Tensor("Relu_2:", shape=(?, 7, 7, 128), dtype=float32)
Tensor("MaxPool_2:", shape=(?, 4, 4, 128), dtype=float32)
Tensor("dropout_2/mul_0:", shape=(?, 4, 4, 128), dtype=float32)
Tensor("Reshape_1:", shape=(?, 2048), dtype=float32)'''

# L4 FC 4x4x128 inputs -> 625 outputs
W4 = tf.get_variable("W4", shape=[128 * 4 * 4, 625],
                     initializer=tf.contrib.layers.xavier_initializer())
b4 = tf.Variable(tf.random_normal([625]))
L4 = tf.nn.relu(tf.matmul(L3, W4) + b4)
L4 = tf.nn.dropout(L4, keep_prob=keep_prob)
'''Tensor("Relu_3:", shape=(?, 625), dtype=float32)
Tensor("dropout_3/mul_0:", shape=(?, 625), dtype=float32)
Tensor("Relu_3_0:", shape=(?, 625), dtype=float32)
Tensor("dropout_3/mul_0:", shape=(?, 625), dtype=float32)'''

# L5 Final FC 625 inputs -> 10 outputs
W5 = tf.get_variable("W5", shape=[625, 10],
                     initializer=tf.contrib.layers.xavier_initializer())
b5 = tf.Variable(tf.random_normal([10]))
hypothesis = tf.matmul(L4, W5) + b5
'''Tensor("add_1:", shape=(?, 10), dtype=float32)'''

```

[lab-11-2-mnist\\_deep\\_cnn.py](#)



```

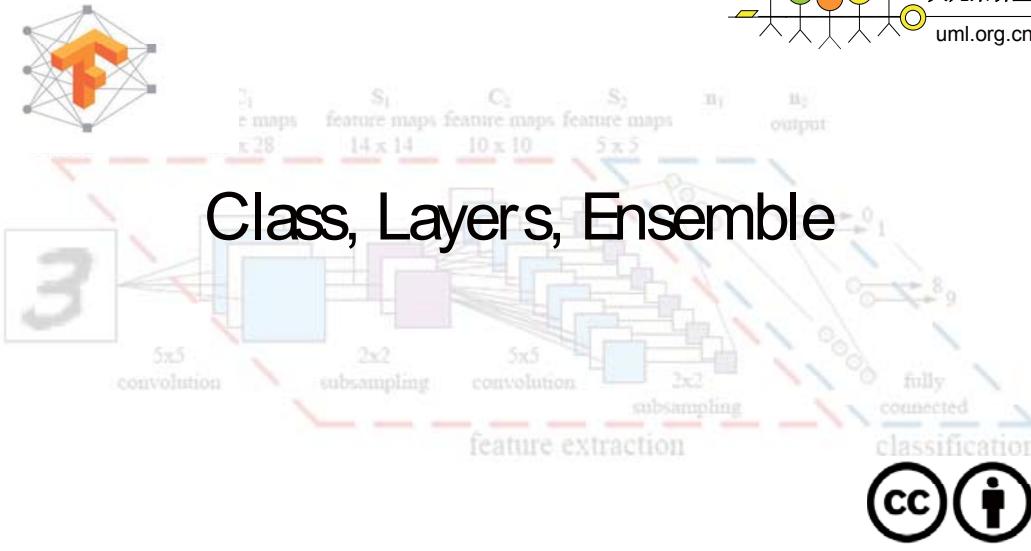
# Test model and check accuracy
correct_prediction = tf.equal(tf.argmax(hypothesis, 1),
                              tf.argmax(Y, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
print('Accuracy:', sess.run(accuracy,
                           feed_dict={X: mnist.test.images,
                                      Y: mnist.test.labels, keep_prob: 1}))

```

Epoch: 0013 cost = 0.027188021  
 Epoch: 0014 cost = 0.023604777  
 Epoch: 0015 cost = 0.024607201  
 Learning Finished!

Accuracy: **0.9938**

With TF 1.0!



火龙果讲堂

uml.org.cn

```

# L1 ImgIn shape=(?, 28, 28, 1)
W1 = tf.Variable(tf.random_normal([3, 3, 1, 32], stddev=0.01))
# Conv -> (?, 28, 28, 32)
# Pool -> (?, 14, 14, 32)
L1 = tf.nn.conv2d(X_img, W1, strides=[1, 1, 1, 1], padding='SAME')
L1 = tf.nn.relu(L1)
L1 = tf.nn.max_pool(L1, ksize=[1, 2, 2, 1],
                    strides=[1, 2, 2, 1], padding='SAME')
L1 = tf.nn.dropout(L1, keep_prob=keep_prob)
'''Tensor("Conv2D_0:", shape=(?, 28, 28, 32), dtype=float32)
Tensor("Relu_0:", shape=(?, 28, 28, 32), dtype=float32)
Tensor("MaxPool_0:", shape=(?, 14, 14, 32), dtype=float32)
Tensor("dropout/mul_0:", shape=(?, 14, 14, 32), dtype=float32)'''

# L2 ImgIn shape=(?, 14, 14, 32)
W2 = tf.Variable(tf.random_normal([3, 3, 32, 64], stddev=0.01))
# Conv -> (?, 14, 14, 64)
# Pool -> (?, 7, 7, 64)
L2 = tf.nn.conv2d(L1, W2, strides=[1, 1, 1, 1], padding='SAME')
L2 = tf.nn.relu(L2)
L2 = tf.nn.max_pool(L2, ksize=[1, 2, 2, 1],
                    strides=[1, 2, 2, 1], padding='SAME')
L2 = tf.nn.dropout(L2, keep_prob=keep_prob)
'''Tensor("Conv2D_1:", shape=(?, 14, 14, 64), dtype=float32)
Tensor("Relu_1:", shape=(?, 14, 14, 64), dtype=float32)
Tensor("MaxPool_1:", shape=(?, 7, 7, 64), dtype=float32)
Tensor("dropout_1/mul_0:", shape=(?, 7, 7, 64), dtype=float32)'''

# L3 ImgIn shape=(?, 7, 7, 64)
W3 = tf.Variable(tf.random_normal([3, 3, 64, 128], stddev=0.01))
# Conv -> (?, 7, 7, 128)
# Pool -> (?, 4, 4, 128)
# Reshape -> (?, 4 * 4 * 128) # Flatten them for FC
L3 = tf.nn.conv2d(L2, W3, strides=[1, 1, 1, 1], padding='SAME')
L3 = tf.nn.relu(L3)
L3 = tf.nn.max_pool(L3, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1],
                    padding='SAME')
L3 = tf.nn.dropout(L3, keep_prob=keep_prob)
L3 = tf.reshape(L3, [-1, 128 * 4 * 4])
'''Tensor("Conv2D_2:", shape=(?, 7, 7, 128), dtype=float32)
Tensor("Relu_2:", shape=(?, 7, 7, 128), dtype=float32)
Tensor("MaxPool_2:", shape=(?, 4, 4, 128), dtype=float32)
Tensor("dropout_2/mul_0:", shape=(?, 4, 4, 128), dtype=float32)
Tensor("Reshape_1:", shape=(?, 2048), dtype=float32)'''

# L4 FC 4x4x128 inputs -> 625 outputs
W4 = tf.get_variable("W4", shape=[128 * 4 * 4, 625],
                     initializer=tf.contrib.layers.xavier_initializer())
b4 = tf.Variable(tf.random_normal([625]))
L4 = tf.nn.relu(tf.matmul(L3, W4) + b4)
L4 = tf.nn.dropout(L4, keep_prob=keep_prob)
'''Tensor("Relu_3:", shape=(?, 625), dtype=float32)
Tensor("dropout_3/mul_0:", shape=(?, 625), dtype=float32)
Tensor("Relu_3_0:", shape=(?, 625), dtype=float32)
Tensor("dropout_3/mul_0:", shape=(?, 625), dtype=float32)'''

# L5 Final FC 625 inputs -> 10 outputs
W5 = tf.get_variable("W5", shape=[625, 10],
                     initializer=tf.contrib.layers.xavier_initializer())
b5 = tf.Variable(tf.random_normal([10]))
hypothesis = tf.matmul(L4, W5) + b5
'''Tensor("add_1:", shape=(?, 10), dtype=float32)'''

```

[lab-11-2-mnist\\_deep\\_cnn.py](#)



```

# Test model and check accuracy
correct_prediction = tf.equal(tf.argmax(hypothesis, 1),
                              tf.argmax(Y, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
print('Accuracy:', sess.run(accuracy,
                           feed_dict={X: mnist.test.images,
                                      Y: mnist.test.labels, keep_prob: 1}))

```

Epoch: 0013 cost = 0.027188021  
 Epoch: 0014 cost = 0.023604777  
 Epoch: 0015 cost = 0.024607201  
 Learning Finished!

Accuracy: **0.9938**

```

class Model:
    def __init__(self, sess, name):
        self.sess = sess
        self.name = name
        self._build_net()

    def _build_net(self):
        with tf.variable_scope(self.name):
            # input place holders
            self.X = tf.placeholder(tf.float32, [None, 784])
            # img 28x28x1 (black/white)
            X_img = tf.reshape(self.X, [-1, 28, 28, 1])
            self.Y = tf.placeholder(tf.float32, [None, 10])

            # L1 ImgIn shape=(?, 28, 28, 1)
            W1 = tf.Variable(tf.random_normal([3, 3, 1, 32],
                stddev=0.01))

            ...
            ...

    def predict(self, x_test, keep_prop=1.0):
        return self.sess.run(self.logits,
            feed_dict={self.X: x_test, self.keep_prob: keep_prop})

    def get_accuracy(self, x_test, y_test, keep_prop=1.0):
        return self.sess.run(self.accuracy,
            feed_dict={self.X: x_test, self.Y: y_test, self.keep_prob: keep_prop})

    def train(self, x_data, y_data, keep_prop=0.7):
        return self.sess.run([self.cost, self.optimizer], feed_dict={
            self.X: x_data, self.Y: y_data, self.keep_prob: keep_prop})

```

[lab-11-3-mnist\\_cnn\\_class.py](#)

**Python Class** 火龙果讲堂 uml.org.cn

```

# initialize
sess = tf.Session()
m1 = Model(sess, "m1")

sess.run(tf.global_variables_initializer())

print('Learning Started!')

# train my model
for epoch in range(training_epochs):
    avg_cost = 0
    total_batch = int(mnist.train.num_examples / batch_size)

    for i in range(total_batch):
        batch_xs, batch_ys = mnist.train.next_batch(batch_size)
        c, _ = m1.train(batch_xs, batch_ys)
        avg_cost += c / total_batch

```

**tf.layers** 火龙果讲堂 uml.org.cn

**average\_pooling1d(...)**: Average Pooling layer for 1D inputs.  
**average\_pooling2d(...)**: Average pooling layer for 2D inputs (e.g. images).  
**average\_pooling3d(...)**: Average pooling layer for 3D inputs (e.g. volumes).  
**batch\_normalization(...)**: Functional interface for the batch normalization layer.  
**conv1d(...)**: Functional interface for 1D convolution layer (e.g. temporal convolution).  
**conv2d(...)**: Functional interface for the 2D convolution layer.  
**conv2d\_transpose(...)**: Transposed convolution layer (sometimes called Deconvolution).  
**conv3d(...)**: Functional interface for the 3D convolution layer.  
**dense(...)**: Functional interface for the densely-connected layer.  
**dropout(...)**: Applies Dropout to the input.  
**max\_pooling1d(...)**: Max Pooling layer for 1D inputs.  
**max\_pooling2d(...)**: Max pooling layer for 2D inputs (e.g. images).  
**max\_pooling3d(...)**: Max pooling layer for 3D inputs (e.g. volumes).  
**separable\_conv2d(...)**: Functional interface for the depthwise separable 2D convolution layer.

[https://www.tensorflow.org/api\\_docs/python/tf/layers](https://www.tensorflow.org/api_docs/python/tf/layers)

**tf.layers** 火龙果讲堂 uml.org.cn

```

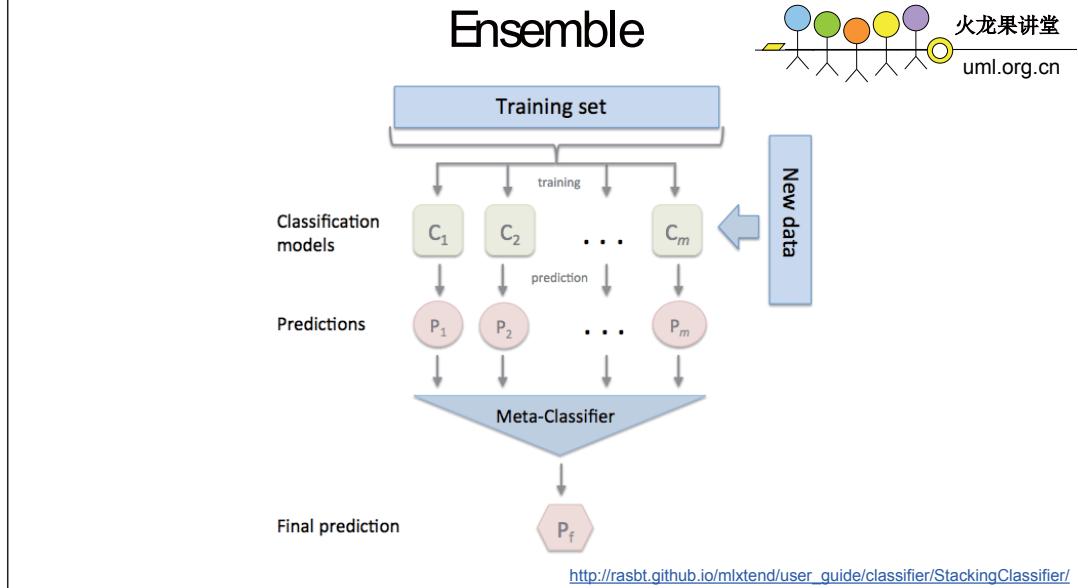
# L1 ImgIn shape=(?, 28, 28, 1)
W1 = tf.Variable(tf.random_normal([3, 3, 1, 32], stddev=0.01))
# Conv      -> (?, 28, 28, 32)
# Pool      -> (?, 14, 14, 32)
L1 = tf.nn.conv2d(X_img, W1, strides=[1, 1, 1, 1], padding='SAME')
L1 = tf.nn.relu(L1)
L1 = tf.nn.max_pool(L1, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding='SAME')
L1 = tf.nn.dropout(L1, keep_prob=self.keep_prob)
...
# L2 ImgIn shape=(?, 14, 14, 32)
W2 = tf.Variable(tf.random_normal([3, 3, 32, 64], stddev=0.01))

# Convolutional Layer #1
conv1 = tf.layers.conv2d(inputs=X_img, filters=32, kernel_size=[3,3], padding="SAME", activation=tf.nn.relu)
pool1 = tf.layers.max_pooling2d(inputs=conv1, pool_size=[2, 2], padding="SAME", strides=2)
dropout1 = tf.layers.dropout(inputs=pool1, rate=0.7, training=self.training)

# Convolutional Layer #2
conv2 = tf.layers.conv2d(inputs=dropout1, filters=64, kernel_size=[3,3], padding="SAME", activation=tf.nn.relu)
...
flat = tf.reshape(dropout3, [-1, 128 * 4 * 4])
dense4 = tf.layers.dense(inputs=flat, units=625, activation=tf.nn.relu)
dropout4 = tf.layers.dropout(inputs=dense4, rate=0.5, training=self.training)
...

```

[lab-11-4-mnist\\_cnn\\_layers.py](#)



# Ensemble training

```
class Model:
    def __init__(self, sess, name):
        self.sess = sess
        self.name = name
        self._build_net()

    def _build_net(self):
        with tf.variable_scope(self.name):
            ...

```

```
models = []
num_models = 7
for m in range(num_models):
    models.append(Model(sess, "model" + str(m)))

sess.run(tf.global_variables_initializer())
print('Learning Started!')

# train my model
for epoch in range(training_epochs):
    avg_cost_list = np.zeros(len(models))
    total_batch = int(mnist.train.num_examples / batch_size)
    for i in range(total_batch):
        batch_xs, batch_ys = mnist.train.next_batch(batch_size)

        # train each model
        for m_idx, m in enumerate(models):
            c, _ = m.train(batch_xs, batch_ys)
            avg_cost_list[m_idx] += c / total_batch

    print('Epoch:', '%04d' % (epoch + 1), 'cost =', avg_cost_list)

print('Learning Finished!')
lab-11-5-mnist_cnn_ensemble_layers.py
```



# Ensemble prediction

$C_1$

$C_2$

$C_m$

# Ensemble prediction



$C_1$

0	1	2	3	4	5	6	7	8	9
0.1	0.01	0.02	0.8	...	...	...	...	...	...

$C_2$

0.01	0.5	0.02	0.4	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...	...

$C_m$

0.01	0.01	0.1	0.7	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...	...

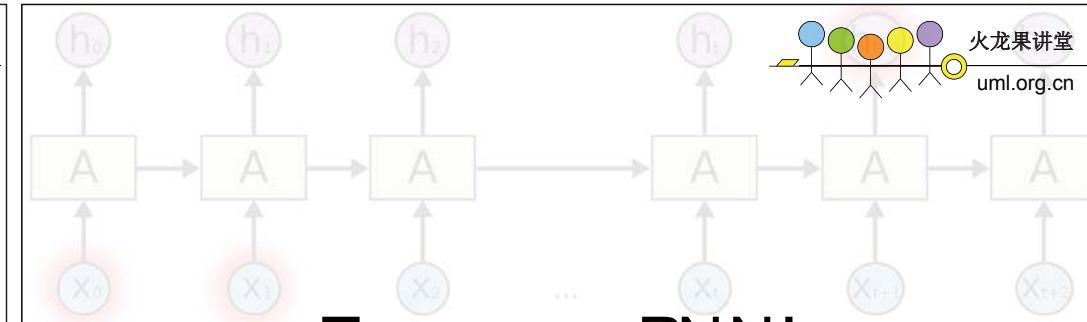
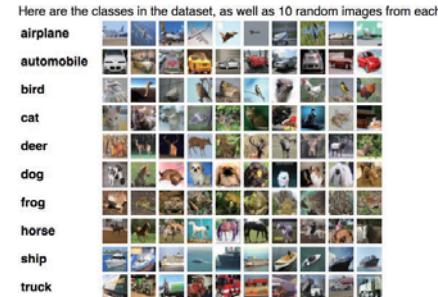
Sum

0.12	0.52	0.14	1.9	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...	...

↑ argmax

## Exercise

- Deep & Wide?
- CIFAR 10
- ImageNet



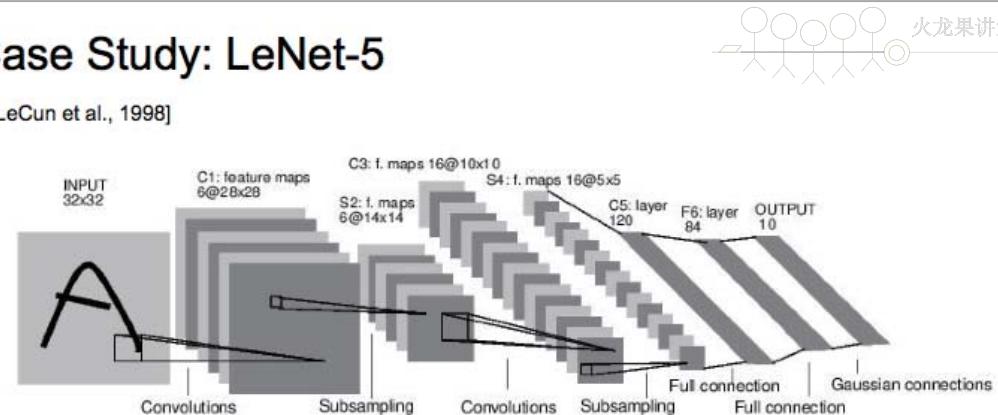
## Tomorrow RNN!

- Time series prediction (stock market)
- Character/sentence analysis
- Neural translation (sequence to sequence)

## Backup slides

## Case Study: LeNet-5

[LeCun et al., 1998]

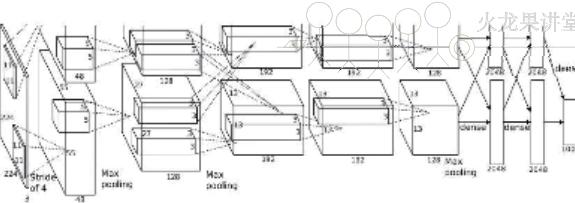


Conv filters were 5x5, applied at stride 1

Subsampling (Pooling) layers were 2x2 applied at stride 2  
i.e. architecture is [CONV-POOL-CONV-POOL-CONV-FC]

## Case Study: AlexNet

[Krizhevsky et al. 2012]



Input: 227x227x3 images

**First layer (CONV1):** 96 11x11 filters applied at stride 4

=>

Output volume [55x55x96]

Q: What is the total number of parameters in this layer?

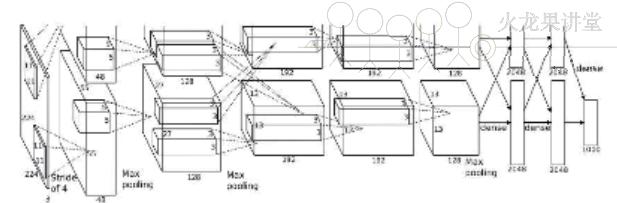
Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 62

27 Jan 2016

## Case Study: AlexNet

[Krizhevsky et al. 2012]



Input: 227x227x3 images

**First layer (CONV1):** 96 11x11 filters applied at stride 4

=>

Q: what is the output volume size? Hint:  $(227-11)/4+1 = 55$

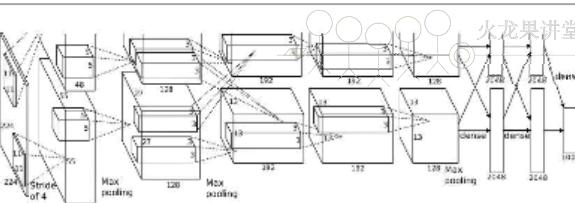
Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 61

27 Jan 2016

## Case Study: AlexNet

[Krizhevsky et al. 2012]



Input: 227x227x3 images

**First layer (CONV1):** 96 11x11 filters applied at stride 4

=>

Output volume [55x55x96]

Parameters:  $(11 \times 11 \times 3) \times 96 = 35K$

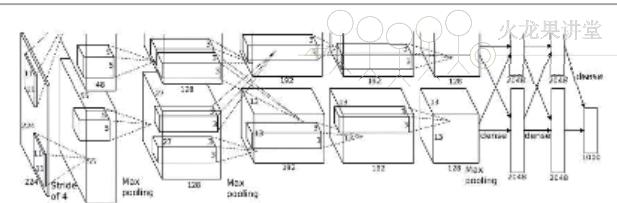
Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 63

27 Jan 2016

## Case Study: AlexNet

[Krizhevsky et al. 2012]



Input: 227x227x3 images

After CONV1: 55x55x96

**Second layer (POOL1):** 3x3 filters applied at stride 2

Output volume: 27x27x96

Parameters: 0!

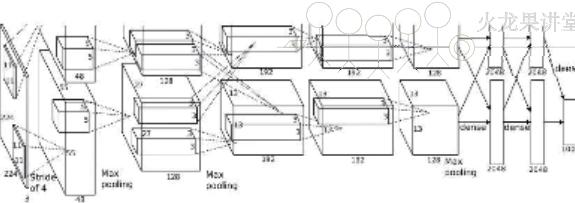
Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 66

27 Jan 2016

## Case Study: AlexNet

[Krizhevsky et al. 2012]



Input: 227x227x3 images

After CONV1: 55x55x96

After POOL1: 27x27x96

...

Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 67

27 Jan 2016

## Case Study: AlexNet

[Krizhevsky et al. 2012]

Full (simplified) AlexNet architecture:

[227x227x3] INPUT

[55x55x96] CONV1: 96 11x11 filters at stride 4, pad 0

[27x27x96] MAX POOL1: 3x3 filters at stride 2

[27x27x96] NORM1: Normalization layer

[27x27x256] CONV2: 256 5x5 filters at stride 1, pad 1

[13x13x256] MAX POOL2: 3x3 filters at stride 2

[13x13x256] NORM2: Normalization layer

[13x13x384] CONV3: 384 3x3 filters at stride 1, pad 1

[13x13x384] CONV4: 384 3x3 filters at stride 1, pad 1

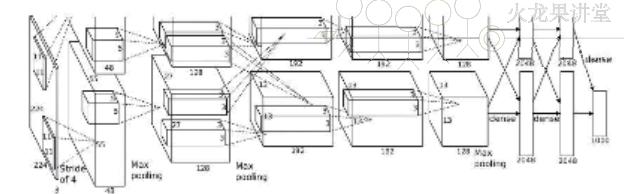
[13x13x256] CONV5: 256 3x3 filters at stride 1, pad 1

[6x6x256] MAX POOL3: 3x3 filters at stride 2

[4096] FC6: 4096 neurons

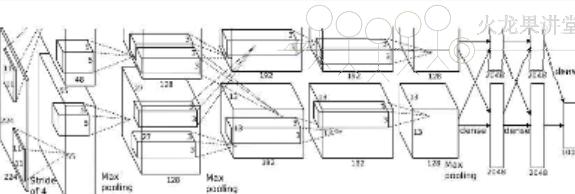
[4096] FC7: 4096 neurons

[1000] FC8: 1000 neurons (class scores)



## Case Study: AlexNet

[Krizhevsky et al. 2012]



Full (simplified) AlexNet architecture:

[227x227x3] INPUT

[55x55x96] CONV1: 96 11x11 filters at stride 4, pad 0

[27x27x96] MAX POOL1: 3x3 filters at stride 2

[27x27x96] NORM1: Normalization layer

[27x27x256] CONV2: 256 5x5 filters at stride 1, pad 2

[13x13x256] MAX POOL2: 3x3 filters at stride 2

[13x13x256] NORM2: Normalization layer

[13x13x384] CONV3: 384 3x3 filters at stride 1, pad 1

[13x13x384] CONV4: 384 3x3 filters at stride 1, pad 1

[13x13x256] CONV5: 256 3x3 filters at stride 1, pad 1

[6x6x256] MAX POOL3: 3x3 filters at stride 2

[4096] FC6: 4096 neurons

[4096] FC7: 4096 neurons

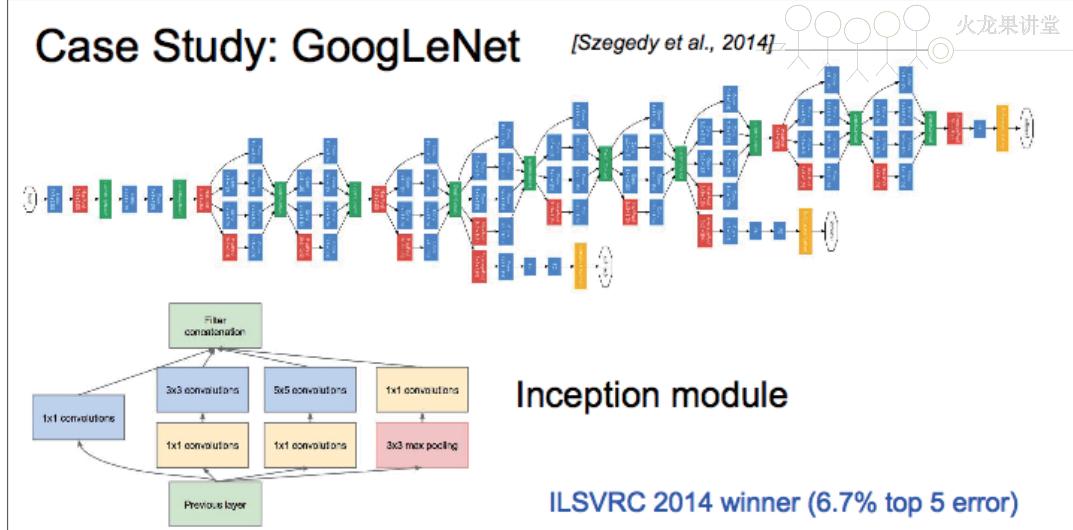
[1000] FC8: 1000 neurons (class scores)

### Details/Retrospectives:

- first use of ReLU
- used Norm layers (not common anymore)
- heavy data augmentation
- dropout 0.5
- batch size 128
- SGD Momentum 0.9
- Learning rate 1e-2, reduced by 10 manually when val accuracy plateaus
- L2 weight decay 5e-4
- 7 CNN ensemble: 18.2% -> 15.4%

## Case Study: GoogLeNet

[Szegedy et al., 2014]



### Inception module

ILSVRC 2014 winner (6.7% top 5 error)

Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 69

27 Jan 2016

Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 75

27 Jan 2016

## Case Study: ResNet

[He et al., 2015]

ILSVRC 2015 winner (3.6% top 5 error)



MSRA @ ILSVRC & COCO 2015 Competitions

- 1st places in all five main tracks

- ImageNet Classification: "Ultra-deep" (quote Yann) 152-layer nets
- ImageNet Detection: 16% better than 2nd
- ImageNet Localization: 27% better than 2nd
- COCO Detection: 11% better than 2nd
- COCO Segmentation: 12% better than 2nd

\*Improvements are relative numbers



Kaiming He, Xiangyu Zhang, Shaoqing Ren, & Jian Sun. "Deep Residual Learning for Image Recognition". arXiv 2015.

Slide from Kaiming He's recent presentation <https://www.youtube.com/watch?v=1PGLj-uKT1w>

Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 77

27 Jan 2016

## Case Study: ResNet

[He et al., 2015]

ILSVRC 2015 winner (3.6% top 5 error)

Revolution of Depth

AlexNet, 8 layers  
(ILSVRC 2012)

VGG, 19 layers  
(ILSVRC 2014)

ResNet, 152 layers  
(ILSVRC 2015)



Kaiming He, Xiangyu Zhang, Shaoqing Ren, & Jian Sun. "Deep Residual Learning for Image Recognition". arXiv 2015.

(slide from Kaiming He's recent presentation)

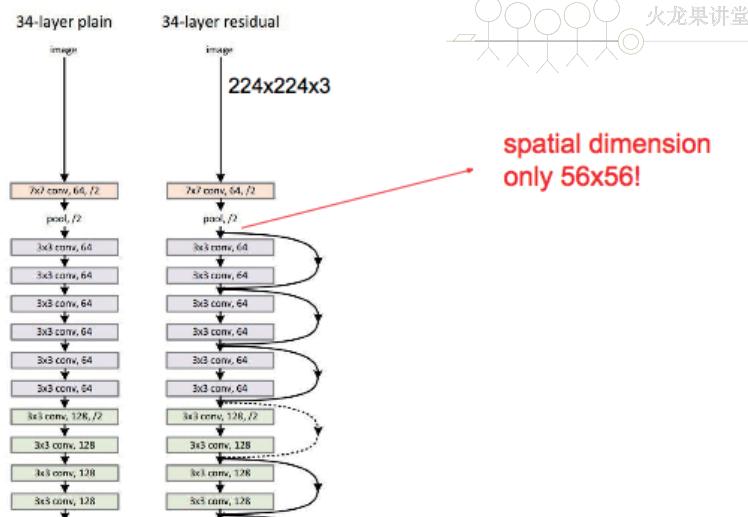
Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 80

27 Jan 2016

## Case Study: ResNet

[He et al., 2015]



Fei-Fei Li & Andrej Karpathy & Justin Johnson

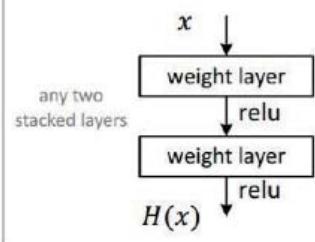
Lecture 7 - 81

27 Jan 2016

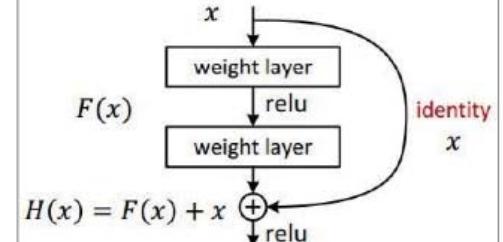
## Case Study: ResNet

[He et al., 2015]

- Plain net



- Residual net



Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 82

27 Jan 2016

## Case Study: ResNet

[He et al., 2015]



火龙果讲堂

- Batch Normalization after every CONV layer
- Xavier/2 initialization from He et al.
- SGD + Momentum (0.9)
- Learning rate: 0.1, divided by 10 when validation error plateaus
- Mini-batch size 256
- Weight decay of 1e-5
- No dropout used

Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 83

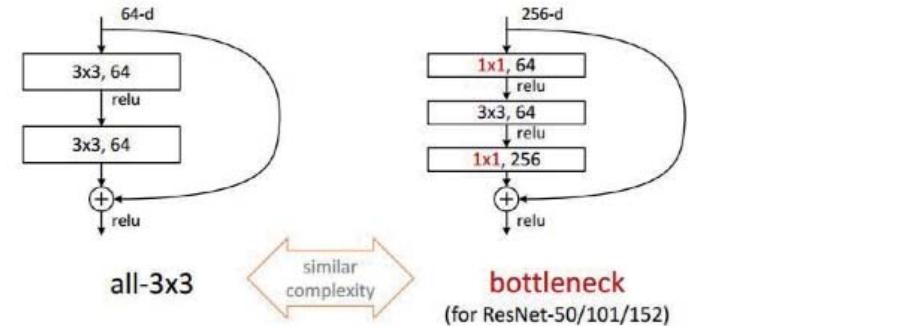
27 Jan 2016

## Case Study: ResNet

[He et al., 2015]



火龙果讲堂



Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 84

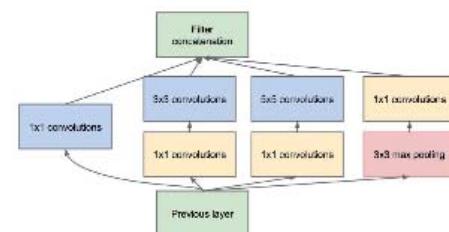
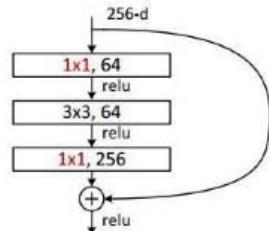
27 Jan 2016

## Case Study: ResNet

[He et al., 2015]



火龙果讲堂



(this trick is also used in GoogLeNet)

Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 85

27 Jan 2016

## Case Study: ResNet

[He et al., 2015]



火龙果讲堂

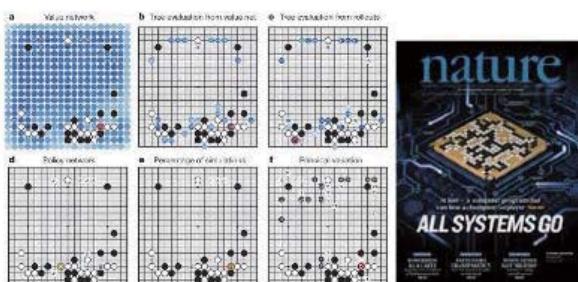
layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	112×112					
				7×7, 64, stride 2		
					3×3 max pool, stride 2	
conv2_x	56×56	$\left[ \begin{matrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{matrix} \right] \times 2$	$\left[ \begin{matrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{matrix} \right] \times 3$	$\left[ \begin{matrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{matrix} \right] \times 3$	$\left[ \begin{matrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{matrix} \right] \times 3$	$\left[ \begin{matrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{matrix} \right] \times 3$
conv3_x	28×28	$\left[ \begin{matrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{matrix} \right] \times 2$	$\left[ \begin{matrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{matrix} \right] \times 4$	$\left[ \begin{matrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{matrix} \right] \times 4$	$\left[ \begin{matrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{matrix} \right] \times 4$	$\left[ \begin{matrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{matrix} \right] \times 8$
conv4_x	14×14	$\left[ \begin{matrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{matrix} \right] \times 2$	$\left[ \begin{matrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{matrix} \right] \times 6$	$\left[ \begin{matrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{matrix} \right] \times 6$	$\left[ \begin{matrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{matrix} \right] \times 23$	$\left[ \begin{matrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{matrix} \right] \times 36$
conv5_x	7×7	$\left[ \begin{matrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{matrix} \right] \times 2$	$\left[ \begin{matrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{matrix} \right] \times 3$	$\left[ \begin{matrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{matrix} \right] \times 3$	$\left[ \begin{matrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{matrix} \right] \times 3$	$\left[ \begin{matrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{matrix} \right] \times 3$
	1×1				average pool, 1000-d fc, softmax	
		FLOPs	$1.8 \times 10^9$	$3.6 \times 10^9$	$3.8 \times 10^9$	$7.6 \times 10^9$
						$11.3 \times 10^9$

Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 86

27 Jan 2016

## Case Study Bonus: DeepMind's AlphaGo



Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 87

27 Jan 2016

The input to the policy network is a  $19 \times 19 \times 48$  image stack consisting of 48 feature planes. The first hidden layer zero pads the input into a  $23 \times 23$  image, then convolves  $k$  filters of kernel size  $5 \times 5$  with stride 1 with the input image and applies a rectifier nonlinearity. Each of the subsequent hidden layers 2 to 12 zero pads the respective previous hidden layer into a  $21 \times 21$  image, then convolves  $k$  filters of kernel size  $3 \times 3$  with stride 1, again followed by a rectifier nonlinearity. The final layer convolves 1 filter of kernel size  $1 \times 1$  with stride 1, with a different bias for each position, and applies a softmax function. The match version of AlphaGo used  $k = 192$  filters; Fig. 2b and Extended Data Table 3 additionally show the results of training with  $k = 128, 256$  and  $384$  filters.

### policy network:

[ $19 \times 19 \times 48$ ] Input

CONV1: 192  $5 \times 5$  filters , stride 1, pad 2 => [ $19 \times 19 \times 192$ ]

CONV2..12: 192  $3 \times 3$  filters, stride 1, pad 1 => [ $19 \times 19 \times 192$ ]

CONV: 1  $1 \times 1$  filter, stride 1, pad 0 => [ $19 \times 19$ ] (*probability map of promising moves*)

Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 88

27 Jan 2016



## RNN

循环神经网络  
深度学习网络架构



## Gifts from Google and Line

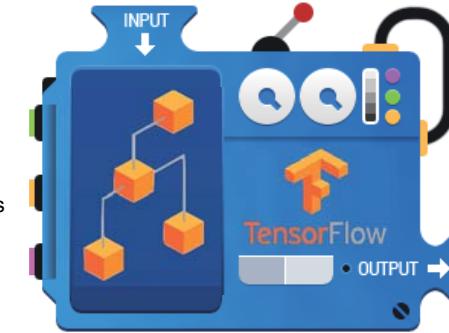


# Feedback



## TensorFlow Mechanics

- 1 Build graph using TensorFlow operations
- 2 feed data and run graph (operation)  
`sess.run (op, feed_dict={x: x_data})`



WWW.MATHWAREHOUSE.COM

## Machine Learning Basics



- Linear Regression
- Logistic Regression (Binary classification)
- Softmax Classification
- Neural Networks
- CNN

## Linear Regression

- 1  $H(x) = Wx + b$

```
W = tf.Variable(tf.random_normal([1]), name='weight')
b = tf.Variable(tf.random_normal([1]), name='bias')
# Our hypothesis XW+b
hypothesis = X * W + b
```

- 2  $\text{cost}(W, b) = \frac{1}{m} \sum_{i=1}^m (H(x^{(i)}) - y^{(i)})^2$

```
# cost/loss function
cost = tf.reduce_mean(tf.square(hypothesis - y_train))
```

- 3 GradientDescent  $W := W - \alpha \frac{\partial}{\partial W} \text{cost}(W)$

```
# Minimize
train = tf.train.GradientDescentOptimizer(learning_rate=0.01).minimize(cost)
```

# Logistic Regression

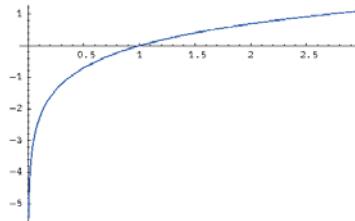


火龙果讲堂  
uml.org.cn

## 1 Model? (Hypothesis?)

$$H(X) = \text{sigmoid}(XW) = \frac{1}{1 + e^{-XW}}$$

```
hypothesis = tf.sigmoid(tf.matmul(X, W) + b)
```



## 2 Cost?

$$\text{cost}(W) = -\frac{1}{m} \sum y \log(H(x)) + (1 - y) \log(1 - H(x))$$

# cost/loss function

```
cost = -tf.reduce_mean(Y * tf.log(hypothesis) + (1 - Y) * tf.log(1 - hypothesis))
```

## 3 GradientDescent $W := W - \alpha \frac{\partial}{\partial W} \text{cost}(W)$

```
train = tf.train.GradientDescentOptimizer(learning_rate=0.01).minimize(cost)
```

# Softmax Classifier



## 1 Model? (Hypothesis?)

```
logits = tf.matmul(X, W) + b
hypothesis = tf.nn.softmax(logits)
```

## 2 Cost?

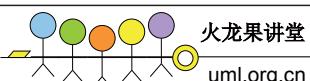
# Cross entropy cost/loss

```
cost = tf.reduce_mean(-tf.reduce_sum(Y * tf.log(hypothesis), axis=1))
```

## 3 GradientDescent $W := W - \alpha \frac{\partial}{\partial W} \text{cost}(W)$

```
train = tf.train.GradientDescentOptimizer(learning_rate=0.01).minimize(cost)
```

# Softmax Classifier



## 1 Model? (Hypothesis?)

```
logits = tf.matmul(X, W) + b
hypothesis = tf.nn.softmax(logits)
```

## 2 Cost?

# Cross entropy cost/loss

```
cost = tf.reduce_mean(-tf.reduce_sum(Y * tf.log(hypothesis), axis=1))
```

```
cost = tf.reduce_mean(
    tf.nn.softmax_cross_entropy_with_logits(logits=logits, labels=Y))
```

## 3 GradientDescent $W := W - \alpha \frac{\partial}{\partial W} \text{cost}(W)$

```
train = tf.train.GradientDescentOptimizer(learning_rate=0.01).minimize(cost)
```

# (Deep) Neural Nets



# input placeholders

```
X = tf.placeholder(tf.float32, [None, 784])
Y = tf.placeholder(tf.float32, [None, 10])
```

# weights & bias for nn Layers

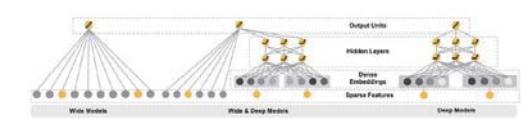
```
W1 = tf.Variable(tf.random_normal([784, 256]))
b1 = tf.Variable(tf.random_normal([256]))
L1 = tf.nn.relu(tf.matmul(X, W1) + b1)
```

```
W2 = tf.Variable(tf.random_normal([256, 256]))
b2 = tf.Variable(tf.random_normal([256]))
L2 = tf.nn.relu(tf.matmul(L1, W2) + b2)
```

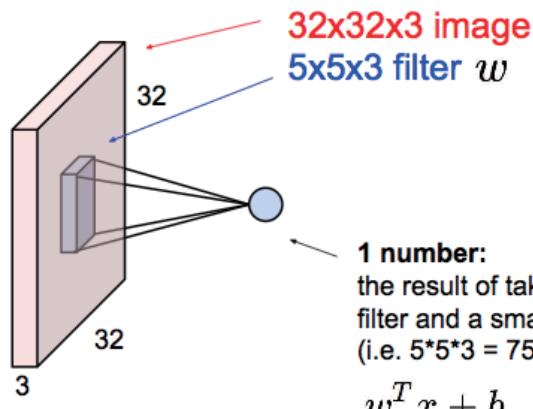
```
W3 = tf.Variable(tf.random_normal([256, 10]))
b3 = tf.Variable(tf.random_normal([10]))
hypothesis = tf.matmul(L2, W3) + b3
```

# define cost/loss & optimizer

```
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(
    logits=hypothesis, labels=Y))
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost)
```



## Convolution Layer



Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 13

27 Jan 2016

## Simple CNN

```

# img 28x28x1 (black/white), Input Layer
X_img = tf.reshape(self.X, [-1, 28, 28, 1])
self.Y = tf.placeholder(tf.float32, [None, 10])

# Convolutional Layer and pooling layer #1
conv1 = tf.layers.conv2d(inputs=X_img, filters=32, kernel_size=[3, 3], padding="SAME", activation=tf.nn.relu)
pool1 = tf.layers.max_pooling2d(inputs=conv1, pool_size=[2, 2], padding="SAME", strides=2)
dropout1 = tf.layers.dropout(inputs=pool1, rate=0.7, training=self.training)

# Convolutional Layer and Pooling Layer #2
conv2 = tf.layers.conv2d(inputs=dropout1, filters=64, kernel_size=[3, 3], padding="SAME", activation=tf.nn.relu)
pool2 = tf.layers.max_pooling2d(inputs=conv2, pool_size=[2, 2], padding="SAME", strides=2)
dropout2 = tf.layers.dropout(inputs=pool2, rate=0.7, training=self.training)

# Dense Layer with Relu
flat = tf.reshape(dropout2, [-1, 64 * 7 * 7])
dense4 = tf.layers.dense(inputs=flat, units=625, activation=tf.nn.relu)
dropout4 = tf.layers.dropout(inputs=dense4, rate=0.5, training=self.training)

# Logits (no activation) Layer: L5 Final FC 625 inputs -> 10 outputs
self.logits = tf.layers.dense(inputs=dropout4, units=10)

# define cost/loss & optimizer
self.cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=self.logits, labels=self.Y))
self.optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(self.cost)

```

## NN tips

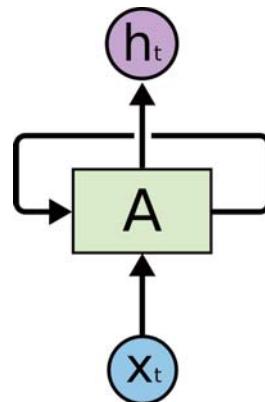
- Initializing weights
- Activation functions
- Regularization
- Optimizers

## Sequence data

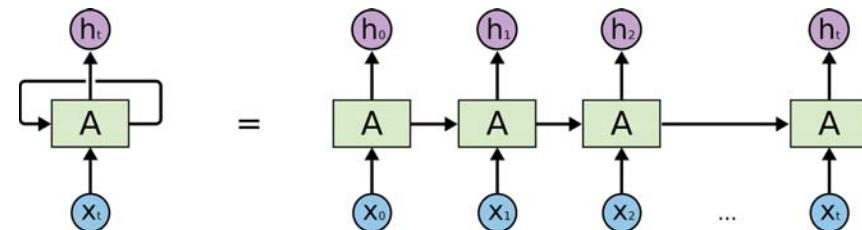
- Hard to understand the context from single word
- We understand based on the previous words + this word. (time series)
- NN/CNN cannot do this

## Sequence data

- Hard to understand the context from single word
- We understand based on the previous words + this word. (time series)
- NN/CNN cannot do this

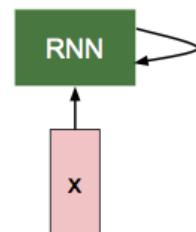


<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

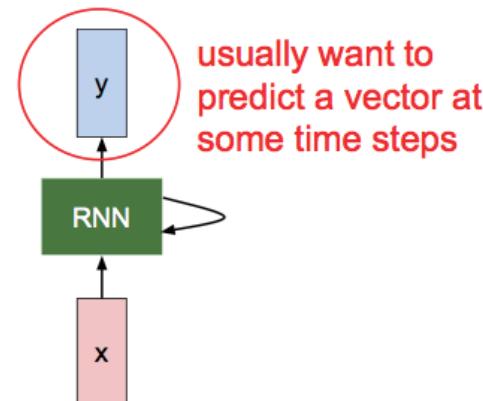


<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

## Recurrent Neural Network



## Recurrent Neural Network

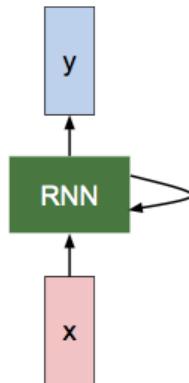


# Recurrent Neural Network

We can process a sequence of vectors  $\mathbf{x}$  by applying a recurrence formula at every time step:

$$h_t = f_W(h_{t-1}, x_t)$$

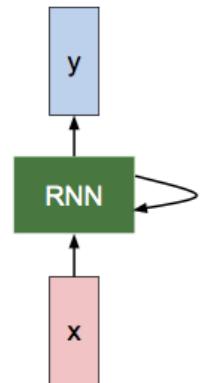
new state      /      old state      input vector at some time step  
 some function with parameters W



# Recurrent Neural Network

We can process a sequence of vectors  $\mathbf{x}$  by applying a recurrence formula at every time step:

$$h_t = f_W(h_{t-1}, x_t)$$



## (Vanilla) Recurrent Neural Network

The state consists of a single “hidden” vector  $\mathbf{h}$ :

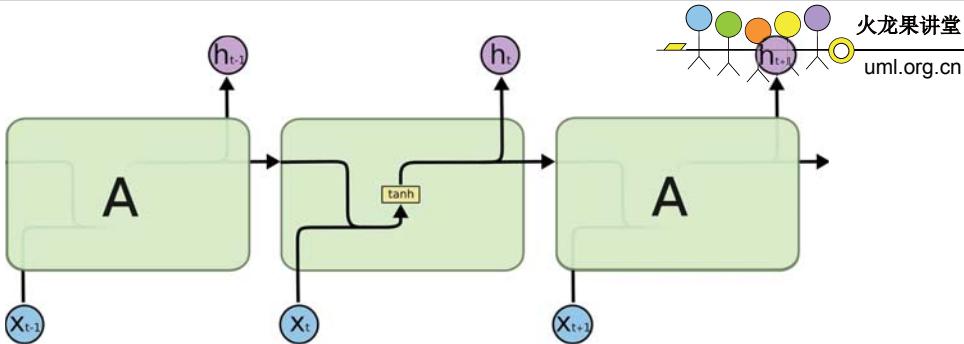
$$h_t = f_W(h_{t-1}, x_t)$$

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

$$y_t = W_{hy}h_t$$



Notice: the same function and the same set of parameters are used at every time step.



Given list of word **vectors**:  $x_1, \dots, x_{t-1}, x_t, x_{t+1}, \dots, x_T$

$$\text{At a single time step: } h_t = \sigma(W^{(hh)}h_{t-1} + W^{(hx)}x_{[t]})$$

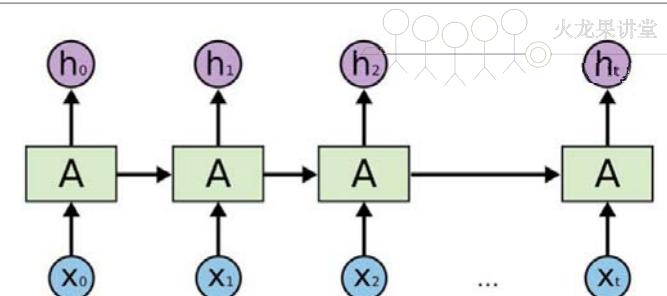
$$\hat{y}_t = \text{softmax}(W^{(S)}h_t)$$

$$P(x_{t+1} = v_j \mid x_t, \dots, x_1) = \hat{y}_{t,j}$$

## Character-level language model example

Vocabulary:  
[h,e,l,o]

Example training sequence:  
“hello”



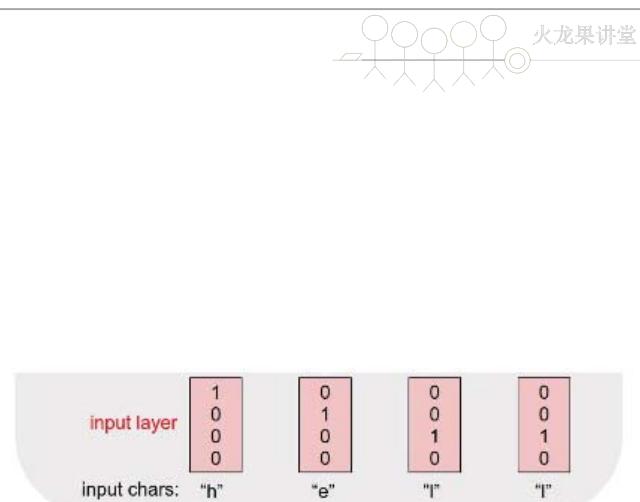
Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 10 - 18 8 Feb 2016

## Character-level language model example

Vocabulary:  
[h,e,l,o]

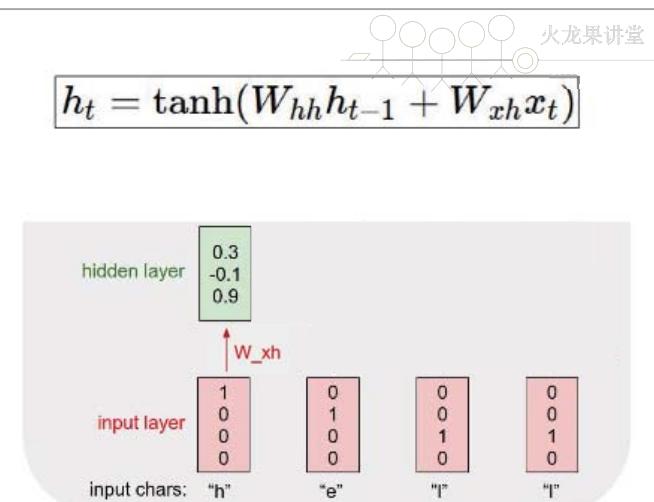
Example training sequence:  
“hello”



## Character-level language model example

Vocabulary:  
[h,e,l,o]

Example training sequence:  
“hello”



Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 10 - 19 8 Feb 2016

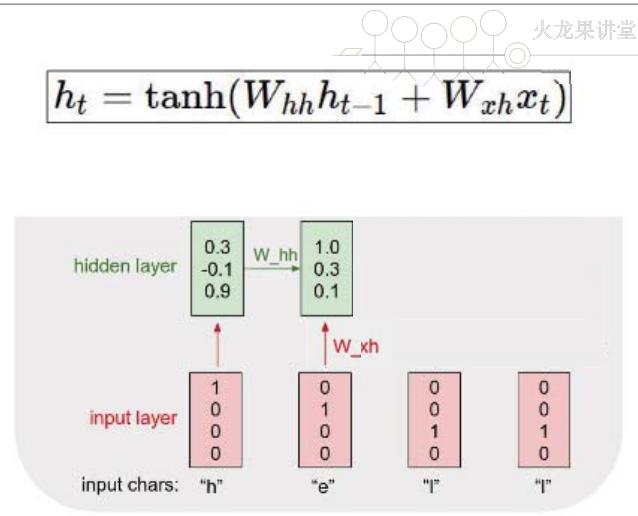
Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 10 - 20 8 Feb 2016

## Character-level language model example

Vocabulary:  
[h,e,l,o]

Example training  
sequence:  
“hello”



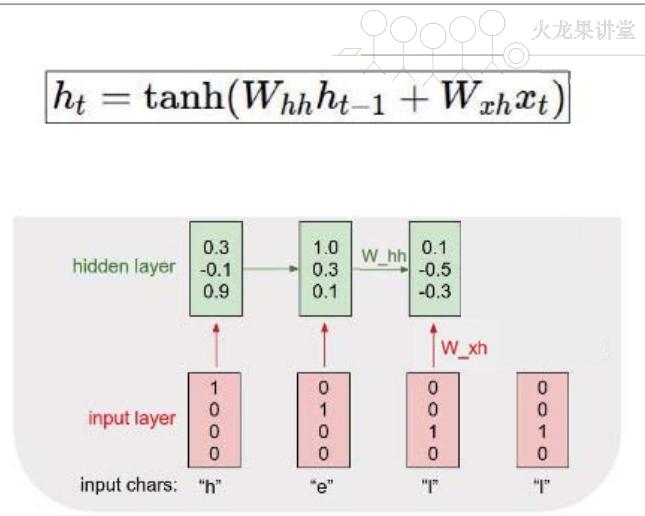
Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 10 - 20 8 Feb 2016

## Character-level language model example

Vocabulary:  
[h,e,l,o]

Example training  
sequence:  
“hello”



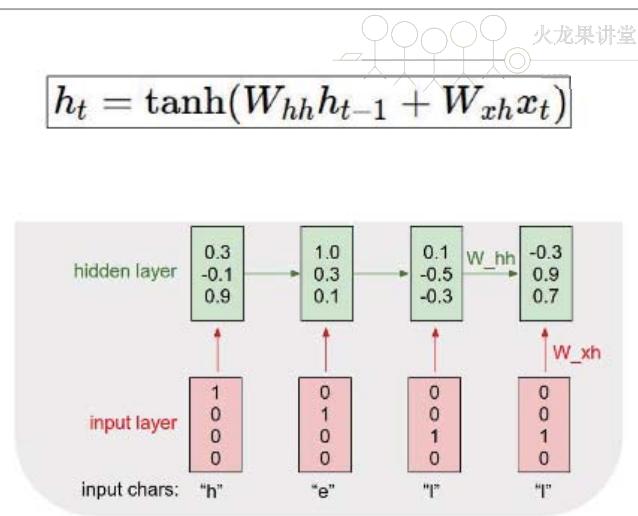
Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 10 - 20 8 Feb 2016

## Character-level language model example

Vocabulary:  
[h,e,l,o]

Example training  
sequence:  
“hello”



Fei-Fei Li & Andrej Karpathy & Justin Johnson

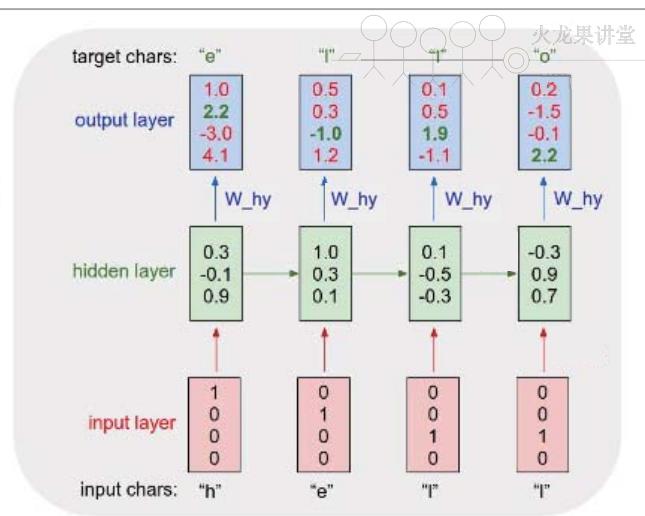
Lecture 10 - 20 8 Feb 2016

## Character-level language model example

$$y_t = W_{hy}h_t$$

Vocabulary:  
[h,e,l,o]

Example training  
sequence:  
“hello”



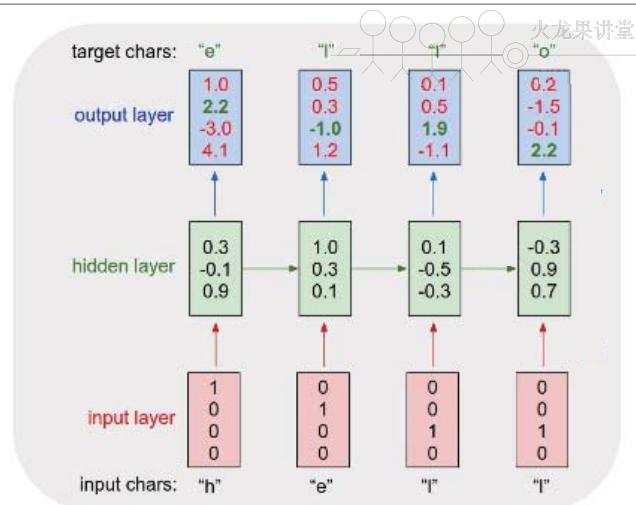
Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 10 - 21 8 Feb 2016

## Character-level language model example

Vocabulary:  
[h,e,l,o]

Example training  
sequence:  
“hello”



Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 10 - 21 8 Feb 2016

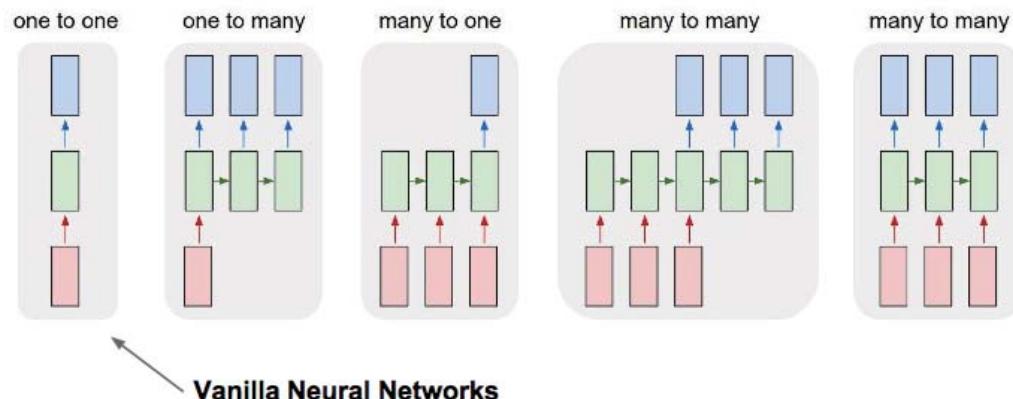
## RNN applications

[https://github.com/TensorFlowKR/awesome\\_tensorflow\\_implementations](https://github.com/TensorFlowKR/awesome_tensorflow_implementations)

- Language Modeling
- Speech Recognition
- Machine Translation
- Conversation Modeling/Question Answering
- Image/Video Captioning
- Image/Music/Dance Generation

<https://kjh0612.github.io/awesome-rnn/>

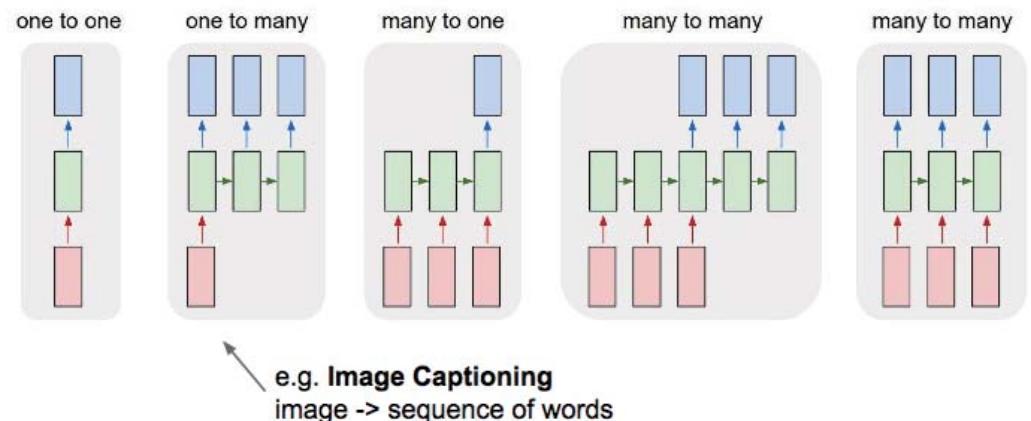
## Recurrent Networks offer a lot of flexibility:



Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 10 - 6 8 Feb 2016

## Recurrent Networks offer a lot of flexibility:

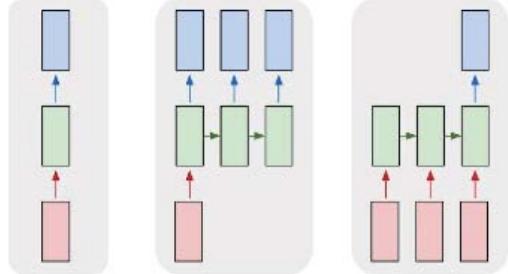


Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 10 - 7 8 Feb 2016

## Recurrent Networks offer a lot of flexibility:

one to one      one to many      many to one



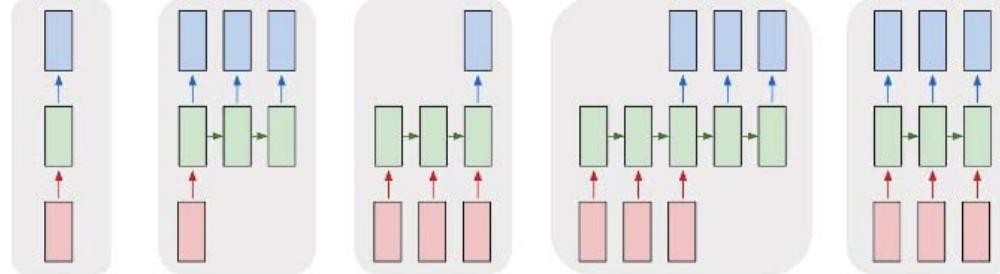
e.g. **Sentiment Classification**  
sequence of words -> sentiment

Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 10 - 8 8 Feb 2016

## Recurrent Networks offer a lot of flexibility:

one to one      one to many      many to one      many to many      many to many



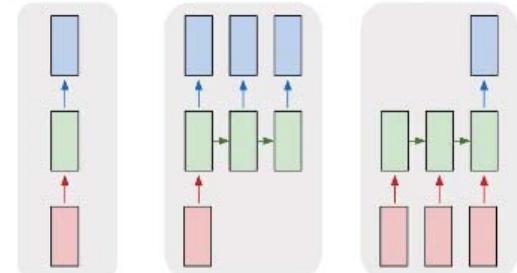
e.g. **Machine Translation**  
seq of words -> seq of words

Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 10 - 9 8 Feb 2016

## Recurrent Networks offer a lot of flexibility:

one to one      one to many      many to one

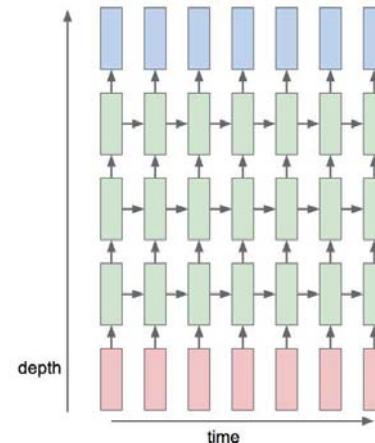


e.g. **Video classification on frame level**

Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 10 - 10 8 Feb 2016

## Multi-Layer RNN



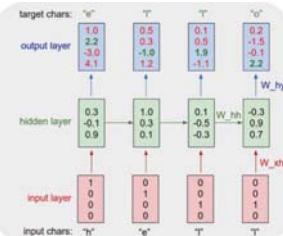
# Training RNNs is challenging

- Several advanced models
  - Long Short Term Memory (LSTM)
  - GRU by Cho et al. 2014



火龙果讲堂  
uml.org.cn

## char-rnn



### Shakespeare

It looks like we can learn to spell English words. But how about if there is more structure and style in the data? To examine this I downloaded all the works of Shakespeare and concatenated them into a single (4.4MB) file. We can now afford to train a larger network, in this case lets try a 3-layer RNN with 512 hidden nodes on each layer. After we train the network for a few hours we obtain samples such as:

**PANDARUS:**  
Alas, I think he shall be come approached and the day  
When little strain would be attain'd into being never fed,  
And who is but a chain and subjects of his death,  
I should not sleep.

**Second Senator:**  
They are away this miseries, produced upon my soul,  
Breaking and strongly should be buried, when I perish  
The earth and thoughts of many states.

**DUKE VINCENTIO:**  
Well, your wit is in the care of side and that.

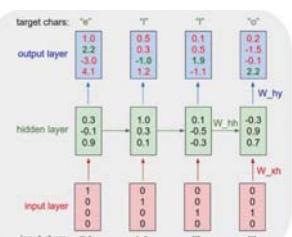
**Second Lord:**  
They would be ruled after this chamber, and  
my fair nues begun out of the fact, to be conveyed,  
Whose noble souls I'll have the heart of the wars.

**Clown:**  
Come, sir, I will make did behold your worship.

**VIOLA:**  
I'll drink it.



Linux Source Code  
I wanted to push structured data to its limit, so for the final challenge I decided to use code. In particular, I took all the source and header files found in the Linux repo on Github, concatenated all of them in a single giant file (474MB of C code) I was originally going to train only on the kernel but that by itself is only ~16MB). Then I trained several as-large-as-fits-on-my-GPU 3-layer LSTMs over a period of a few days. These models have about 10 million parameters, which is still on the lower end for RNN models. The results are superfun:



```
/*
 * Increment the size file of the new incorrect UI_FILTER group information
 * of the size generatively.
 */
static int indicate_policy(void)
{
    int error;
    if (fd == NASH_EPT) {
        /* The kernel blank will coold it to user space. */
        if (ss->segment < mem_total)
            unblock_graph_and_set_blocked();
        else
            ret = 1;
        goto bail;
    }
    seqaddr = in_BB(in.addr);
    selector = seq / 16;
    setup_works = true;
    for (i = 0; i < blocks; i++) {
        seq = buf[i++];
        bpf = bd->bd.next + i * search;
        if (fd) {
            current = blocked;
        }
    }
    rv->name = "Getjbbreg";
    bprm_self_clear(&rv->version);
    reg->new = blocks[(BPF_STATS << info->historidac)] | PPFMS_CLOGBATHINC_SECONDS << 13;
    return segtable;
}
```

## Google Translate



Translate

English Korean Spanish Detect language

If you want to build a ship, don't drum up people together to collect wood and don't assign them tasks and work, but rather teach them to long for the endless immensity of the sea.

如果要建造一艘船，不要一起鼓勵人們收集木材，  
不要分配任務和工作，而應該教他們漫長的海洋無  
限遠。

Rúguō yào jiànzhào yí sōu chuán, bùyào yīqǐ gǔlì rénmen shōují mùcái, bùyào fēnpèi rénwù hé  
gōngzuò, ér yǐnggāi jiāo tāmén mǎncháng dí háiyáng wúxiān yuán.

## Multilingual Machine Translation with RNN

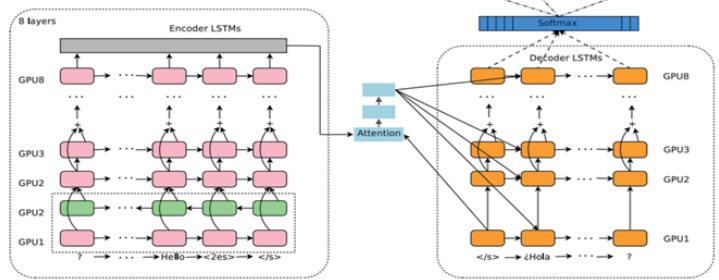


Figure 1: The model architecture of the Multilingual GNMT system. In addition to what is described in [24], our input has an artificial token to indicate the required target language. In this example, the token "<2es>" indicates that the target sentence is in Spanish, and the source sentence is reversed as a processing step. For most of our experiments we also used direct connections between the encoder and decoder although we later found out that the effect of these connections is negligible (however, once you train with those they have to be present for inference as well). The rest of the model architecture is the same as in [24].

Source: <https://arxiv.org/pdf/1611.04558.pdf>

Google's Multilingual Neural Machine Translation System: Enabling Zero-Shot Translation

## Question Answering with RNN

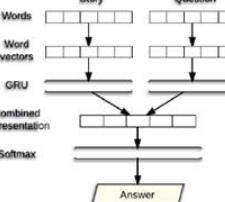


Figure 1: GRU baseline

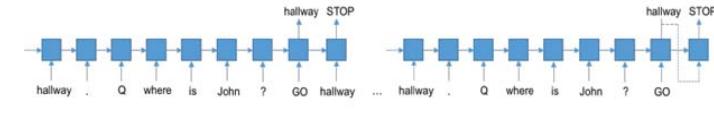
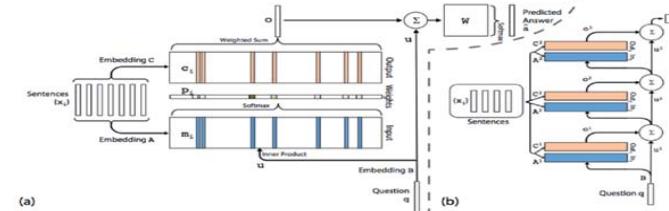


Figure 2: Sequence-to-sequence baseline: training (left); validation / testing (right)



Source: <https://cs224d.stanford.edu/reports/StrohMathur.pdf>

Question Answering Using Deep Learning

## Handwriting generation

recurrent neural network handwriting generation demo

Type a message into the text box, and the network will try to write it out longhand (this paper explains how it works, source code is available [here](#)). Be patient, it can take a while!

Text — up to 100 characters, lower case letters work best

Style — either let the network choose a writing style at random or prime it with a real sequence to make it mimic that writer's style.

- Take the bath, very short they are
- He discussed his idea
- prison welfare officer complement
- She looked closely at the
- cat. His fur is very soft for
- random style

Bias — increasing the bias makes the samples more legible but less diverse. Using a high bias and a printing sequence makes the network write in a neater version of the original style.

Samples

x write

Hello  
Hello  
Hello

<https://www.cs.toronto.edu/~graves/handwriting.html>

## Sketch-RNN

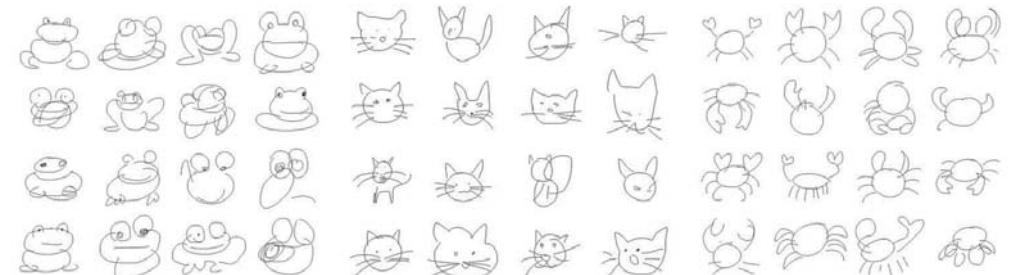
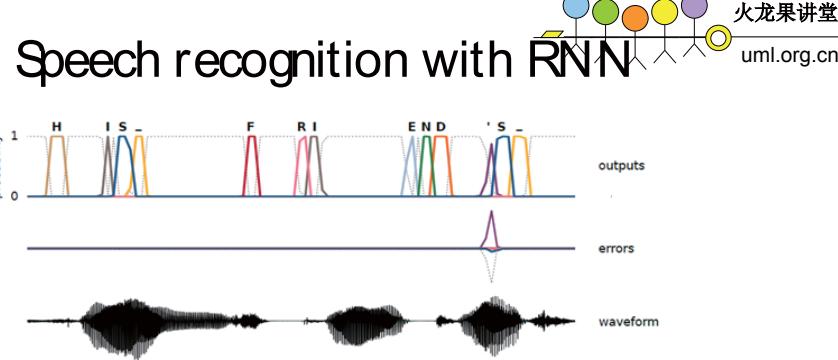


Figure 20: Unconditional generated sketches of frogs, cats, and crabs at  $\tau = 0.8$ .

<https://arxiv.org/pdf/1704.03477v1.pdf>

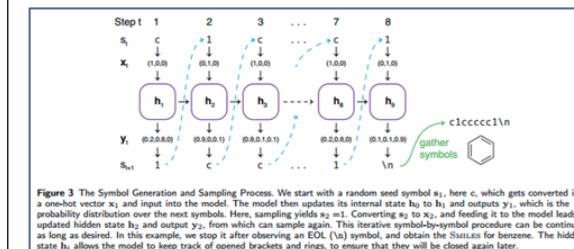


**Figure 4. Network outputs.** The figure shows the frame-level character probabilities emitted by the CTC layer (different colour for each character, dotted grey line for ‘blanks’), along with the corresponding training errors, while processing an utterance. The target transcription was ‘HIS.FRIENDS\_’, where the underscores are end-of-word markers. The network was trained with WER loss, which tends to give very sharp output decisions, and hence sparse error signals (if an output probability is 1, nothing else can be sampled, so the gradient is 0 even if the output is wrong). In this case the only gradient comes from the extraneous apostrophe before the ‘S’. Note that the characters in common sequences such as ‘IS’, ‘RI’ and ‘END’ are emitted very close together, suggesting that the network learns them as single sounds.

Source: <http://proceedings.mlr.press/v32/graves14.pdf>

Towards End-to-End Speech Recognition with Recurrent Neural Networks

## Generate Molecular structures with RNN



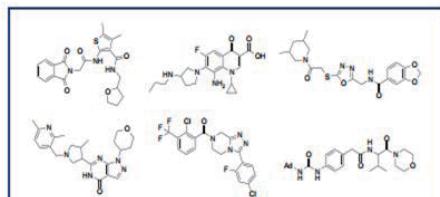
**Figure 3** The Symbol Generation and Sampling Process. We start with a random seed symbol  $x_1$ , here ‘c’, which gets converted into a one-hot vector  $x_1$  and input into the model. The model then updates its internal state  $h_1$  to  $h_2$  and outputs  $y_1$ , which is the probability distribution over the next symbols. Here, sampling yields  $x_2 = \text{c}$ . Converting  $x_2$  to  $x_3$ , and feeding it to the model leads to  $x_3 = \text{c}$  again. This process continues until the model reaches an EOS symbol ‘l’. The generated sequence can be extended as long as desired. In this example, we stop it after observing an EOS (‘l’) symbol, and obtain the SMILES for benzene. The hidden state  $h_i$  allows the model to keep track of opened brackets and rings, to ensure that they will be closed again later.

**Table 1** Molecules sampled during training.

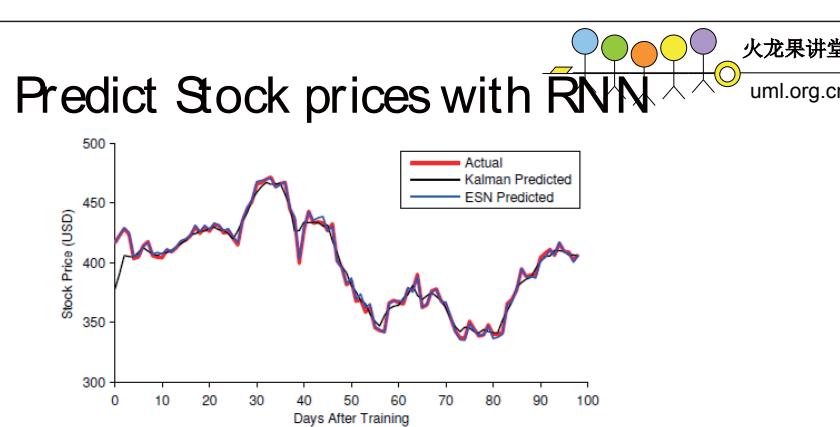
Batch	Generated Example	valid
0	Oc.BK5i%ur+7oAfc7L3T#F885e=nn)CS6RCTAR((0Wcp1Capb)	no
1000	OF=CCC20CCCC(C2)C1CN2CCCCCCCCCCCCCCCCCCCC	no
2000	O=C(N)C(=O)N(c1ccc1C)c2cccc2C0	yes
3000	O=C1c=2N(c3cc(cc3OC2CCC1)CCc4cn(c5c(Cl)cccc54)C)C	yes

Source: <https://arxiv.org/pdf/1701.01329.pdf>

Generating Focussed Molecule Libraries for Drug Discovery with Recurrent Neural Networks



**Figure 4** A few randomly selected, generated molecules. Ad = Adamantyl



**Figure 5:** Google stock price prediction for ESN and Kalman filter. ESN predicts rapid changes in stock price whereas Kalman filter tends to smooth rapid changes. Test error,  $\frac{(y - y_{target})^2}{\sigma_{y_{target}}^2}$ , for the ESN is 0.0027. In contrast, test error for the Kalman filter is 0.2135.

Source: <http://cs229.stanford.edu/proj2012/BernalFokPidiaparthi-FinancialMarketTimeSeriesPredictionwithRecurrentNeural.pdf>  
Financial Market Time Series Prediction with Recurrent Neural Networks

## Weather prediction with Graphical RNN

**Algorithm 1** Graphical RNN model: forward pass

```

Require:  $V, \mathcal{C}, X, Y, E, F, \mathcal{L}, \Theta, inroll = I$ 
1:  $\forall u \in V, h_{u,0,t} = \vec{0} \in \mathbb{R}^d$ 
2:  $loss = 0$ 
3: for  $t = [1, 2, 3, \dots, T]$  do
4:   for all  $u \in V$  do
5:      $h_{u,t,0} = h_{u,t-1,I}$ 
6:   for  $i = [1, 2, 3, \dots, I]$  do
7:     for all  $u \in V$  do
8:       for all  $k \in [K]$  do
9:          $s_{u,k,t,i} = f_k(\{h_{v,t,i-1} | v \in N_{u,k}\})$ 
10:         $s_{u,t,i} = concat(\{s_{u,k,t,i}\}_{k=1}^K)$ 
11:         $x_{u,t,i} = concat(X_{g(u)}(u, t), s_{u,t,i})$ 
12:        for all  $u \in V$  do
13:           $(\hat{y}_{u,t,i}, h_{u,t,i}) = RNN_{\theta_g(u)}(x_{u,t,i}, h_{u,t,i-1})$ 
14:        for all  $u \in V$  do
15:           $loss = loss + \mathcal{L}_{g(u)}(Y_{g(u)}(u, t), \hat{y}_{u,t,i})$ 

```

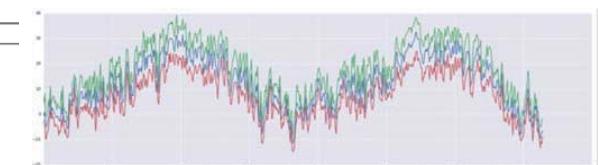


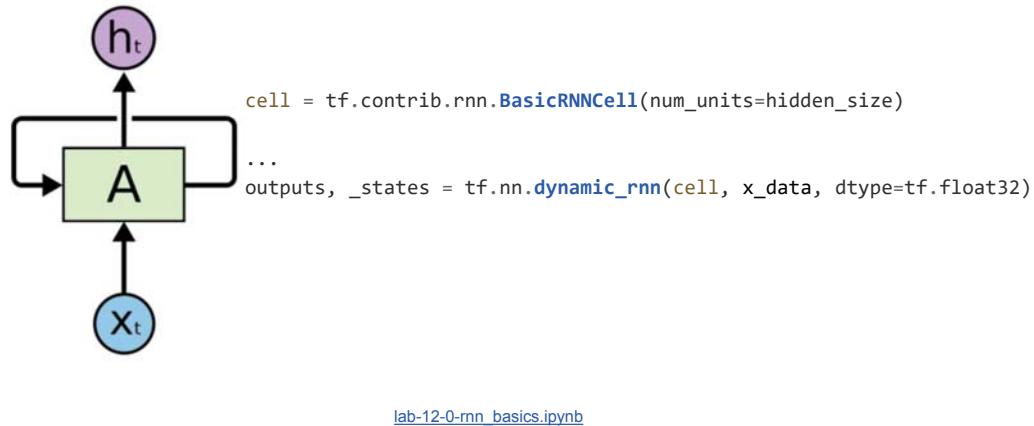
Figure 2: Sample input temperatures

Model type	Details	Train L2	Test L2
Simple	Steady state prediction	100.00	100.00
	Linear	97.22	97.34
RNN	LSTM, tanh	90.48	92.75
	iRNN, ReLU	87.62	91.01
gRNN	iRNN, ReLU, summation summary	84.41	89.38
	iRNN, ReLU, average summary	81.80	87.33

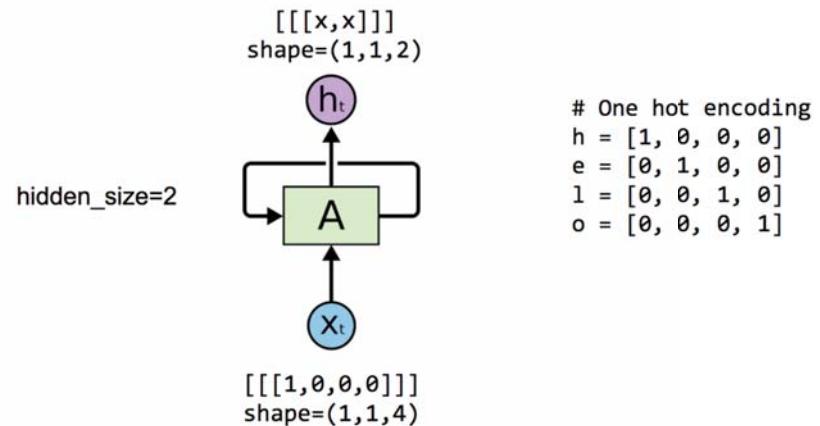
Table 2: Results on weather prediction task (% of steady state baseline).

Source: <https://arxiv.org/pdf/1612.05054.pdf>  
Graphical RNN Models

## RNN in TensorFlow



One node: 4 (*input\_dim*) in 2 (*hidden\_size*)



One node: 4 (*input\_dim*) in 2 (*hidden\_size*)

```

# One cell RNN input_dim (4) -> output_dim (2)
hidden_size = 2
cell = tf.contrib.rnn.BasicRNNCell(num_units=hidden_size)

x_data = np.array([[1,0,0,0]], dtype=np.float32)
outputs, _states = tf.nn.dynamic_rnn(cell, x_data, dtype=tf.float32)

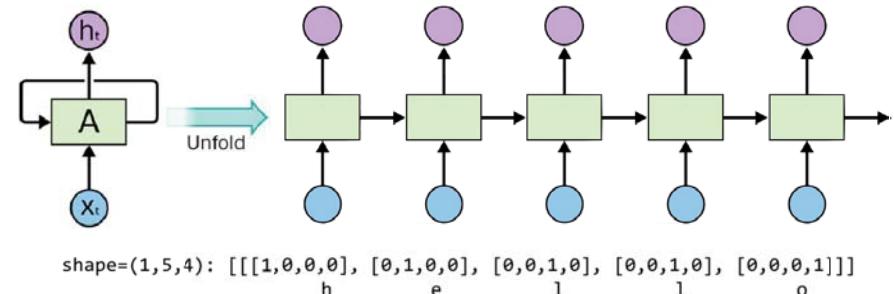
sess.run(tf.global_variables_initializer())
pp.pprint(outputs.eval())
  
```

```
array([[-0.42409304,  0.64651132]])
```

[lab-12-0-rnn\\_basics.ipynb](#)

Hidden\_size=2  
sequence\_length=5

shape=(1,5,2):  $[[[x, x], [x, x], [x, x], [x, x], [x, x]]]$



Unfolding to n sequences

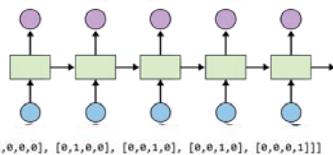
[lab-12-0-rnn\\_basics.ipynb](#)

## Unfolding to n sequences

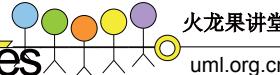
```
# One cell RNN input_dim (4) -> output_dim (2). sequence: 5
hidden_size = 2
cell = tf.contrib.rnn.BasicRNNCell(num_units=hidden_size)
x_data = np.array([[h, e, l, l, o]], dtype=np.float32)
print(x_data.shape)
pp pprint(x_data)
outputs, states = tf.nn.dynamic_rnn(cell, x_data, dtype=tf.float32)
sess.run(tf.global_variables_initializer())
pp pprint(outputs.eval())
```

Hidden\_size=2  
sequence\_length=5

shape=(1,5,2): [[[x,x], [x,x], [x,x], [x,x], [x,x]]]



shape=(1,5,4): [[[1,0,0,0], [0,1,0,0], [0,0,1,0], [0,0,1,0], [0,0,0,1]]]  
lab-12-0-rnn\_basics.ipynb



火龙果讲堂  
uml.org.cn

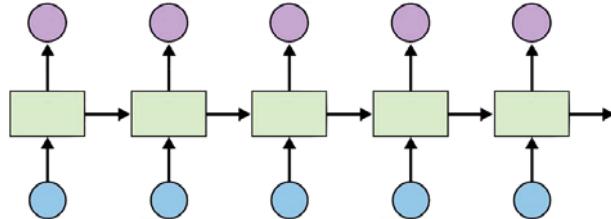
```
# One hot encoding
h = [1, 0, 0, 0]
e = [0, 1, 0, 0]
l = [0, 0, 1, 0]
o = [0, 0, 0, 1]
```

```
X_data = array
([[[ 1., 0., 0., 0.],
 [ 0., 1., 0., 0.],
 [ 0., 0., 1., 0.],
 [ 0., 0., 0., 1.],
 [ 0., 0., 0., 1.]]], dtype=float32)
```

```
Outputs = array
([[[ 0.19709368, 0.24918222,
 [-0.11721198, 0.1784237 ],
 [-0.35297349, -0.66278851],
 [-0.70915914, -0.58334434],
 [-0.38886023, 0.47304463]], dtype=float32])
```

Hidden\_size=2  
sequence\_length=5  
batch\_size=3

```
shape=(3,5,2): [[[x,x], [x,x], [x,x], [x,x], [x,x]],
 [[x,x], [x,x], [x,x], [x,x], [x,x]],
 [[x,x], [x,x], [x,x], [x,x], [x,x]]]
```



shape=(3,5,4): [[[1,0,0,0], [0,1,0,0], [0,0,1,0], [0,0,1,0], [0,0,0,1]], # hello
 [[0,1,0,0], [0,0,0,1], [0,0,1,0], [0,0,1,0], [0,0,1,0]], # eolll
 [[0,0,1,0], [0,0,1,0], [0,1,0,0], [0,1,0,0], [0,0,1,0]]] # lleel

## Batching input



火龙果讲堂  
uml.org.cn

## Batching input

```
# One cell RNN input_dim (4) -> output_dim (2). sequence: 5, batch 3
```

```
# 3 batches 'hello', 'eolll', 'lleel'
x_data = np.array([[h, e, l, l, o],
 [e, o, l, l, l],
 [l, l, e, e, l]], dtype=np.float32)
pp pprint(x_data)

cell = rnn.BasicLSTMCell(num_units=2, state_is_tuple=True)
outputs, _states = tf.nn.dynamic_rnn(cell, x_data,
 dtype=tf.float32)
sess.run(tf.global_variables_initializer())
pp pprint(outputs.eval())
```

shape=(3,5,2): [[[x,x], [x,x], [x,x], [x,x], [x,x]],
 [[x,x], [x,x], [x,x], [x,x], [x,x]],
 [[x,x], [x,x], [x,x], [x,x], [x,x]]]

Hidden\_size=2  
sequence\_length=5  
batch\_size=3

shape=(1,5,4): [[[1,0,0,0], [0,1,0,0], [0,0,1,0], [0,0,1,0], [0,0,0,1]], # hello
 [[0,1,0,0], [0,0,0,1], [0,0,1,0], [0,0,1,0], [0,0,1,0]], # eolll
 [[0,0,1,0], [0,0,1,0], [0,1,0,0], [0,1,0,0], [0,0,1,0]]] # lleel

lab-12-0-rnn\_basics.ipynb

```
array([[[ 1., 0., 0., 0.],
 [ 0., 1., 0., 0.],
 [ 0., 0., 1., 0.],
 [ 0., 0., 1., 0.],
 [ 0., 0., 0., 1.]]])
```

```
[[[ 0., 1., 0., 0.],
 [ 0., 0., 0., 1.],
 [ 0., 0., 1., 0.],
 [ 0., 0., 1., 0.],
 [ 0., 0., 1., 0.]]]
```

```
[[[ 0., 0., 1., 0.],
 [ 0., 0., 1., 0.],
 [ 0., 1., 0., 0.],
 [ 0., 1., 0., 0.],
 [ 0., 0., 1., 0.]]]
```

## Batching input



火龙果讲堂  
uml.org.cn

```
# One cell RNN input_dim (4) -> output_dim (2). sequence: 5, batch 3
```

```
# 3 batches 'hello', 'eolll', 'lleel'
x_data = np.array([[h, e, l, l, o],
 [e, o, l, l, l],
 [l, l, e, e, l]], dtype=np.float32)
pp pprint(x_data)
```

array([[[ 1., 0., 0., 0.]]])

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

[

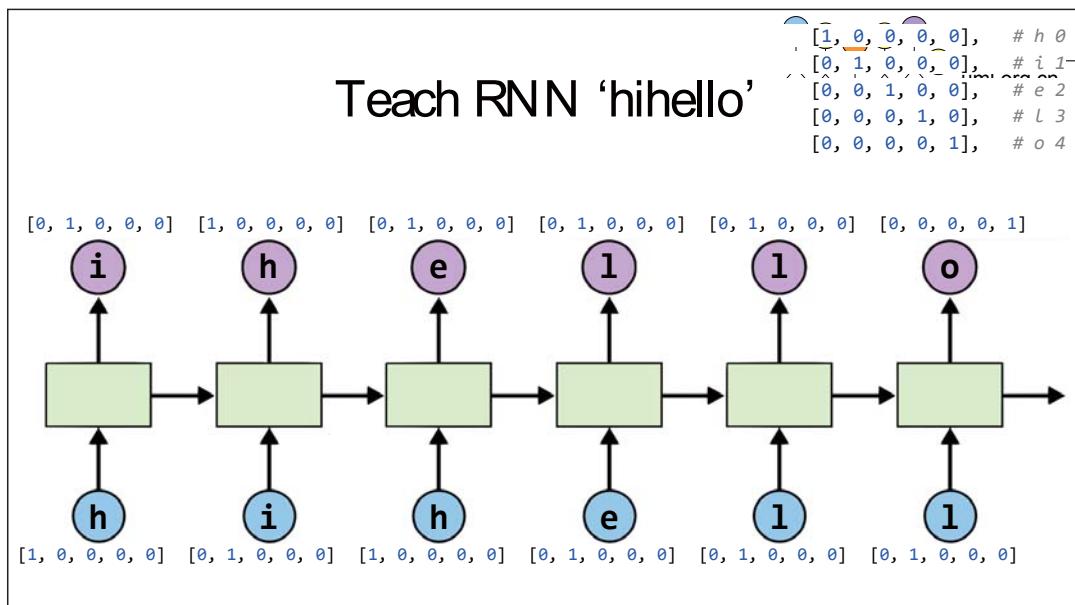
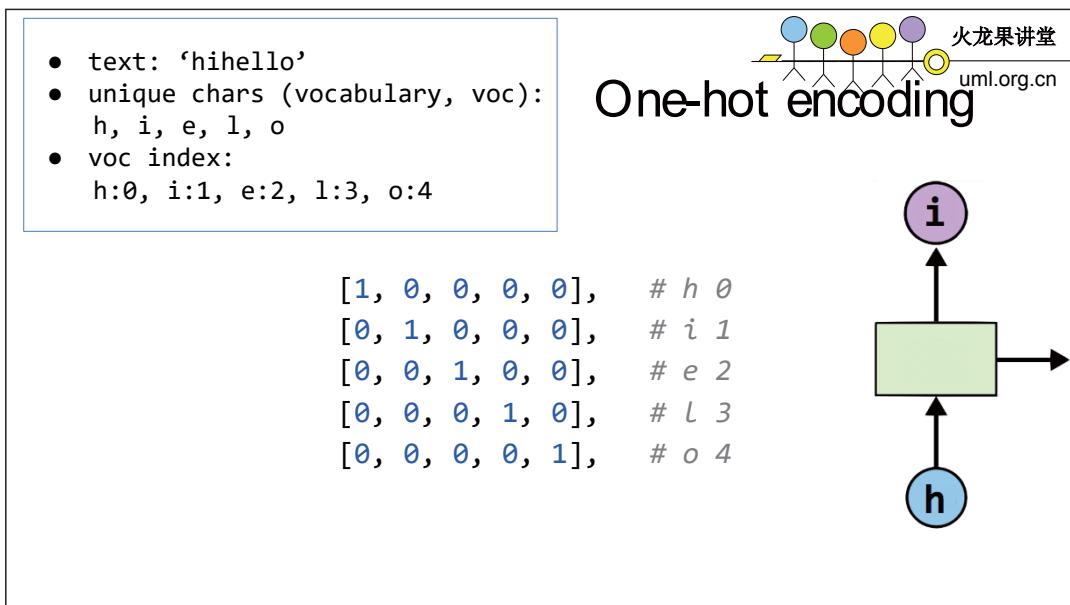
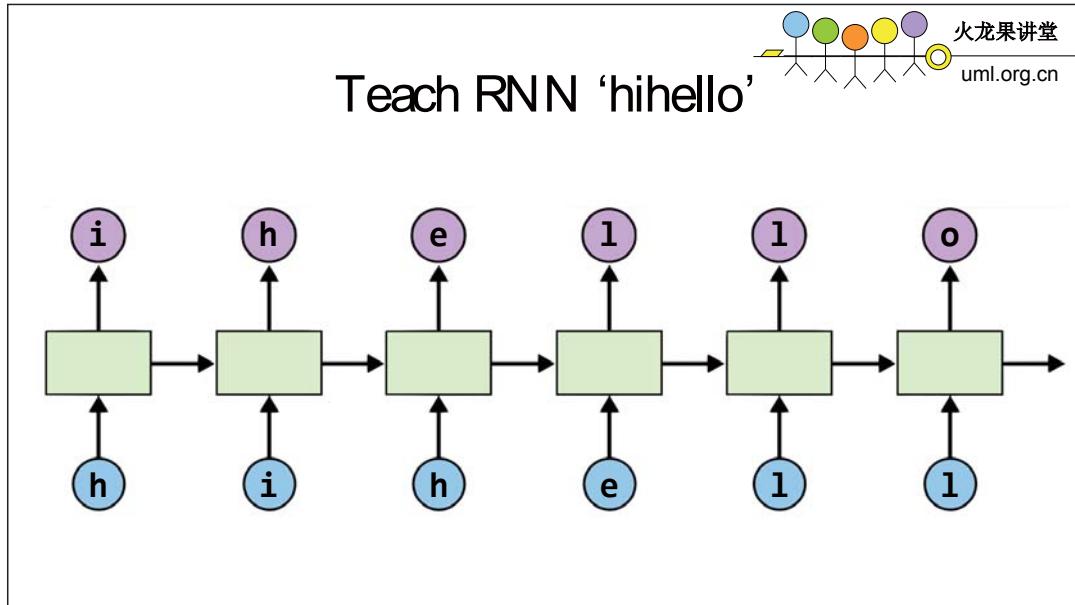
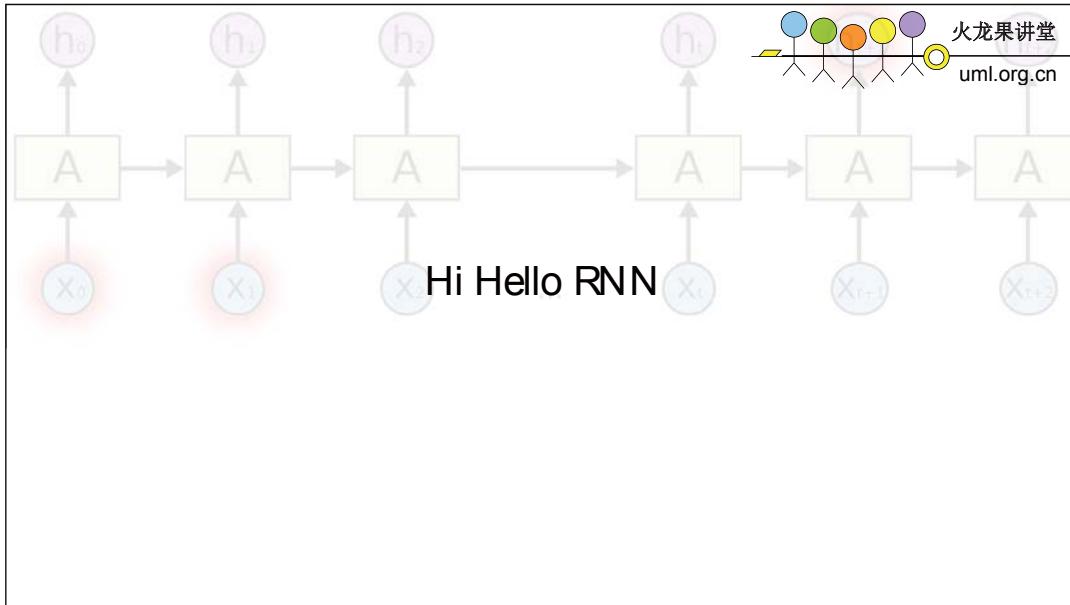
[

[

[

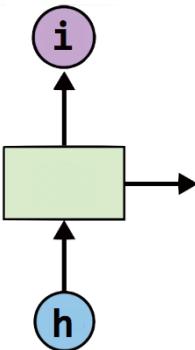
[

[



## Creating rnn cell

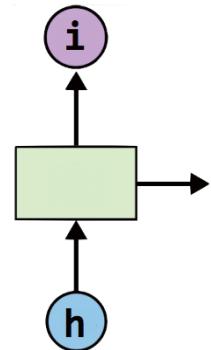
```
# RNN model  
rnn_cell = rnn_cell.BasicRNNCell(rnn_size)
```



## Creating rnn cell

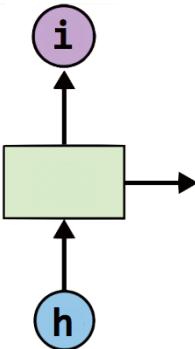
```
# RNN model  
rnn_cell = rnn_cell.BasicRNNCell(rnn_size)
```

```
rnn_cell = rnn_cell.BasicLSTMCell(rnn_size)  
rnn_cell = rnn_cell.GRUCell(rnn_size)
```



## Execute RNN

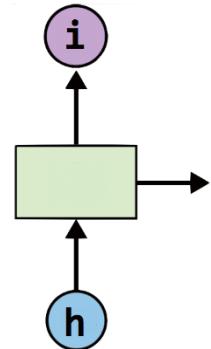
```
# RNN model  
rnn_cell = rnn_cell.BasicRNNCell(rnn_size) ①  
  
outputs, _states = tf.nn.dynamic_rnn(  
    rnn_cell,  
    x, ②  
    initial_state=initial_state,  
    dtype=tf.float32)
```



- text: 'helloworld'
- unique chars (vocabulary, voc): h, i, e, l, o
- voc index:  
h:0, i:1, e:2, l:3, o:4

## RNN parameters

```
hidden_size = 5      # output from the LSTM  
input_dim = 5        # one-hot size  
batch_size = 1       # one sentence  
sequence_length = 6  # |ihello| == 6
```



## Data creation

```

idx2char = ['h', 'i', 'e', 'l', 'o'] # h=0, i=1, e=2, l=3, o=4
x_data = [[0, 1, 0, 2, 3, 3]] # hihell
x_one_hot = [[[1, 0, 0, 0, 0], # h 0
              [0, 1, 0, 0, 0], # i 1
              [1, 0, 0, 0, 0], # h 0
              [0, 0, 1, 0, 0], # e 2
              [0, 0, 0, 1, 0], # l 3
              [0, 0, 0, 1, 0]]] # l 3

y_data = [[1, 0, 2, 3, 3, 4]] # ihello
X = tf.placeholder(tf.float32,
                   [None, sequence_length, input_dim]) # X one-hot
Y = tf.placeholder(tf.int32, [None, sequence_length]) # Y Label

```

[lab-12-1-hello-rnn.py](#)



火龙果讲堂  
uml.org.cn

## Feed to RNN

```

x_one_hot = [[[1, 0, 0, 0, 0], # h 0
              [0, 1, 0, 0, 0], # i 1
              [1, 0, 0, 0, 0], # h 0
              [0, 0, 1, 0, 0], # e 2
              [0, 0, 0, 1, 0], # l 3
              [0, 0, 0, 1, 0]]] # l 3

y_data = [[1, 0, 2, 3, 3, 4]] # ihello

X = tf.placeholder(
    tf.float32, [None, sequence_length, hidden_size]) # X one-hot
Y = tf.placeholder(tf.int32, [None, sequence_length]) # Y Label

cell = tf.contrib.rnn.BasicLSTMCell(num_units=hidden_size,
                                     state_is_tuple=True)
initial_state = cell.zero_state(batch_size, tf.float32)
outputs, _states = tf.nn.dynamic_rnn(
    cell, X, initial_state=initial_state, dtype=tf.float32)

```

[lab-12-1-hello-rnn.py](#)

## Cost: sequence\_loss

```

# [batch_size, sequence_length]
y_data = tf.constant([[1, 1, 1]])

# [batch_size, sequence_length, emb_dim ]
prediction = tf.constant([[0, 1], [1, 0], [0, 1]]], dtype=tf.float32)

# [batch_size * sequence_length]
weights = tf.constant([[1, 1, 1]]], dtype=tf.float32)

sequence_loss = tf.contrib.seq2seq.sequence_loss(prediction, y_data, weights)
sess.run(tf.global_variables_initializer())
print("Loss: ", sequence_loss.eval())

```

Loss: 0.646595

[lab-12-0-rnn basics.ipynb](#)



火龙果讲堂  
uml.org.cn

## Cost: sequence\_loss

```

# [batch_size, sequence_length]
y_data = tf.constant([[1, 1, 1]])

# [batch_size, sequence_length, emb_dim ]
prediction1 = tf.constant([[0, 1], [0, 1], [0, 1]]], dtype=tf.float32)
prediction2 = tf.constant([[1, 0], [1, 0], [1, 0]]], dtype=tf.float32)
prediction3 = tf.constant([[0, 1], [1, 0], [0, 1]]], dtype=tf.float32)

# [batch_size * sequence_length]
weights = tf.constant([[1, 1, 1]]], dtype=tf.float32)

sequence_loss1 = tf.contrib.seq2seq.sequence_loss(prediction1, y_data, weights)
sequence_loss2 = tf.contrib.seq2seq.sequence_loss(prediction2, y_data, weights)
sequence_loss3 = tf.contrib.seq2seq.sequence_loss(prediction3, y_data, weights)

sess.run(tf.global_variables_initializer())
print("Loss1: ", sequence_loss1.eval(),
      "Loss2: ", sequence_loss2.eval(),
      "Loss3: ", sequence_loss3.eval())

```

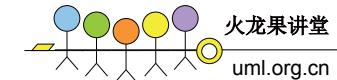
Loss1: 0.313262 Loss2: 1.31326 Loss3: 0.646595

[lab-12-0-rnn basics.ipynb](#)

## Cost: sequence\_loss

```
outputs, _states = tf.nn.dynamic_rnn(  
    cell, X, initial_state=initial_state, dtype=tf.float32)  
weights = tf.ones([batch_size, sequence_length])  
  
sequence_loss = tf.contrib.seq2seq.sequence_loss(  
    logits=outputs, targets=Y, weights=weights)  
loss = tf.reduce_mean(sequence_loss)  
train = tf.train.AdamOptimizer(learning_rate=0.1).minimize(loss)
```

[lab-12-1-hello-rnn.py](#)



## Training

```
prediction = tf.argmax(outputs, axis=2)  
  
with tf.Session() as sess:  
    sess.run(tf.global_variables_initializer())  
    for i in range(2000):  
        l, _ = sess.run([loss, train], feed_dict={X: x_one_hot, Y: y_data})  
        result = sess.run(prediction, feed_dict={X: x_one_hot})  
        print(i, "loss:", l, "prediction: ", result, "true Y: ", y_data)  
  
        # print char using dic  
        result_str = [idx2char[c] for c in np.squeeze(result)]  
        print("\tPrediction str: ", ''.join(result_str))
```

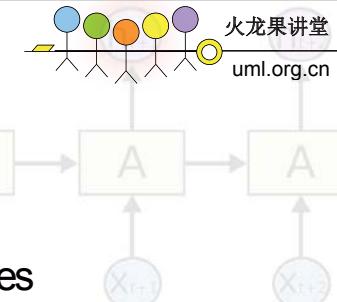
[lab-12-1-hello-rnn.py](#)

## Results

```
prediction = tf.argmax(outputs, axis=2)  
  
with tf.Session() as sess:  
    sess.run(tf.global_variables_initializer())  
    for i in range(2000):  
        l, _ = sess.run([loss, train], feed_dict={X: x_one_hot, Y: y_data})  
        result = sess.run(prediction, feed_dict={X: x_one_hot})  
        print(i, "loss:", l, "prediction: ", result, "true Y: ", y_data)  
  
        # print char using dic  
        result_str = [idx2char[c] for c in np.squeeze(result)]  
        print("\tPrediction str: ", ''.join(result_str))
```

```
0 loss: 1.55474 prediction: [[3 3 3 3 4 4]] true Y: [[1, 0, 2, 3, 3, 4]] Prediction str: 111100  
1 loss: 1.55081 prediction: [[3 3 3 3 4 4]] true Y: [[1, 0, 2, 3, 3, 4]] Prediction str: 111100  
2 loss: 1.54704 prediction: [[3 3 3 3 4 4]] true Y: [[1, 0, 2, 3, 3, 4]] Prediction str: 111100  
3 loss: 1.54342 prediction: [[3 3 3 3 4 4]] true Y: [[1, 0, 2, 3, 3, 4]] Prediction str: 111100  
...  
1998 loss: 0.75305 prediction: [[1 0 2 3 3 4]] true Y: [[1, 0, 2, 3, 3, 4]] Prediction str: ihello  
1999 loss: 0.752973 prediction: [[1 0 2 3 3 4]] true Y: [[1, 0, 2, 3, 3, 4]] Prediction str:  
ihello
```

[lab-12-1-hello-rnn.py](#)



RNN with long sequences

## Manual data creation

```

idx2char = [ 'h', 'i', 'e', 'l', 'o' ]
x_data = [[0, 1, 0, 2, 3, 3]]      # hihell
x_one_hot = [[[1, 0, 0, 0, 0],      # h 0
              [0, 1, 0, 0, 0],      # i 1
              [1, 0, 0, 0, 0],      # h 0
              [0, 0, 1, 0, 0],      # e 2
              [0, 0, 0, 1, 0],      # l 3
              [0, 0, 0, 1, 0]]]]    # l 3

y_data = [[1, 0, 2, 3, 3, 4]]    # ihello

```

lab-12-1-hello-rnn.py

# Better data creation

```
sample = " if you want you"
idx2char = list(set(sample)) # index -> char
char2idx = {c: i for i, c in enumerate(idx2char)} # char -> idx

sample_idx = [char2idx[c] for c in sample] # char to index
x_data = [sample_idx[:-1]] # X data sample (0 ~ n-1) hello: hell
y_data = [sample_idx[1:]] # Y label sample (1 ~ n) hello: ello

X = tf.placeholder(tf.int32, [None, sequence_length]) # X data
Y = tf.placeholder(tf.int32, [None, sequence_length]) # Y label

X one hot = tf.one hot(X, num classes) # one hot: 1 -> 0 1 0 0 0 0 0 0 0 0 0
```

lab-12-2-char-seq-rnn.py

## Hyper parameters

```

sample = " if you want you"
idx2char = list(set(sample)) # index -> char
char2idx = {c: i for i, c in enumerate(idx2char)} # char -> idx

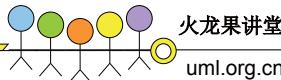
# hyper parameters
dic_size = len(char2idx) # RNN input size (one hot size)
rnn_hidden_size = len(char2idx) # RNN output size
num_classes = len(char2idx) # final output size (RNN or softmax, etc.)
batch_size = 1 # one sample data, one batch
sequence_length = len(sample) - 1 # number of lstm unfolding (unit #)

```

# LSTM and Loss

[lab-12-2-char-seq-rnn.py](#)

## Training and Results



```
with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    for i in range(3000):
        l, _ = sess.run([loss, train], feed_dict={X: x_data, Y: y_data})
        result = sess.run(prediction, feed_dict={X: x_data})
        # print char using dic
        result_str = [idx2char[c] for c in np.squeeze(result)]
        print(i, "loss:", l, "Prediction:", ''.join(result_str))
```

```
0 loss: 2.29895 Prediction: nnuffuunnuuuyuy
1 loss: 2.29675 Prediction: nnuffuunnuuuyuy
...
1418 loss: 1.37351 Prediction: if you want you
1419 loss: 1.37331 Prediction: if you want you
```

[lab-12-2-char-seq-rnn.py](#)

## Really long sentence?



```
sentence = ("if you want to build a ship, don't drum up people together to "
            "collect wood and don't assign them tasks and work, but rather "
            "teach them to long for the endless immensity of the sea.")
```

[lab-12-4-mn\\_long\\_char.py](#)

## Really long sentence?



```
sentence = ("if you want to build a ship, don't drum up people together to "
            "collect wood and don't assign them tasks and work, but rather "
            "teach them to long for the endless immensity of the sea.")
```

```
# training dataset
0 if you wan -> f you want
1 f you want -> you want
2 you want -> you want t
3 you want t -> ou want to
```

```
168 of the se -> of the sea
169 of the sea -> f the sea.
```

[lab-12-4-rnn\\_long\\_char.py](#)

## Making dataset

```
char_set = list(set(sentence))
char_dic = {w: i for i, w in enumerate(char_set)}
```

```
dataX = []
dataY = []
for i in range(0, len(sentence) - seq_length):
    x_str = sentence[i:i + seq_length]
    y_str = sentence[i + 1: i + seq_length + 1]
    print(i, x_str, '->', y_str)

    x = [char_dic[c] for c in x_str] # x str to index
    y = [char_dic[c] for c in y_str] # y str to index

    dataX.append(x)
    dataY.append(y)
```

```
# training dataset
0 if you wan -> f you want
1 f you want -> you want
2 you want -> you want t
3 you want t -> ou want to

168 of the se -> of the sea
169 of the sea -> f the sea.
```

[lab-12-4-mn\\_long\\_char.py](#)

## RNN parameters

```
char_set = list(set(sentence))
char_dic = {w: i for i, w in enumerate(char_set)}

data_dim = len(char_set)
hidden_size = len(char_set)
num_classes = len(char_set)
seq_length = 10 # Any arbitrary number

batch_size = len(datax)
```

[lab-12-4-rnn\\_long\\_char.py](#)



```
# training dataset
0 if you wan -> f you want
1 f you want -> you want
2 you want -> you want t
3 you want t -> ou want to

168 of the se -> of the sea
169 of the sea -> f the sea.
```

## Stacked RNN

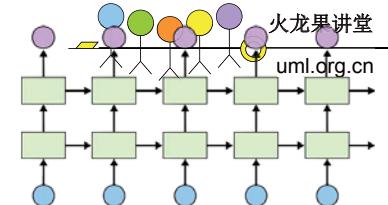
```
X = tf.placeholder(tf.int32, [None, seq_length])
Y = tf.placeholder(tf.int32, [None, seq_length])

# One-hot encoding
X_one_hot = tf.one_hot(X, num_classes)
print(X_one_hot) # check out the shape

# Make a lstm cell with hidden_size (each unit output vector size)
cell = rnn.BasicLSTMCell(hidden_size, state_is_tuple=True)
cell = rnn.MultiRNNCell([cell] * 2, state_is_tuple=True)

# outputs: unfolding size x hidden size, state = hidden size
outputs, _states = tf.nn.dynamic_rnn(cell, X_one_hot, dtype=tf.float32)
```

[lab-12-4-mn\\_long\\_char.py](#)



## Softmax (FC) in Deep CNN

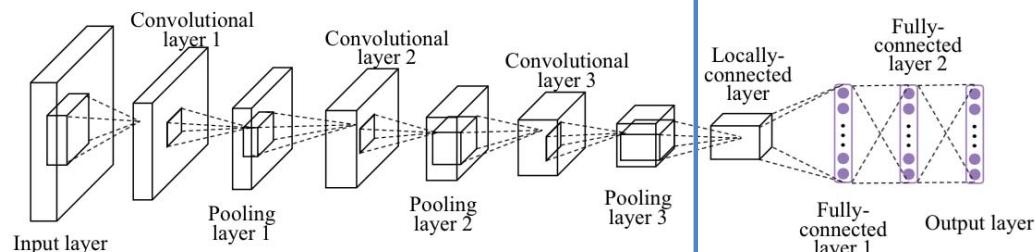
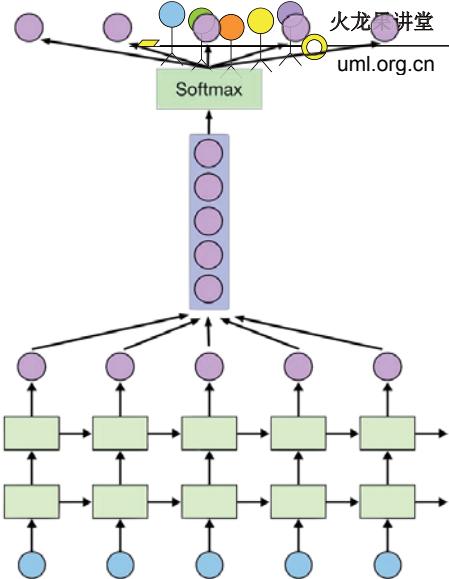


Image credit: [http://personal.ie.cuhk.edu.hk/~ccloy/project\\_target\\_code/index.html](http://personal.ie.cuhk.edu.hk/~ccloy/project_target_code/index.html)

## Softmax

```
outputs, _ = tf.nn.dynamic_rnn(cell, X_one_hot)
```

[lab-12-4-mn\\_long\\_char.py](#)



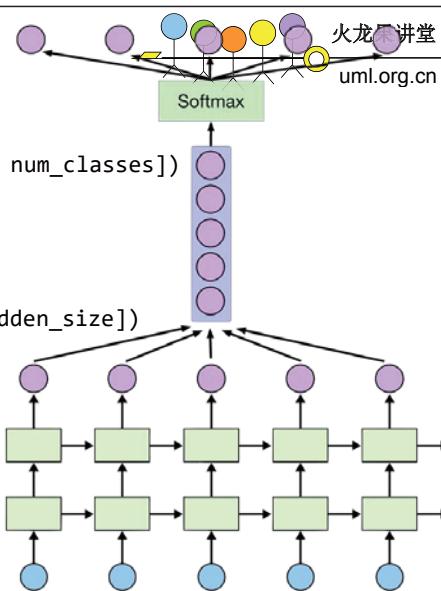
## Softmax

```
outputs = tf.reshape(outputs,
    [batch_size, seq_length, num_classes])
```

```
X_for_softmax = tf.reshape(outputs,
    [-1, hidden_size])
```

```
outputs, _ = tf.nn.dynamic_rnn(cell, X_one_hot)
```

[lab-12-4-rnn\\_long\\_char.py](#)



## Softmax

```
# (optional) softmax layer
```

```
X_for_softmax = tf.reshape(outputs, [-1, hidden_size])
```

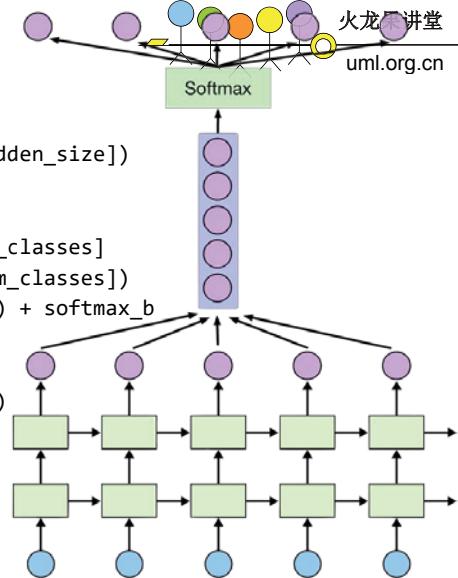
```
softmax_w = tf.get_variable("softmax_w",
    [hidden_size, num_classes])
```

```
softmax_b = tf.get_variable("softmax_b", [num_classes])
```

```
outputs = tf.matmul(X_for_softmax, softmax_w) + softmax_b
```

```
outputs = tf.reshape(outputs,
    [batch_size, seq_length, num_classes])
```

[lab-12-4-mn\\_long\\_char.py](#)



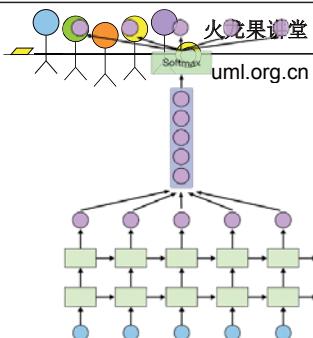
## Loss

```
# reshape out for sequence_loss
outputs = tf.reshape(outputs,
    [batch_size, seq_length, num_classes])
# All weights are 1 (equal weights)
weights = tf.ones([batch_size, seq_length])

sequence_loss = tf.contrib.seq2seq.sequence_loss(
    logits=outputs, targets=Y, weights=weights)
mean_loss = tf.reduce_mean(sequence_loss)

train_op =
    tf.train.AdamOptimizer(learning_rate=0.1).minimize(mean_loss)
```

[lab-12-4-rnn\\_long\\_char.py](#)



## Training and print results

```
sess = tf.Session()
sess.run(tf.global_variables_initializer())
```

```
for i in range(500):
    _, l, results = sess.run(
        [train_op, mean_loss, outputs],
        feed_dict={X: dataX, Y: dataY})
```

```
for j, result in enumerate(results):
    index = np.argmax(result, axis=1)
    print(i, j, ''.join([char_set[t] for t in index]), l)
```

0 167 tttttttt 3.23111

0 168 tttttttt 3.23111

0 169 tttttttt 3.23111

499 167 oof the se 0.229306

499 168 tf the sea 0.229306

499 169 n the sea. 0.229306

[lab-12-4-mn\\_long\\_char.py](#)

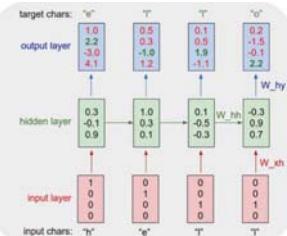
# Training and print results

```
# Let's print the last char of each result to check it works
results = sess.run(outputs, feed_dict={X: dataX})
for j, result in enumerate(results):
    index = np.argmax(result, axis=1)
    if j is 0: # print all for the first result to make a sentence
        print(''.join([char_set[t] for t in index]), end=' ')
    else:
        print(char_set[index[-1]], end='')
```

g you want to build a ship, don't drum up people together to collect wood and don't assign them tasks and work, but rather teach them to long for the endless immensity of the sea.

[lab-12-4-rnn\\_long\\_char.py](#)

## char-rnn



### Shakespeare

It looks like we can learn to spell English words. But how about if there is more structure and style in the data? To examine this I downloaded all the works of Shakespeare and concatenated them into a single (4.4MB) file. We can now afford to train a larger network, in this case lets try a 3-layer RNN with 512 hidden nodes on each layer. After we train the network for a few hours we obtain samples such as:

PANDARUS:  
Alas, I think he shall be come approached and the day  
When little strain would be attain'd into being never fed,  
And who is but a chain and subjects of his death,  
I should not sleep.

Second Senator:  
They are away this miseries, produced upon my soul,  
Breaking and strongly should be buried, when I perish  
The earth and thoughts of many states.

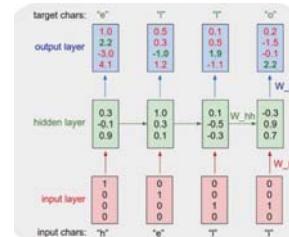
DUKE VINCENTIO:  
Well, your wit is in the care of side and that.

Second Lord:  
They would be ruled after this chamber, and  
my fair nues begun out of the fact, to be conveyed,  
Whose noble souls I'll have the heart of the wars.

Clown:  
Come, sir, I will make did behold your worship.

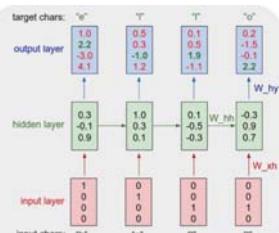
VIOLA:  
I'll drink it.

## char/word rnn (char/word level n to n model)



[char-rnn-tensorflow](#)

[word-rnn-tensorflow](#)



### Linux Source Code

I wanted to push structured data to its limit, so for the final challenge I decided to use code. In particular, I took all the source and header files found in the Linux repo on GitHub, concatenated all of them in a single giant file (474MB of C code) it was originally going to train only on the kernel but that by itself is only ~16MB). Then I trained several as-large-as-fits-on-my-GPU 3-layer LSTMs over a period of a few days. These models have about 10 million parameters, which is still on the lower end for RNN models. The results are superfun:

```
/*
 * Increment the size file of the new incorrect UI_FILTER group information
 * of the size generatively.
 */
static int indicate_policy(void)
{
    int error;
    if (fd == MARK_EPT) {
        /* The kernel blank will coold it to user space. */
        if (ss->segment < mem_total)
            unblock_graph_and_set_blocked();
        else
            iwt = 1;
        goto bail;
    }
    seqaddr = in_BB(in.addr);
    selector = seq / 16;
    setup_works = true;
    for (i = 0; i < blocks; i++) {
        seq = buf[i++];
        bpf = bd->bd.next + i * search;
        if (fd) {
            current = blocked;
        }
    }
    rv->name = "Getjbbreg";
    bprm_self_clearl(&rv->version);
    reg->new = blocks[(BPF_STATS << info->historidac)] | PPFMS_CLOABTHINC_SECONDS << 13;
    return seqtable;
}
```

RNN with time series data (stock)

Time series data

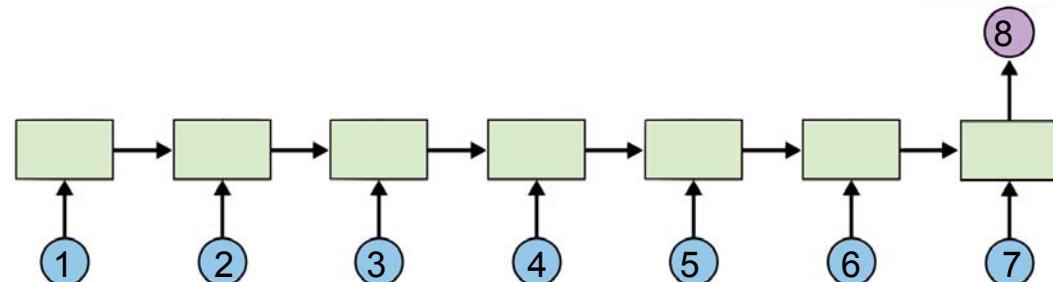


Time series data

Open	High	Low	Volume	Close
828.659973	833.450012	828.349976	1247700	831.659973
823.02002	828.070007	821.655029	1597800	828.070007
819.929993	824.400024	818.97998	1281700	824.159973
819.359985	823	818.469971	1304000	818.97998
819	823	816	1053600	820.450012
816	820.958984	815.48999	1198100	819.23999
811.700012	815.25	809.780029	1129100	813.669983
809.51001	810.659973	804.539978	989700	809.559998
807	811.840027	803.190002	1155300	808.380005

'data-02-stock\_daily.csv'

Many to one



lab-12-5-rnn\_stock\_prediction.py

```

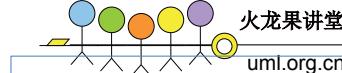
timesteps = seq_length = 7
data_dim = 5
output_dim = 1
# Open,High,Low,Close,Volume
xy = np.loadtxt('data-02-stock_daily.csv', delimiter=',')
xy = xy[::-1] # reverse order (chronically ordered)
xy = MinMaxScaler(xy)
x = xy
y = xy[:, [-1]] # Close as Label

dataX = []
dataY = []
for i in range(0, len(y) - seq_length):
    _x = x[i:i + seq_length]
    _y = y[i + seq_length] # Next close price
    print(_x, "->", _y)
    dataX.append(_x)
    dataY.append(_y)

lab-12-5-rnn_stock_prediction.py

```

## Reading data



火龙果讲堂 uml.org.cn

```

[ 0.18667876 0.20948057 0.20878184 0.
 0.21744815]

[ 0.30697388 0.31463414 0.21899367
 0.01247647 0.21698189]

[ 0.21914211 0.26390721 0.2246864
 0.45632338 0.22496747]

[ 0.23312993 0.23641916 0.16268272
 0.57017119 0.14744274]

[ 0.13431201 0.15175877 0.11617252
 0.39380658 0.13289962]

[ 0.13973232 0.17060429 0.15860382
 0.28173344 0.18171679]

[ 0.18933069 0.20057799 0.19187983
 0.29783096 0.2086465]

-> [ 0.14106001]

```



火龙果讲堂 uml.org.cn

```

# split to train and testing
train_size = int(len(dataY) * 0.7)
test_size = len(dataY) - train_size
trainX, testX = np.array(dataX[0:train_size]),
                  np.array(dataX[train_size:len(dataX)])
trainY, testY = np.array(dataY[0:train_size]),
                  np.array(dataY[train_size:len(dataY)])

# input placeholders
X = tf.placeholder(tf.float32, [None, seq_length, data_dim])
Y = tf.placeholder(tf.float32, [None, 1])

lab-12-5-rnn_stock_prediction.py

```

## Training and test datasets



火龙果讲堂 uml.org.cn

## LSTM and Loss

```

# input placeholders
X = tf.placeholder(tf.float32, [None, seq_length, data_dim])
Y = tf.placeholder(tf.float32, [None, 1])

cell = tf.contrib.rnn.BasicLSTMCell(num_units=output_dim, state_is_tuple=True)
outputs, _states = tf.nn.dynamic_rnn(cell, X, dtype=tf.float32)
Y_pred = outputs[:, -1] # We use the last cell's output

# cost/loss
loss = tf.reduce_sum(tf.square(Y_pred - Y)) # sum of the squares
# optimizer
optimizer = tf.train.GradientDescentOptimizer(0.01)
train = optimizer.minimize(loss)

lab-12-5-rnn_stock_prediction.py

```



火龙果讲堂 uml.org.cn

```

sess = tf.Session()
sess.run(tf.global_variables_initializer())

for i in range(1000):
    _, l = sess.run([train, loss],
                   feed_dict={X: trainX, Y: trainY})
    print(i, l)

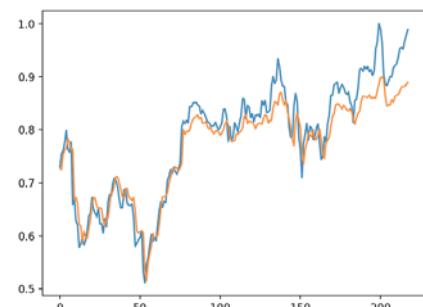
testPredict = sess.run(Y_pred, feed_dict={X: testX})

import matplotlib.pyplot as plt
plt.plot(testY)
plt.plot(testPredict)
plt.show()

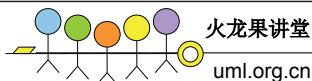
lab-12-5-rnn_stock_prediction.py

```

## Training and Results



## Exercise

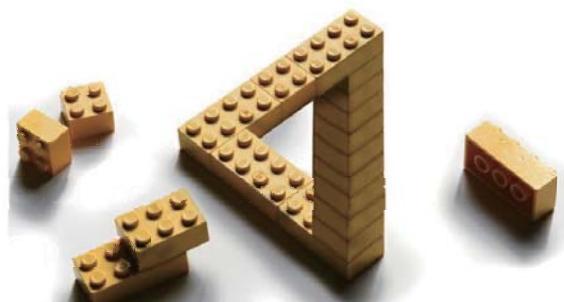


- Implement stock prediction using linear regression only
- Improve stock prediction results using more features such as keywords and/or sentiments in top news
- Implement sequence to sequence mode (translate a language to another)
  - Keras/klab-12-5-seq2seq.py

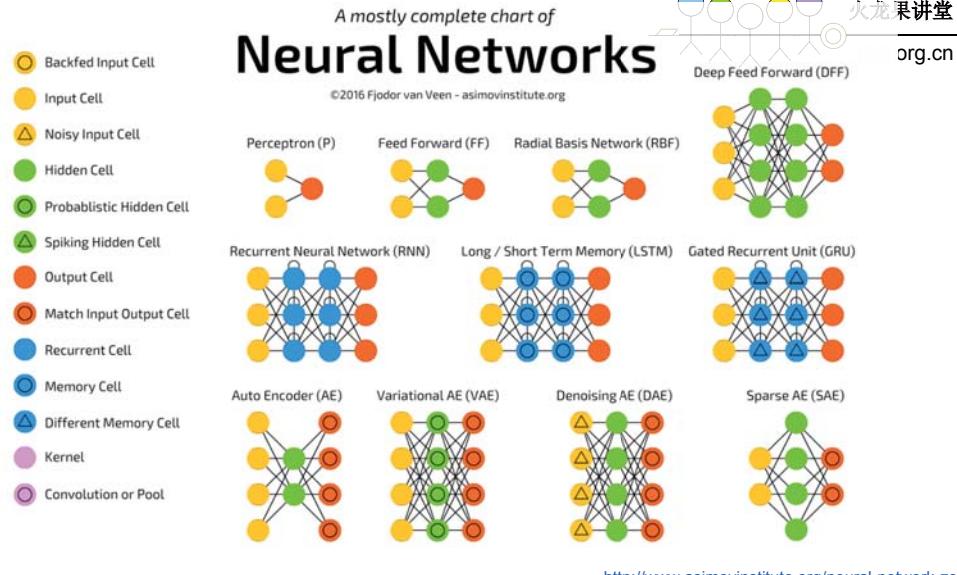
## Deep Learning

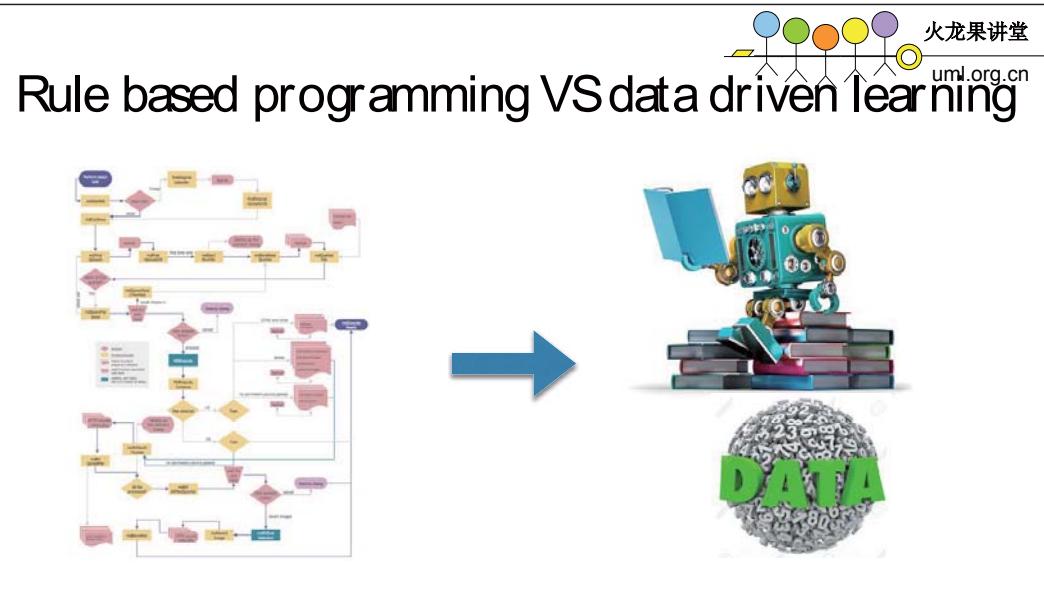
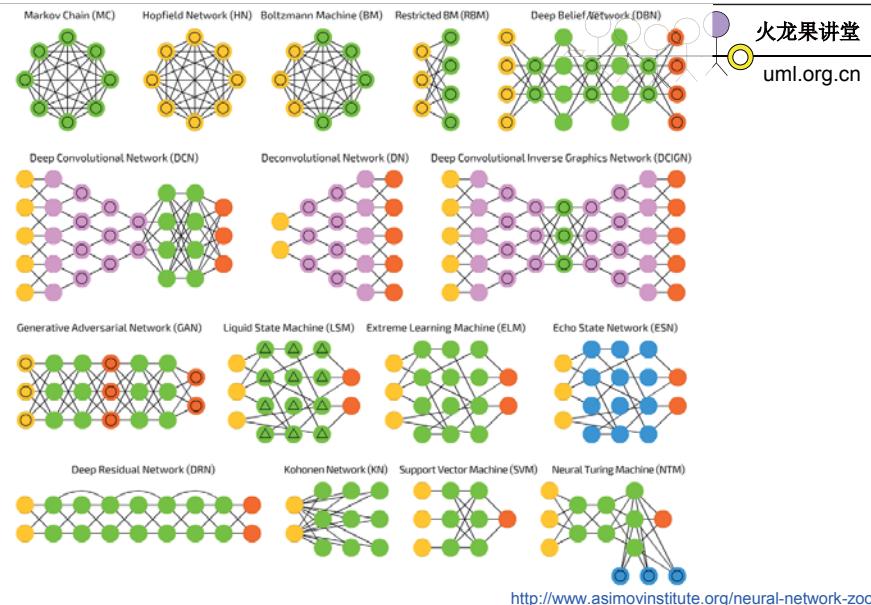
- DNN
- CNN
- RNN

'The only limit is your imagination'

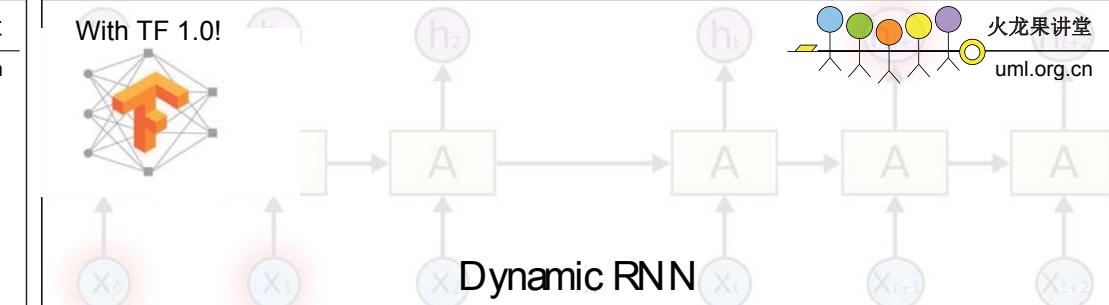


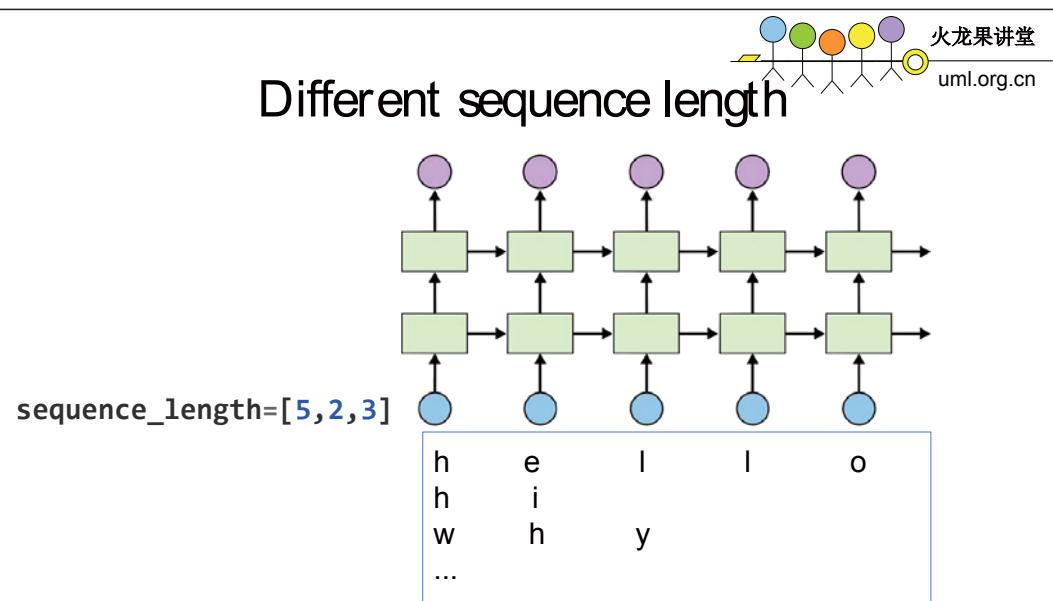
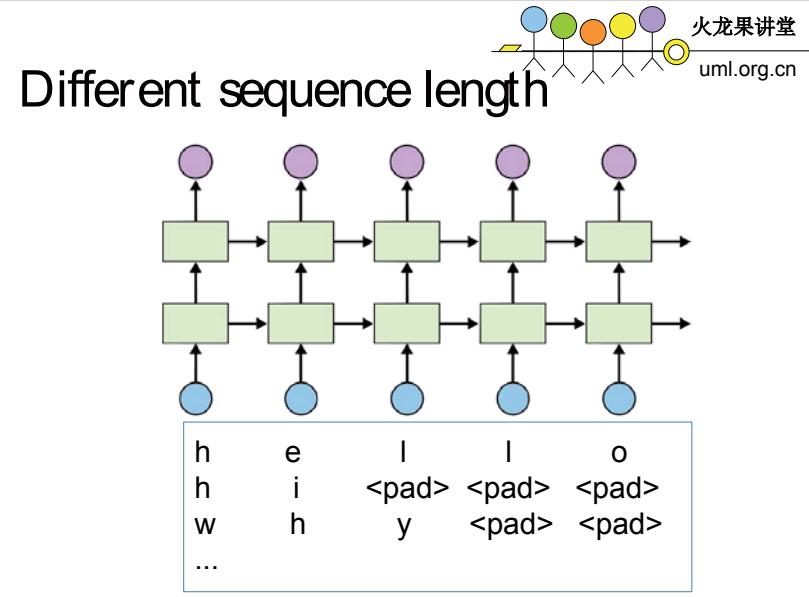
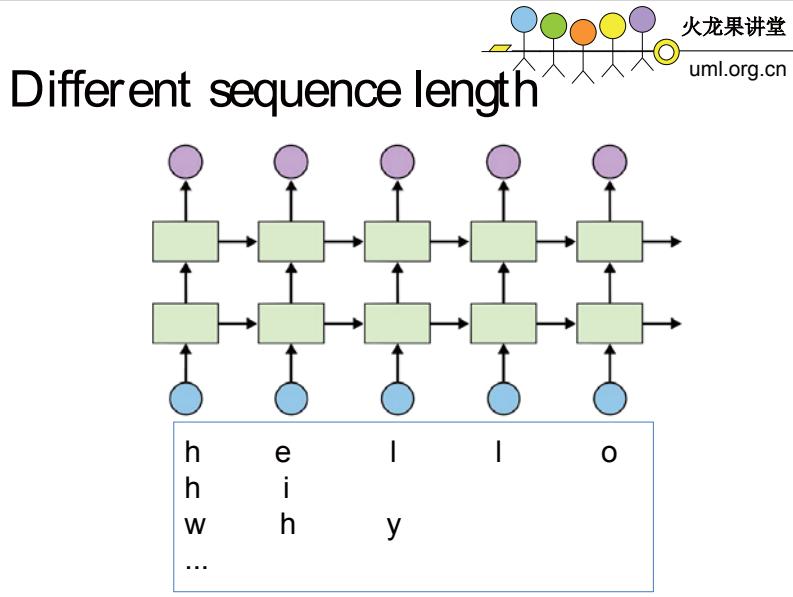
<http://itchyi.squarespace.com/thelatest/2012/5/17/the-only-limit-is-your-imagination.html>





## Backup Slides





火龙果讲堂  
uml.org.cn

### Dynamic RNN

```
# 3 batches 'hello', 'eolll', 'lleel'
x_data = np.array([[[[...]]], dtype=np.float32)

hidden_size = 2
cell = rnn.BasicLSTMCell(num_units=hidden_size,
                         state_is_tuple=True)
outputs, _states = tf.nn.dynamic_rnn(
    cell, x_data, sequence_length=[5,3,4],
    dtype=tf.float32)
sess.run(tf.global_variables_initializer())
print(outputs.eval())
```

```
array([[[[-0.17904168, -0.08053244],
        [-0.01294809,  0.01660814],
        [-0.05754048, -0.1368292 ],
        [-0.08655578, -0.20553185],
        [ 0.07297077, -0.21743253]],

       [[ 0.10272847,  0.06519825],
        [ 0.20188759, -0.05027055],
        [ 0.09514933, -0.16452041],
        [ 0.        ,  0.        ],
        [ 0.        ,  0.        ]],

       [[-0.04893036, -0.14655617],
        [-0.07947272, -0.20996611],
        [ 0.06466491, -0.02576563],
        [ 0.15087658,  0.05166111],
        [ 0.        ,  0.        ]]]
```

[lab-12-0-rnn\\_basics.ipynb](#)